

STAT395: Sampling

Richard Arnold et al.

2026-04-24

Contents

1	Introduction	7
1.1	Lecture Plan	7
1.2	Recommended reading	8
1.3	R Packages	9
2	Sampling	11
2.1	Random numbers	11
2.2	Sampling from categorical random variables	15
2.3	Sampling from discrete numerical random variables	20
2.4	Sampling from continuous random variables	24
2.5	Sampling from a finite collection	29
3	Sample Surveys	35
3.1	The Survey Process	35
3.2	Populations and Frames	37
3.3	Approaches to sampling	40
3.4	Example datasets	40
4	Statistical Inference	43
4.1	Inference in infinite populations	43
4.2	Inference in finite populations	47
4.3	The survey inference process	49
4.4	Example	53
5	Sample Designs	55
5.1	Census	55
5.2	Simple Random Sampling Without Replacement (SRSWOR)	55
5.3	Simple Random Sampling With Replacement (SRSWR)	59
5.4	Probability proportion to size Sampling (PPSWR)	60
5.5	Linear Systematic Random Sample (LSRS)	62
5.6	Stratified Simple Random Sample (STSRS)	65
5.7	One stage cluster sampling (1SC)	70
5.8	Two stage cluster sampling (2SC)	74
5.9	Complex multistage designs	77
6	Sample Designs in R	81
6.1	Census	82
6.2	Simple Random Sample WithOut Replacement (SRSWOR)	83
6.3	Probability Proportional to Size With Replacement (PPSWR)	84
6.4	Stratified Sample	84
6.5	One stage cluster design (1SC)	85

6.6	Two stage cluster design (2SC)	86
6.7	Stratified two stage cluster design	86
7	Survey Analysis	87
7.1	Continuous variables	89
7.2	Categorical Variables	99
8	Weight adjustments	107
8.1	Calibration	107
8.2	Replicate weights	110
8.3	Demonstration of ps SE improvement	120
9	Imputation	123
10	Package mice	153
11	Designing Surveys	155
11.1	Another interpretation of the Deff	155
11.2	Computing the Required Sample Size	156
	References	157
A	Placeholder Appendix	159

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

##
## Attaching package: 'kableExtra'

## The following object is masked from 'package:dplyr':
##
##   group_rows

## Loading required package: grid
## Loading required package: Matrix
## Loading required package: survival

##
## Attaching package: 'survey'

## The following object is masked from 'package:graphics':
##
##   dotchart

##
## Attaching package: 'srvyr'

## The following object is masked from 'package:kableExtra':
##
```

```
##      group_rows
## The following object is masked from 'package:stats':
##
##      filter
##
## Attaching package: 'sampling'
## The following objects are masked from 'package:survival':
##
##      cluster, strata
##
## Attaching package: 'mice'
## The following object is masked from 'package:stats':
##
##      filter
## The following objects are masked from 'package:base':
##
##      cbind, rbind
```


Chapter 1

Introduction

These are notes for the Sampling section of STAT395/455. The course overlaps to some extent with STAT392/439, but there is a greater emphasis on R coding, computation, summary and display.

1.1 Lecture Plan

Course Learning Objectives

- Describe the steps in the survey process, including the sources of survey error;
- Draw samples from a sample frame under a variety of sample designs, and compute survey weights;
- Carry out, interpret and present statistical inference procedures on survey data, including adjustments for non-response.

Lecture plan (3 lectures + 1 tutorial per week)

Bring a laptop to the Friday (tutorial) lecture.

Week 1 – The Survey Process

- The survey process, sources of survey error
- Populations and Frames
- Sampling in finite and infinite populations; Basics of SRSWOR inference; sample weights;
- **Tutorial:** Sampling basics; Population and Frames

Week 2 – Sampling

Use of `sample()` and other simulation methods in R for a variety of designs [`sampler` package (does power calculations, allocation, SRS and STSRS, but not cluster)] [`sampling` package – can do cluster samples]

- Simple random sampling, with and without replacement
- Stratified sampling
- Cluster sampling
- Sampling in complex designs
- Sample size calculation and allocation
- Survey weights – create estimates from weights without standard errors

- **Tutorial:** using R to draw samples and compute survey weights

Week 3 - Inference with the R survey package

- Inference in simple and complex designs, including regression
- Variance estimation
- Definition and interpretation of the design effect
- **Tutorial:** using R to specify designs and draw inferences

Week 4 – Non-response and complex designs

- Post-stratification for non-response adjustment
- Imputation (`simputation` package)
- Jackknife and bootstrap inference
- **Tutorial:** raking adjustments to weights; imputation; jackknife inference

Assignment 4 (10%) [due end of Week 12: covers material from Weeks 9-11]

- Survey process, populations and frames
- Describe a real survey, comment on design
- Propose a sample design, frame, comment on scope and coverage
- Draw and analyse a sample under various designs
- Compute sample sizes and allocate samples in stratified sampling
- Compute and interpret the design effect
- Analyse data and present results from a survey

Test 2 (30%) [assessment period] [1.5 hours, closed book]

- Survey process, populations and frames
- Describe sample designs
- Compute and interpret sample weights
- Use sample weights for point estimates
- Interpret output from survey inference in R
- Compare results, interpret design effect
- Describe methods for non-response adjustment (imputation and post-stratification); comment on bias reduction and assumptions being made
- Describe the jackknife method of variance estimation

1.2 Recommended reading

There are lots of books on sample surveys.

- Thomas Lumley’s book ‘Complex Surveys’ (Lumley (2010)) is a good introduction to sample survey analysis with R, and accompanies the **survey** package (Lumley (2004)).
- Groves et al. ‘Survey Methodology’ (groves.etal) is a good general summary of the important issues in survey design.
- Sarndal, Swensson and Wretman is a classic in the field of the theory of sampling (Särndal et al. (1992))
- Likewise the books by Cochran (1977) and Kish (1965)
- Little and Rubin have a great book on missing data (Little and Rubin (2002))
- More recent are the books by Lohr (1999) and Scheaffer et al. (2006)
- And in 1994 StatsNZ published ‘A Guide to Good Survey Design’ (Zealand (1995))

- Notes from the STAT392 course on Sample Surveys https://homepages.ecs.vuw.ac.nz/~rarnold/STAT392/SampleSurveysBook/_book/index.html
- A crash guide to using iNZight for sample surveys https://homepages.ecs.vuw.ac.nz/~rarnold/STAT392/iNZightManual/_book/index.html (iNZight is free software, using the **survey** package in the back end.)

And also

- CRAN site on R resources for Official Statistics <https://cran.r-project.org/web/views/OfficialStatistics.html>
- Online book on using R for survey analysis: https://bookdown.org/jimr1603/Intermediate_R_-_R_for_Survey_Analysis/
- This gallery of R graphs <https://r-graph-gallery.com/>

Some simple lectures on Sample Surveys

- https://r-project.ro/conference2017/presentations/Camelia_Goga_sondages.pdf
- Notes for epidemiologists: https://www.epirhandbook.com/en/new_pages/survey_analysis.html
- Notes on imputation: <https://www.r-bloggers.com/2023/01/imputation-in-r-top-3-ways-for-imputing-missing-data/>
- An online book about spatial sampling <https://dickbrus.github.io/SpatialSamplingwithR/>
- From Qualtrics <https://www.qualtrics.com/articles/strategy-research/sampling-methods/>

Data

- UC Irvine Machine Learning data repository: <http://archive.ics.uci.edu/>
- Some free microdata sets: <https://asdfree.com/>
- Kaggle data repository: <https://www.kaggle.com/>
 - Airline dataset from Kaggle: <https://www.kaggle.com/datasets/thedevastator/airlines-traffic-passenger-statistics> [distributed under the Creative Commons Licence 4.0: <https://creativecommons.org/licenses/by-nc-sa/4.0/>]
- Address file from LINZ: <https://data.linz.govt.nz/layer/123113-nz-addresses/>
- List of NZ schools <https://www.educationcounts.govt.nz/directories/list-of-nz-schools>
- Data on Population, GDP and Military Expenditure from ‘Our World in Data’ <https://ourworldindata.org/>

1.3 R Packages

- **survey** package (Lumley (2004)) with this handy summary of command uses <https://stats.oarc.ucla.edu/r/seminars/survey-data-analysis-with-r/>
- **srvyr** a dplyr wrapper for the **survey** package and **jtools** which also provides some extensions
- **gtsummary** creates outputs ready for publication (difficult to build/load)
- **sampling** package for drawing samples according to various designs
- **mice** package for imputation
- **simputation** package for imputation (difficult to build/load)
- **haven** package to import SAS, Stata and SPSS data files

Chapter 2

Sampling

In statistics the data we analyse is often a subset, i.e. a **sample** of members of a **population**.

For example we might have a sample of households whose members respond to a questionnaire, or a sample of plants from a field.

In order to select from a population we first enumerate (list) the population in some way and then make a random selection of our subset members from that list. The method by which we do so is called the **sample design**.

Our focus in statistics is often the **estimation** of population characteristics (known as **population parameters**) on the basis of sample characteristics (which are called **sample statistics**). We are always aware that the sample we drew from the population is not the only sample we could have drawn, and therefore our conclusions are subject to uncertainty, which we call **sampling error**.

This uncertainty associated with sampling is the reason we have confidence intervals on our estimates, and why we can neither **prove** nor **disprove** hypotheses.

(One exception is when we do a **census** where we select and get fully accurate data from every member of the population. Then there's no uncertainty left at all, and our sampling errors are zero.)

When generating **simulated** datasets we also generate samples. Simulation is often done as a means of testing out statistical ideas and methods, or (in the case of the bootstrap method) when conventional ways of estimating standard errors (and therefore confidence intervals) don't work.

Before we get into detail of estimation and inference, let's spend some time thinking about how sampling works.

2.1 Random numbers

In both settings (sampling and simulation) we need to start with a method of generating random numbers. We'll start by assuming we have some means of generating a random sequence of numbers uniformly distributed in the interval $[0,1]$:

$$U_i \sim \text{Uniform}(0,1) \quad \text{for } i = 1, 2, 3, \dots$$

Here 'uniformly distributed' means that all values in the range $[0,1]$ are equally likely to occur.




Figure 2.1: Statistical Inference

In R this is done using the `runif()` function. For example to generate 5 random draws from `Uniform(0,1)`:

```
runif(5)
```

```
## [1] 0.2875775 0.7883051 0.4089769 0.8830174 0.9404673
```

Another draw will a different set of random numbers:

```
runif(5)
```

```
## [1] 0.0455565 0.5281055 0.8924190 0.5514350 0.4566147
```

If we draw a histogram of 100 draws it looks a bit lumpy:

```
hist(runif(100), xlab="x", ylab="Frequency")
```

But a histogram of 100,000 draws is very uniform:

```
hist(runif(100000), xlab="x", ylab="Frequency")
```

Although the numbers appear to be random they are in fact generated using formulae, so if one knows the formula the sequence of random numbers generated by `runif()` is completely predictable. However the formula is complex enough that there are no easily discernible patterns in sequences of many millions of draws. These numbers are **psuedo-random**, but they are random enough for our purposes.

Underlying the random number generator is a **seed** value, which initiates the sequence. If we set the seed to a particular value, the sequence of random numbers that will be generated after doing so will always be the same.

```
set.seed(111)
```

```
runif(5)
```




Figure 2.2: 100 draws from Uniform(0,1)




Figure 2.3: 100000 draws from Uniform(0,1)

```
## [1] 0.5929813 0.7264811 0.3704220 0.5149238 0.3776632
```

```
runif(5)
```

```
## [1] 0.41833733 0.01065785 0.53229524 0.43216062 0.09368152
```

```
runif(5)
```

```
## [1] 0.55577991 0.59022849 0.06714114 0.04754785 0.15620252
```

```
set.seed(8)
```

```
runif(5)
```

```
## [1] 0.4662952 0.2078233 0.7996580 0.6518713 0.3215092
```

```
set.seed(111)
```

```
runif(5)
```

```
## [1] 0.5929813 0.7264811 0.3704220 0.5149238 0.3776632
```

This property can be very useful when testing and comparing algorithms: we can generate an identical sequence of random numbers and check (for example) that two methods of estimation give the same results on two identical randomly generated datasets.

2.2 Sampling from categorical random variables

Draws from a uniform distribution are all we need to simulate/sample from any distribution.

Heads or tails: if the probability of flipping heads is 0.5, we can generate n random numbers, and then label any values less than 0.5 as Heads and the remainder as tails.

```
n <- 10
```

```
u <- runif(n)
```

```
u
```

```
## [1] 0.2875775 0.7883051 0.4089769 0.8830174 0.9404673 0.0455565 0.5281055
```

```
## [8] 0.8924190 0.5514350 0.4566147
```

```
y <- ifelse(u<=0.5,"Heads","Tails")
```

```
y
```

```
## [1] "Heads" "Tails" "Heads" "Tails" "Tails" "Heads" "Tails" "Tails" "Tails"
```

```
## [10] "Heads"
```

```
table(y)
```

```
## y
```

```
## Heads Tails
```

```
##      4      6
```

```
table(ifelse(runif(10000)<=0.5,"Heads","Tails"))
```

```
##
```

```
## Heads Tails
```

```
## 5059 4941
```

We can extend this idea to any set of categories.

Red, Green or Blue: Say we have a six sided die: three sides are red, two are green and one is blue. The probabilities of rolling the three colours are therefore $\frac{1}{2}$, $\frac{1}{3}$, $\frac{1}{6}$ respectively. These probabilities add to one.

The random variable Y is the colour that the die shows when we roll it.

```
pp <- c(Red=1/2, Green=1/3, Blue=1/6)
sum(pp)
```

```
## [1] 1
```

In general we might have m categories, where the probability of that Y category k is p_k : $\Pr(Y = k) = p_k$. and

$$\sum_{k=1}^m p_k = 1$$

We then compute the **cumulative probabilities** c_k , the partial sums of these probabilities:

$$c_k = \sum_{j=1}^k p_j$$

The final cumulative probability c_m is always 1 by definition.

Implicitly we also have a value $c_0 = 0$, which is the the empty sum and always equal to zero.

These cumulative probabilities c_k are 'the probability that random variable Y has a category label less than or equal to k ':

$$\Pr(Y \leq k) = c_k = \sum_{j=1}^k p_j$$

Red, Green or Blue: the cumulative probabilities are

$$\begin{aligned} c_0 &= 0.0000 \\ c_1 &= p_1 = 0.5000 \\ c_2 &= p_1 + p_2 = 0.5000 + 0.3333 = 0.8333 \\ c_3 &= p_1 + p_2 + p_3 = 0.5000 + 0.3333 + 0.1667 = 1.0000 \end{aligned}$$

```
m <- length(pp) # Number of categories
cp <- cumsum(pp)
cp
```

```
##      Red      Green      Blue
## 0.5000000 0.8333333 1.0000000
```

We then take a set of uniform random numbers $u_1, \dots, u_n$. We then want to find which of the m labels corresponds to each random number u_i .

The rule is that if u_i is less than cumulative probability c_k but greater than c_{k-1} , then we assign the label k .

There are various ways to do this, here's a version using `case_when()` from the `dplyr` package:

```
set.seed(112)
u <- runif(8)
u
```

```
## [1] 0.37667253 0.92201944 0.99187774 0.97723602 0.23629728 0.10706678 0.03915213
## [8] 0.72802690
```

```
y <- case_when(u <= cp[1] ~ "Red",
               u <= cp[2] ~ "Green",
               .default="Blue") # here .default means "otherwise"
data.frame(round(u,5),y)
```

```
##   round.u..5.    y
## 1    0.37667   Red
## 2    0.92202  Blue
## 3    0.99188  Blue
## 4    0.97724  Blue
## 5    0.23630   Red
## 6    0.10707   Red
## 7    0.03915   Red
## 8    0.72803  Green
```

```
table(y)
```

```
## y
## Blue Green  Red
##    3     1     4
```

```
table(y)/length(y)
```

```
## y
## Blue Green  Red
## 0.375 0.125 0.500
```

Note that the categories in the table haven't come out in the Red-Green-Blue order that we want - instead they're in the default alphabetical order.

We can fix this by converting `y` to a `factor` object:

```
y <- factor(y,levels=names(pp))
table(y)
```

```
## y
## Red Green  Blue
##    4     1     3
```

Here's what is really going on: we if we plot the cumulative probabilities as a staircase – a **cumulative distribution function (CDF)** then we draw u_i from a Uniform(0,1) distribution, and then find where a horizontal line drawn at the level of u_i intersects the CDF: and that is the value of the categorical variable y_i .

The larger the probability p_k the bigger the jump in the staircase, and the larger the probability that category k will be selected.

```
set.seed(112)
n <- 10000
u <- runif(n)
y <- case_when(u <= cp[1] ~ "Red",
               u <= cp[2] ~ "Green",
               .default="Blue")
y <- factor(y, levels=names(pp))
table(y)
```

```
## y
```




Figure 2.4: Drawing a random category




Figure 2.5: Drawing a random category

```
##   Red Green  Blue
##  4976  3336  1688
```

The `sample()` function implements this conveniently: select from the integers 1 to `m`

```
n <- 10000
y <- sample(m, size=n, prob=pp, replace=TRUE)
table(y)
```

```
## y
##   1    2    3
## 5014 3350 1636
```

(need `replace=TRUE` here, see below).

... or state the vector `1:m` explicitly

```
y <- sample(1:m, size=n, prob=pp, replace=TRUE)
table(y)
```

```
## y
##   1    2    3
## 5010 3371 1619
```

... or directly from the vector of labels ...

```
y <- sample(c("Red", "Green", "Blue"), size=n, prob=pp, replace=TRUE)
y <- factor(y, levels=names(pp))
table(y)
```

```
## y
##   Red Green  Blue
##  4980  3325  1695
```

... get the labels from `pp` ...

```
y <- sample(names(pp), size=n, prob=pp, replace=TRUE)
y <- factor(y, levels=names(pp))
table(y)
```

```
## y
##   Red Green  Blue
##  4994  3389  1617
```

2.3 Sampling from discrete numerical random variables

The ideas above carry straight over to discrete numerical random variables.

The Binomial distribution $Y \sim \text{Bin}(m, p)$ is the number of successes in m trials where the probability of success on each trial is the same p , and all the trials are independent.

Y takes values 0 to m , and those probabilities can be calculated by

$$P(Y = k) = \binom{m}{k} p^k (1-p)^{m-k} = \frac{m!}{k!(m-k)!} p^k (1-p)^{m-k}$$

R provides these probabilities in the function `dbinom()`

e.g. $m = 10$ rolls of a six-sided die, and Y is the number of times we get a 1, so $p = 1/6$

```

m <- 10
p <- 1/6
pp <- dbinom(0:m, m, p)
pp

## [1] 1.615056e-01 3.230112e-01 2.907100e-01 1.550454e-01 5.426588e-02
## [6] 1.302381e-02 2.170635e-03 2.480726e-04 1.860544e-05 8.269086e-07
## [11] 1.653817e-08

barplot(pp, names=0:m,
        xlab="y", ylab="Pr(Y=y)",
        main=bquote("Bin("*. (m)*", ".*. (sprintf(" %.4f", p))*")"))

```




Figure 2.6: Binomial probabilities

The CDF is given by the function `pbinom()`

```

cp <- pbinom(0:m, m, p)

plot(0:m, cp, type="s",
     xlab="y", ylab="Pr(Y<=y)",
     main=bquote("Bin("*. (m)*", ".*. (sprintf(" %.4f", p))*")"))

```

Drawing a single value of Y proceeds the same way as for a categorical variable:

So we can draw multiple values of Y from $\text{Bin}(10, \frac{1}{6})$ using the `sample()` function:




Figure 2.7: Binomial CDF




Figure 2.8: Drawing from a Binomial(10,1/6) distribution

```
n <- 10000
pp <- dbinom(0:m, m, p)
y <- sample(0:m, n, prob=pp, replace=TRUE)
barplot(table(y))
```




Figure 2.9: 10000 draws from a Binomial(10,1/6)

Of course R provides a specific function for doing the same thing: `rbinom()`

```
n <- 10000
y <- rbinom(n, m, p)
barplot(table(y))
```

But that function is just doing what we've described above using `sample()`.

A particular common example is where we have m objects and want to select one of them at random, with each object having equal probability of being selected: i.e. each has probability $1/m$.

2.4 Sampling from continuous random variables

Continuous random variables are simulated using the same principle.

The cumulative distribution function (CDF) of a continuous random variable Y is defined to




Figure 2.10: 10000 draws from a Binomial(10,1/6)

be

$$F(y) = \Pr(Y \leq y) = \int_{-\infty}^y f(y') \, dy'$$




Figure 2.11: Normal distribution density and CDF

We can use $F(y)$ to draw a random value from the distribution of Y by first drawing a random value U from a $\text{Uniform}(0,1)$ distribution.

We then solve the equation

$$U = F(y)$$

for y : i.e. find the value of y corresponding to a CDF value of U .

For many standard distributions (including the Normal and Gamma distributions) there is no analytic formula which gives the solution for y given u : but there are numerical approximations which can be used in those settings.

One simple case which does have an analytic solution is the Exponential distribution $Y \sim \text{Exp}(\lambda)$. This is a one parameter distribution where Y is non-negative ($Y \geq 0$), and can be used to model cases such the length of time it takes for an event to occur (e.g. the length of time for a radioactive nucleus to decay). The single parameter λ is the **rate** (e.g. number of radioactive decay events per unit time).

The exponential distribution has probability density function

$$f(y) = \lambda e^{-\lambda y} \quad \text{where } y \geq 0$$




Figure 2.12: Exponential probability distribution

and has mean $E[Y] = 1/\lambda$.

The CDF is the integral of the probability density function

$$\Pr(Y \leq y) = F(y) = \int_0^y f(y) dy$$

which for the Exponential distribution is

$$F(y) = \int_0^y \lambda e^{-\lambda t} dt = 1 - e^{-\lambda y}$$




Figure 2.13: Exponential cumulative probability distribution

Solving

$$u = F(y) = 1 - e^{-\lambda y}$$

for y leads to

$$y = -\frac{1}{\lambda} \log(1 - u)$$

Using this function we can easily draw directly from an Exponential distribution.

```
n <- 10000
u <- runif(n)
lambda <- 0.5
y <- -(1/lambda)*log(1-u)
mean(y)
```

```
## [1] 1.990456
```

```
hist(y, main="Histogram of draws from Exp(0.5)")
```




Figure 2.14: 10000 draws from an Exponential distribution

2.5 Sampling from a finite collection

The simulations discussed so far all draw from **infinite** populations: we can set the sample size n to any value. However it often occurs that we have only a finite population - e.g. when drawing a random sample of students from a class, or households from a suburb for interview. Then the population size N is finite.

In the general case we have a set of N **units** (i.e. students or households) in a population of finite size, and we want to draw a sample of n of these units.

We usually start by listing or **enumerating** the units. This list is called a **frame** or a **sample frame**.

Consider the case where we have a small street containing $N = 10$ houses, numbered 1-10. We want to select 5 of these houses with equal probability. There are a number of ways we could do this.

2.5.1 Sampling without replacement

If the houses have equal probability then we could proceed as follows:

- We draw a random number from 1-10, e.g. house number 2, and mark it as selected
- We draw another random number from 1-10, e.g. house number 5, and mark it as selected
- We repeat this process, but if we choose a house that is already selected, we discard that selection;
- The process continues until we have a set of 5 distinct houses.

At the start we take our frame (i.e. list) of houses, and initially mark them all as unselected. Then at each step where we successfully select a previously unselected house we mark that house as selected.

```
bigN <- 10 # population size
frame <- data.frame(house=1:bigN, selected=FALSE)
n <- 5 # sample size
nselected <- 0
while(nselected < n) {
  i <- sample(bigN, 1) # sample a single number at random from 1:bigN with equal prob
  if(!frame$selected[i]) {
    frame$selected[i] <- TRUE
    nselected <- nselected + 1
  }
}
frame$house[frame$selected]
```

```
## [1] 1 6 7 8 10
```

It's not guaranteed how many steps this process is going to take, since the number of times we select a house that has already been selected is random. If the sample size n is small and the population size N is large then it will take n steps, or possibly just a few more. If n is close to N and N is large, it could take a very very long time.

A simpler way to solve this is to assign each house a distinct random number between 0 and 1, and then sort the frame in order of the random numbers.

```
bigN <- 10 # population size
frame <- data.frame(house=1:bigN, selected=FALSE)
n <- 5 # sample size
frame$u <- runif(bigN)
frame <- frame[order(frame$u),]
frame$selected[1:n] <- TRUE
frame
```

```
##   house selected      u
## 9      9      TRUE 0.0812427
## 4      4      TRUE 0.3739315
## 2      2      TRUE 0.4137952
## 6      6      TRUE 0.4341604
## 10     10      TRUE 0.4909159
## 3      3     FALSE 0.5850114
## 8      8     FALSE 0.8708720
## 5      5     FALSE 0.8837804
## 1      1     FALSE 0.9278206
## 7      7     FALSE 0.9475266
```

```
frame <- frame[order(frame$house),]
frame$house[frame$selected]
```

```
## [1] 2 4 6 9 10
```

Of course, the simplest way of all is just to use the `sample()` function

```
bigN <- 10 # population size
n <- 5 # sample size
sample(bigN, n)
```

```
## [1] 10 4 6 3 2
```

It's very often that we do have a frame, a list of objects with names, and we want to select from this at random.

```
frame <- data.frame(Name=c("Mark","Chloe","Sam","Caitlin","Harry","Nina"),
                    Age=c(20,19,18,19,19,21),
                    Major=c("STAT","STAT","DATA","COMP","MATH","STAT"))
frame
```

```
##      Name Age Major
## 1   Mark  20  STAT
## 2  Chloe  19  STAT
## 3    Sam  18  DATA
## 4 Caitlin  19  COMP
## 5   Harry  19  MATH
## 6    Nina  21  STAT
```

We can select, say, 3 names

```
n <- 3
samplenames <- sample(frame$Name, n)
samplenames
```

```
## [1] "Nina" "Sam"  "Mark"
```

and then subset the frame

```
samplemembers <- frame[frame$Name%in%samplenames,]
samplemembers
```

```
##      Name Age Major
## 1 Mark  20  STAT
## 3 Sam  18  DATA
## 6 Nina  21  STAT
```

If you like dplyr

```
samplemembers <- frame |>
  filter(Name%in%samplenames)
samplemembers
```

```
##      Name Age Major
## 1 Mark  20  STAT
## 2 Sam  18  DATA
## 3 Nina  21  STAT
```

Note we can find out where in the frame particular individuals are located

```
match(c('Mark','Harry'),frame$Name)
```

```
## [1] 1 5
```

If we're sampling without replacement, note (of course) that the sample size can be no bigger than the population size $n \leq N$.

2.5.2 Sampling with replacement

When selecting a sample of 5 houses it makes sense not to allow the possibility of selecting the same house twice. So when we've selected a house we remove it from the list of houses available to select. When using `sample()` this means that we set `replace=FALSE`, meaning after we've selected a member of the population we don't put it back again.

But there are occasions where we select **with replacement** is the sensible thing to do. This means there is a chance where a unit can be selected more than once, and this occurs in situations where we want to assign different probabilities of selection to different units of the population.

Starting with the the case of equal probabilities, we repeatedly draw units from the population under identical circumstances: on every draw each unit is available for selection, and we count how many times each unit is selected.

```
set.seed(111)
bigN <- 10
n <- 5
replicate(n, sample(bigN, 1))
```

```
## [1] 4 3 9 5 3
```

In this sample we've drawn Unit 3 twice.

If n is small and $bigN$ is large the probability of repeats is small. But if a repeat happens we of course do not make our selected household repeat the questionnaire, instead we count them twice (or as many times as they are selected) in our analysis.

Imagine our sample of 10 houses aren't actually just houses, they're farms, and some of them are very large and some are very small. We want to give the large farms a high chance of selection.

```
set.seed(122)
bigN <- 10
n <- 5
frame <- data.frame(farm=1:10, ncows=c(1000,5000,2000,500,500,
                                       5000,500,5000,30000,100))
frame
```

```
##      farm ncows
## 1      1  1000
## 2      2  5000
## 3      3  2000
## 4      4   500
## 5      5   500
## 6      6  5000
## 7      7   500
## 8      8  5000
## 9      9 30000
## 10     10   100
```

```
sample(frame$farm, n, prob=frame$ncows, replace=TRUE)
```

```
## [1] 2 2 9 9 9
```

Unsurprisingly, the very large farm number 9 has been selected three times here.

It **is** possible to have unequal probability sampling **without** replacement, estimation of standard errors is complicated in such designs, so we won't be using such designs.

For now, just note that we switch to with replacement designs when we're sampling from sets of units where we want to have units with different probabilities of selection.

Chapter 3

Sample Surveys

A sample survey is an exercise in which **data** are collected from a **sample** (a subset) or a **population**. The data collected are used to create **estimates** of the characteristics or **parameters** of the population.

Some examples:

- A sample of voters are telephoned to ask their voting intentions, with the aim of predicting the outcome of the election;
- A sample of sites in a river catchment is tested for the presence of algae to determine the spread of the algae throughout the entire catchment;
- A sample of households is selected, and the residents asked about their employment status, in order to estimate the national unemployment rate.

A **census** is a special case of a sample survey in which every member of the population is surveyed.

3.1 The Survey Process

The different steps in the survey process are shown in Figure 3.1.

- Every survey has a set of **objectives** – the particular populations which are of interest;
- These parameters are properties of the **target population**: the population about which estimates are to be made;
- Not all members of the target population are able to be identified: the ones that can be surveyed make up the **survey population**;
- The **sample frame** is a listing of all members of the survey population;
- The **sample design** is the method of selecting the sample from the frame. The **sample size** is usually decided as a compromise between the required **accuracy of estimates** and the survey **costs** and other constraints (e.g. the time available for the survey);
- **survey instrument** is the means of data collection. It is usually a paper or electronic questionnaire, completed by the respondent or the interviewer/observer. The instrument aims to measure the properties of the sample members which mean that the desired population characteristic can be estimated. The instrument must be **valid** (it must actually measure what it intends to measure) and **reliable** (repeated measurements of the same sample member under identical circumstances should always yield similar results);
- The sample members are **contacted** and **recruited** into the sample. Not all selected members will be able to be contacted (even after strenuous efforts), and some will not




Figure 3.1: The Survey Process

- respond even if contacted;
- The data are collected from the sample members in some **mode** (e.g. face-to-face, telephone, web, observational, ...);
- The data collected are **captured** (stored on a computer); **coded** (converted into standard classification systems); **edited** (checks for data consistency); and **stored** in a final dataset;
- The data may be **adjusted**, e.g. imputation or weight adjustments for nonresponse may be done;
- Estimates** of the parameters of interest are constructed, and other **analysis** of the data is carried out (e.g. regression estimation, comparison with other data etc.);
- The results are summarised in a **report**;
- The original data may be **archived** in some appropriate form, or **destroyed**;
- A post-survey **evaluation** may be made to determine how well the survey met its goals.

For example: The Household Labour Force Survey (HLFS).

- One of the main **objectives** of the HLFS is to estimate the unemployment rate every quarter;
- The **target population** is the working age population of New Zealand;
- The **survey population** is the civilian, non-institutionalised usually resident population of adults aged 15+ living in permanent, private dwellings on the main islands of New Zealand (North Island, South Island and Waiheke Island);
- The **sample frame** is a list of dwellings created at the most recent census and regularly updated to reflect changes;
- The **sample design** is a stratified cluster design. Within each regional government area a sample of small areas (PSUs) is selected. A sample of households is selected from within each PSU. Every adult from the selected households is surveyed. 15000 households and 30000 adults are surveyed every three months, in order to create unemployment estimates accurate to within 0.5%.
- The **survey instrument** is an electronic questionnaire.
- The interviewers contact the households in person or by telephone, making up to 10 call backs to ensure contact is made. **Proxy responses** are permitted (i.e. one household member can respond on behalf of another). The interview takes place by computer assisted personal or telephone interviewing (CAPI or CATI).
- The results are **post-stratified** to match the current population estimates in each local government area;
- Estimates** of the unemployment rate are published within about 6 weeks of the end of the quarter.

3.2 Populations and Frames

The target population defines the **scope** of the survey: the set of units which make up the population in which we are interested.

The sample frame is a method of gaining access to the members of the target population. It is an actual or conceptual list. However it may not be perfect (Figure 3.2).

- Not all members of the target population will appear on the sample frame: the missing units are **out of coverage**;
- Not all units listed on the sample frame are members of the target population: the extra units are **out of scope**.
- The survey population are those units which are both in the target population (in scope) and on the frame (in coverage).

Our aim is to have a sample frame which covers most if not all of the target population, and which has only very few out of scope units.

The reason for this is very important: we can make estimates about units in a population **only** if they are included in the sample frame. This inference is possible **whether or not we select them**. If there is a subpopulation of units that is not on our frame (i.e. they are out of coverage) then we can never select them, we can never learn anything about them, and therefore we cannot make inferences about them.




Figure 3.2: Scope and Coverage

3.2.1 What is a sampling frame?

The sampling frame is a key element in the survey process which provides the basis from which a *random* (probability based) sample can be selected. The purpose of having a sampling frame is to enable each member of the target population a nonzero probability of being selected to participate in the survey. The *quality* of the data, from a survey which has as its basis a comprehensive frame, can be *measured and controlled*.

The term sampling frame is sometimes misunderstood, often being thought of as a list containing each member of the population. In fact that is more a description of an ‘ideal’ sampling frame. The best way to show what a sampling frame can be in practice is to examine some examples:

- **List Frames.** Good up-to-date lists make ideal frames. However lists of units (electoral rolls, telephone listings etc.) are generally formed for purposes other than for use as sampling frames. For this reason they are often not close to ‘ideal’ sampling frames. Some discussion of list frame problems comes later.
- **Multi-stage Sampling Frames.** These frames consist of a set of lists which are arranged in a hierarchy of stages. For example a typical household survey frame will be organised like so: The first list is a set of non-overlapping geographic areas from which a sample of areas is taken. For example we may list all the local regional councils and take a sample of them. Within these areas we might then list all the households

and take samples from the household lists. From each household we would then list the residents and interview selections of them. Note that we do not have to list the entire population. The only comprehensive list needed is the list of the areas which will be quite small. By giving each area a chance of being selected we ensure each person living in each area has a chance of selection.

- **Continuous Sampling Frames.** Consider the population of people entering a country from international flights. Suppose we choose a random number between one and ten ($= k$) and then interview the k th, $(k + 10)$ th, $(k + 20)$ th ... persons who file through customs. We will have a sample where every person has had a 1 in 10 chance of being selected. But what is the frame? There is no list. We call such a situation a ‘conceptual list’, where the people filing into the country effectively formed a long queue of people which we have treated like a listing of the population, although a listing of names was never constructed. \end{description}

One can find many different definitions of sampling frames. The following is a simplified definition that should serve our purposes:

A *sampling frame* is part of the system of rules and procedures used when determining who is to be surveyed and which ensure that each member of the population get some measurable (non-zero) chance of participating in the survey.

The examples above show that we need to think beyond the concept of a list of units, whose construction is a separate exercise from the other survey stages.

3.2.2 What is a *good* sampling frame?

We will often have several frames available and will have to decide between different options. To help us with this decision we need to be a little more specific about what makes a good sampling frame. The first thing we should consider is the effective coverage of the sampling frame. This can be determined by testing it for the following properties:

- **Each unit should be counted at least once.** If there are missing units (out of coverage) then the survey results may be biased if the excluded units tend to be different to the units which are included.
- **Each unit should be counted only once.** The problem is that we usually can not tell how many times a unit is duplicated. This means we can not tell how many chances a unit has of being sampled. This can cause similar bias if the units which are duplicated are different to the other units. The results will be biased towards the duplicated sub-group of the population.
- **Each unit should be distinguishable from the other units on the frame.** If we select a unit from the frame (say Jane Smith) but then can not tell who this refers to, when there are many Jane Smiths in the population, then the frame has not done its job.
- **Should provide up-to-date information about the units.** For example if it claims to provide Jane Smith’s address it should be the current address.
- **No units should be there if they are not meant to be.** Extra units (out of scope) are not a serious problem unless we fall into the elementary error of replacing them whenever we come across them in the sample. This confuses selection probabilities. The correct procedure is to drop any ineligible unit found in the sample and not replace them.
- **Should allow direct access to each unit.** Ideally the frame will ultimately supply us with a list of the units we are interested in analyzing. Sometimes a frame will only

give samples of groups of units e.g. a list of addresses, each of which may contain many people. Further work is needed before a sample of people is realized.

3.3 Approaches to sampling

We're talking here about random sampling, or **probability sampling**. For the purposes of statistical inference such sampling is the only defensible method - it's the only one that leads us to properly constructed confidence intervals and hypothesis tests. Other approaches do exist though, and even though you'll see them used they all have significant defects.

- **Probability sampling** is the gold standard: every member of the population has a known, non-zero chance of selection;
- **Purposive sampling**: the surveyor selects members of the population that they consider to be representative - dangerously biased towards the views and prejudices of the surveyor;
- **Haphazard sampling**: the surveyor uses any members of the population that they come across. This is what happens on the street corner when you're asked questions by a surveyor;
- **Quota sampling**: resembles stratified sample (see later): the surveyor fills quotas of respondents to balance numbers of respondents in categories (e.g. by age/gender/ethnicity), but does not make an accounting for non-response (see later);
- **Snowball sampling**: used in rare populations that are highly connected (e.g. refugees from a particular country): existing sample members recruit new sample members using their contacts;
- **Self-selected sampling**: this is the **worst** of all: the respondents themselves decide whether to be selected: they see an advertisement on a webpage or phone in with their views. Only profiles motivated population members - the silent majority are never measured.

3.4 Example datasets

We'll use and re-use a number of standard datasets.

3.4.1 API

This is a dataset on Student performance in California schools. The dataset is distributed with the `survey` library in R, and there are a number of unit record datasets which are drawn from a population dataset `apipop` under various sample designs. The population units are the schools, and a number of characteristics are available. Full documentation can be found in the `survey` package. The variables we use mostly are:

- `stype` - School type (Elementary/Middle/High School)
- `snum` - School number
- `dnum` - District number
- `api99` - Academic Performance Indicator for the school in the year 1999
- `api00` - Academic Performance Indicator for the school in the year 2000
- `sch.wide` - Did the school meet its school-wide growth target?
- `awards` - Is the school eligible for an awards program
- `meals` - Percentage of students eligible for subsidised meals
- `ell` - Percentage of students that are English Language Learners
- `grad.sch` - Percentage of students with parents with a postgraduate education
- `enroll` - Number of students enrolled
- `api.stu` - Number of students tested in the API

3.4.2 From the dplyr package

- `starwars` - a list of characteristics of 87 characters from the Star Wars movie series;

Chapter 4

Statistical Inference

4.1 Inference in infinite populations

Let's briefly review a few familiar results from earlier statistics courses.

Most of what we've done in earlier courses has assumed that we have a **simple random sample** from an **infinite** population.

We'll use data from the `starwars` dataset (in the `dplyr` package). There are data on 87 characters, including their heights in centimetres.

Six heights are missing, but let's just guess them so we have no missing data (this process of inventing data has a respectable name: **'imputation'**, and we'll come back to it later).

```
library(dplyr)
data(starwars)
sum(is.na(starwars$height))

## [1] 6

starwars <- starwars |>
  mutate(height=case_when(name=="Arvel Crynyd" ~ 172, # same as Luke
    name=="Finn" ~ 172, # same as Luke
    name=="Rey" ~ 150, # same as Leia
    name=="Poe Dameron" ~ 172, # same as Luke
    name=="BB8" ~ 96, # same as R2D2
    name=="Captain Phasma" ~ 167, # same as C3PO
    .default=height))
```

The mean height $\bar{Y}$ of characters in the Star Wars universe is a **population parameter**. We can **estimate** it from a sample of characters.

Note: we're using capital letters for **population parameters** ($\bar{Y}$ is the population mean), and lower case letters for **sample statistics** ($\bar{y}$ is the sample mean).

It's also common to see Greek letters for **population parameters** (e.g. μ for the population mean), upper case latin letters for random variables, and lower case latin letters for sample statistics.

For the purposes of this exercise let's imagine there are only $N = 87$ individuals in the entire Star Wars universe. Let's draw a random sample of $n = 20$ from this population of $N = 87$:

name	height	mass	hair_color	homeworld
Tarfful	234	136	brown	Kashyyyk
Jek Tono Porkins	180	110	brown	Bestine IV
Bossk	190	113	none	Trandosha
Lando Calrissian	177	79	black	Socorro
IG-88	200	140	none	NA
Kit Fisto	196	87	none	Glee Anselm
Biggs Darklighter	183	84	black	Tatooine
Obi-Wan Kenobi	182	77	auburn, white	Stewjon
Quarsh Panaka	183	NA	black	Naboo
Mace Windu	188	84	none	Haruun Kal
Ratts Tyerel	79	15	none	Aleen Minor
Beru Whitesun Lars	165	75	brown	Tatooine
Chewbacca	228	112	brown	Kashyyyk
Jabba Desilijic Tiure	175	1358	NA	Nal Hutta
Rugor Nass	206	NA	none	Naboo
Sly Moore	178	48	none	Umbara
Dormé	165	NA	brown	Naboo
Leia Organa	150	49	brown	Alderaan
Finn	172	NA	black	NA
Darth Maul	175	80	none	Dathomir

```
set.seed(222)
bigN <- nrow(starwars)
n <- 20
samplemembers <- starwars[sample(bigN, n, replace=FALSE),]
samplemembers |>
  dplyr::select(name,height,mass,hair_color,homeworld) |>
  kbl() |>
  kable_styling()
```

Compute the sample mean $\bar{y}$ and standard deviation s_y :

```
ybar <- mean(samplemembers$height)
ybar
```

```
## [1] 180.3
```

```
sd_y <- sd(samplemembers$height)
sd_y
```

```
## [1] 31.13147
```

We know how to make a point estimate $\widehat{\bar{Y}}$ for the population mean $\bar{Y}$ (remember that a hat $\widehat{\cdot}$ over a symbol indicates an estimate of that parameter): the population mean is estimated by the sample mean $\bar{y}$:

$$\widehat{\bar{Y}} = \bar{y} = 180.3$$




Figure 4.1: Histogram of heights of a sample of Star Wars characters (vertical line at the sample mean)

and a 95% confidence interval is given by

$$\widehat{Y} \pm \text{MOE}[\widehat{Y}] = \widehat{Y} \pm T_{n-1}^* \times \text{SE}[\widehat{Y}] \quad (4.1)$$

$$= \bar{y} \pm T_{n-1}^* \times \frac{s_y}{\sqrt{n}} \quad (4.2)$$

$$= 180.3 \pm 2.09 \times \frac{31.1}{\sqrt{20}} \quad (4.3)$$

$$= 180.3 \pm 2.09 \times 6.96 \quad (4.4)$$

$$= 180.3 \pm 14.6 \quad (4.5)$$

$$= (165.7, 194.9) \quad (4.6)$$

Here:

- $\text{SE}[\widehat{Y}]$ is the **standard error** or **sampling error** of the estimator $\widehat{Y}$ of the population mean $\bar{Y}$;
- T_{n-1}^* is the quantile of the t -distribution with $n - 1$ degrees of freedom so that the interval $[-T_{n-1}^*, +T_{n-1}^*]$ encloses 95% of the probability of the t -distribution;
- $\text{MOE}[\widehat{Y}]$ is the **margin of error**, the half-width of the confidence interval, and is the standard error scaled by T_{n-1}^* .

This is what is known as a **model based** estimate, the statistical model being

$$y_i = \mu + \varepsilon_i \quad \text{where } \varepsilon_i \stackrel{\text{iid}}{\sim} N(0, \sigma^2) \quad (4.7)$$

There are just two parameters here: the population mean μ , and the population standard deviation σ .

We can fit this simplest of all linear regression models using the `lm()` function and get the same results:

```
lm1 <- lm(height ~ 1, data=samplemembers)
summary(lm1)

##
## Call:
## lm(formula = height ~ 1, data = samplemembers)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -101.30   -6.05    0.70   11.20   53.70
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  180.300      6.961    25.9 2.77e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 31.13 on 19 degrees of freedom
##
## confint(lm1)
##
##              2.5 % 97.5 %
## (Intercept) 165.73 194.87
```

Here (**Intercept**) is the estimate for μ , the population mean.

The ‘residual standard error’ in the R output is actually what we’d rather call $\hat{\sigma}$, our estimate of the standard deviation σ .

The confidence interval is only valid if the model in equation (4.7) holds: i.e. if the data are normally distributed about the mean. The histogram in Figure 4.1 suggests that this is not the case however, since there is a clear outlier at low height.

The Central Limit Theorem (CLT) tells us that if n is large enough then our estimate for μ is useable, and its confidence interval realistic, even if the model in equation (4.7) doesn’t actually hold.

The CLT enables much of statistical practice to yield meaningful results even without our knowing the true data generation model.

In this course most of the examples will have sample sizes large enough that the CLT will mean our results are valid. In these cases when constructing a confidence interval it will be sufficient to replace T_{n-1}^* with the equivalent $Z^* = 1.96$ from the Normal distribution.

4.2 Inference in finite populations

Recall that the standard error of an estimate is the variability of estimates between different samples: in fact it is defined as the standard deviation of all possible estimates from all possible estimates drawn using the same sample design.

In an infinite population there are an infinite number of possible samples.

However in a finite population that isn’t the case. In a population of size N if we draw a simple random sample without replacement of size n there are

$$\binom{N}{n} = {}^N C_n \quad (4.8)$$

$$= \frac{N!}{n!(N-n)!} \quad (4.9)$$

$$= \frac{N \times (N-1) \times (N-2) \times \dots \times (N-n+1)}{n \times (n-1) \times (n-2) \times \dots \times 1} \quad (4.10)$$

possible samples: i.e. the number of ways it is possible to draw n objects from a collection of N without regard to the order they are selected.

In our example of selecting $n = 20$ from $N = 87$ there are an astronomical number of possible samples:

$$\binom{87}{20} = \frac{87!}{20!67!} = 2.38 \times 10^{19}$$

There is of course only one way to select a census though:

$$\binom{87}{87} = \frac{87!}{87!0!} = 1$$

(using the at first face slightly annoying definition $0!=1$). Since there’s only one possible sample when we’re drawing a census the variation among sample estimates from different samples is zero: i.e. the standard error is zero.

It’s also true that if n is close to N then sample estimates hardly differ from one another.

We clearly need a modification of our standard error formula

$$\mathbf{SE}[\widehat{Y}] = \frac{s_y}{\sqrt{n}} \quad (\text{Infinite population})$$

to allow for sample sizes n being close to the population size.

That modification is the **finite population correction (fpc)**:

$$\mathbf{SE}[\widehat{Y}] = \sqrt{1 - \frac{n}{N}} \times \frac{s_y}{\sqrt{n}} \quad (\text{Finite population})$$

The factor $\sqrt{1 - n/N}$ squashes the standard error to zero in the case that $n = N$, but for very small n has a value close to 1, and makes no difference at all.

In our example we had a standard error of

$$\mathbf{SE}[\widehat{Y}] = \frac{s_y}{\sqrt{n}} \quad (4.11)$$

$$= \frac{31.1}{\sqrt{20}} \quad (4.12)$$

$$= 6.96 \quad (4.13)$$

If we incorporate the fpc this becomes

$$\mathbf{SE}[\widehat{Y}] = \sqrt{1 - \frac{n}{N}} \times \frac{s_y}{\sqrt{n}} \quad (4.14)$$

$$= \sqrt{1 - \frac{20}{87}} \frac{31.1}{\sqrt{20}} \quad (4.15)$$

$$= 0.8776 \times \frac{31.1}{\sqrt{20}} \quad (4.16)$$

$$= 6.11 \quad (4.17)$$

This narrows the confidence interval for the estimate of the population mean height from (165.7,194.9) to (167.5,193.1).

This is a **design based** confidence interval, in which it is the sample design which foremost in our derivation of standard errors.

There is quite a lot of theory behind derivation of standard errors in the complex designs which arise in sample surveys.

The **survey** package in R implements the consequences of sample designs which will mean we can use the theory without having to derive it.

For example here's a call to some **survey** package functions which gives us the effect of the fpc in our example:

```
samplemembers$bigN <- bigN
ss <- svymean(height~1, design=svydesign(ids=~1, fpc=~bigN, data=samplemembers))
ss

##           mean      SE
## height 180.3 6.1089

confint(ss)

##           2.5 %    97.5 %
## height 168.3268 192.2732
```

Table 4.1: Sample Frame (N rows)

Unit ID	Auxiliary	Selection	Selection
	Data	Probability	Indicator
1	$\mathbf{X}_1$	π_1	I_1
2	$\mathbf{X}_2$	π_2	I_2
3	$\mathbf{X}_3$	π_3	I_3
$\vdots$	$\vdots$	$\vdots$	$\vdots$
i	$\mathbf{X}_i$	π_i	I_i
$\vdots$	$\vdots$	$\vdots$	$\vdots$
N	$\mathbf{X}_N$	π_N	I_N

If you're interested in the theory take a look at a text like Cochran (1977) or Särndal et al. (1992).

4.3 The survey inference process

Sampling begins with the sample frame - usually this is a **list** of the members, or **units**, which make up the population.

The list contains at minimum some **identifier** of each unit, which enables it to be found/contacted/measured once it has been selected, and possibly a lot of other information too.

Data that we have on the frame is called **auxiliary information**: it is data that is available about a unit whether or not it has been sampled. Some sample designs rely on this auxiliary data during the process of sample selection.

In probabilistic (random) sampling each unit i has a known nonzero chance of selection π_i , which may depend on the auxiliary data in some way. These probabilities don't have to be equal, just known and nonzero.

In random sampling we generate a random variable I_i for each population member, based on the sample design, to decide whether it is selected into the sample or not:

$$I_i = \begin{cases} 1, & \text{if unit } i \text{ is selected} \\ 0, & \text{otherwise} \end{cases} \quad (4.18)$$

The sample size is just the sum of these random variables

$$n = \sum_{i=1}^N I_i$$

So we start with the N rows of the sample frame (Table 4.1) generate the sample indicators I_i , and reduce to the n rows of data in Table 4.2. For each member k of the sample we know its population identifier i_k , its auxiliary information $\mathbf{x}_k$, selection probability π_k , and we also collect and record the survey data $\mathbf{y}_k$.

The **expected value** of each indicator variable is π_i

$$E[I_i] = 1 \times \Pr(I_i = 1) + 0 \times \Pr(I_i = 0) \quad (4.19)$$

$$= 1 \times \pi_i + 0 \times (1 - \pi_i) \quad (4.20)$$

$$= \pi_i \quad (4.21)$$

Table 4.2: Sample Data (n rows)

Sample Unit	Unit ID	Auxiliary Data	Selection Probability	Sample Weight	Sample Data
1	i_1	$\mathbf{x}_1$	π_1	w_1	$\mathbf{y}_1$
2	i_2	$\mathbf{x}_2$	π_2	w_2	$\mathbf{y}_2$
3	i_3	$\mathbf{x}_3$	π_3	w_3	$\mathbf{y}_3$
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
k	i_k	$\mathbf{x}_k$	π_k	w_k	$\mathbf{y}_k$
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
n	i_n	$\mathbf{x}_n$	π_n	w_n	$\mathbf{y}_n$

Each selected sample member carries a **sample weight** which is the number of units in the population that that unit represents. These weights are defined as the inverses of the selection probabilities.

$$W_i = \frac{1}{\pi_i}$$

Consider some characteristic Y_i which each population member has. e.g. in a survey of households we might be measuring the amount of money Y_i that household i spends on internet streaming services. The total amount of money spent by all households in the population combined is

$$Y = \sum_{i=1}^N Y_i$$

However we only know Y_i for some households: those which are selected, i.e. those which have $I_i = 1$.

The **Horwitz-Thompson** (HT) estimator of the population total is

$$\widehat{Y} = \sum_{i=1}^N I_i W_i Y_i \quad (4.22)$$

$$= \sum_{k=1}^n w_k y_k \quad (4.23)$$

i.e. a weighted sum of the observed household spending values y_k just for the n selected households. ($I_i = 0$ for any unselected households, so the sum of N terms simplifies to a sum of n terms, just over the sample members.)

We can see that the HT estimator is **unbiased**:

$$E[\widehat{Y}] = \sum_{i=1}^N E[I_i] W_i Y_i \quad (4.24)$$

$$= \sum_{i=1}^N \pi_i W_i Y_i \quad (4.25)$$

$$= \sum_{i=1}^N \pi_i \frac{1}{\pi_i} Y_i \quad (4.26)$$

$$= \sum_{i=1}^N Y_i \quad (4.27)$$

$$= Y \quad (4.28)$$

i.e. although the observed value $\widehat{Y}$ from a given sample may not equal the true value Y , when averaged over all possible samples we can see that on average it does equal the true population value.

We also want the standard error $\mathbf{SE}[\widehat{Y}]$ to be small, so that we have precise estimates.

We're all familiar with the idea that the larger the sample size n the smaller the standard error. We now know that the standard error is going to crush down to zero in finite populations the closer we get to a census.

There's a further advantage we can get by making a good choice of **sample design**.

Two different methods of selecting a sample of the **same sample size** n may lead to very different values of the standard error, and thus confidence intervals of very different widths.

Incidentally, the expected value of the sum of the sample weights among sample members is equal to the population size:

$$\sum_{k=1}^n w_k = \sum_{i=1}^N I_i W_i \quad (4.29)$$

$$E \left[\sum_{k=1}^n w_k \right] = \sum_{i=1}^N E[I_i] W_i \quad (4.30)$$

$$= \sum_{i=1}^N \pi_i W_i \quad (4.31)$$

$$= \sum_{i=1}^N \pi_i \frac{1}{\pi_i} \quad (4.32)$$

$$= \sum_{i=1}^N 1 \quad (4.33)$$


$$= N \quad (4.34)$$

Here's one way to think about sample weights. We have a population of size N we want to make estimates about: what we'd really like is a dataset of all N individuals. Instead we only have $n < N$ individuals. So we (conceptually) replicate each of the n sampled individuals by their weight: this makes a new synthetic dataset which does have the N entries we're interested in. Each row is identical to a number of others, however in the real population there are many individuals who sometimes very closely resemble each other. So the properties of our synthetic population may be good enough for estimating the properties of the true population.

The weight of a sample member indicates how many members of the population that sample member represents. If we had a population of size $N = 100$ and a simple random sample of size $n = 20$ then each sample member represents $w = 100/20 = 5$ people in the population.

In many designs the weights differ between individuals: high weights mean that the sample member represents many other individuals. A weight of 1 means that sample member represents themselves alone.

There's no problem with weights being fractions - since we only conceptually but don't **actually** replicate the data.



ID	Address	Value (\$M)
6451	1 Chestnut Place	1.201
6452	2 Chestnut Place	0.780
8773	2A Chestnut Place	0.520
7164	3 Chestnut Place	2.530
6125	4 Chestnut Place	0.900

ID	Address	Value (\$M)	Prob. π_i	Selected, I_i
6451	1 Chestnut Place	1.201	0.4	0
6452	2 Chestnut Place	0.780	0.4	0
8773	2A Chestnut Place	0.520	0.4	1
7164	3 Chestnut Place	2.530	0.4	0
6125	4 Chestnut Place	0.900	0.4	1

4.4 Example

We might have a small street of five houses. Our frame is just the list of addresses. We might know the household rateable value from the City Council list of dwellings.

We might then select a simple random sample of $n = 2$ houses from the $N = 5$ in the small street. The probability of selection is $\pi_i = 2/5$ for each house. We draw our random sample and select house number 2A and house number 5 into our sample.

We visit each selected house and ask how much they spend y_k on all their internet streaming services each month.

Each house in the sample has a weight of $2.5 = 5/2$: i.e. it represents itself and 1.5 other houses in the street. The sum of the weights of the two selected houses is equal to the population size $N = 5$.

Our sample is really too small to do any sensible inference on: but let's do it anyway to show how things work. We calculate the mean and standard deviation of the monthly internet streaming values: $\bar{y} = \$47.00$, and $s_y = \$39.60$.

Our estimate of the mean monthly spend is $\bar{y} = \$47.0$, with a margin of error of

$$\text{MOE}[\hat{Y}] = Z^* \sqrt{1 - \frac{n}{N}} \times \frac{s_y}{\sqrt{n}} \quad (4.35)$$

$$= 1.96 \sqrt{1 - \frac{2}{5}} \times \frac{\$39.60}{\sqrt{2}} \quad (4.36)$$

$$= \$42.51 \quad (4.37)$$

which gives a 95% confidence interval of (\$4.49,\$89.51) (although the use of the large sample value $Z^* = 1.96$ is very much **not** justified here.)

In finite populations we are sometimes interested in population totals - which don't make any sense in infinite populations. Now we have an estimate of the mean spend per house, we can scale it up by the population size to see how much per month is being spent by all of the

	ID	Address	Value (\$M)	Weight w_k	Internet spend, y_k
3	8773	2A Chestnut Place	0.52	2.5	75
5	6125	4 Chestnut Place	0.90	2.5	19

houses combined:

$$\widehat{Y} = N\widehat{\bar{Y}} \quad (4.38)$$

$$= 5 \times \$47.00 \quad (4.39)$$

$$= \$235.00 \quad (4.40)$$

and we can scale the confidence interval up in the same way

$$\mathbf{CI}[\widehat{Y}] = N\mathbf{CI}[\widehat{\bar{Y}}] \quad (4.41)$$

$$= 5 \times (\$4.49, \$89.51) \quad (4.42)$$

$$= (\$22.45, \$447.55) \quad (4.43)$$

Chapter 5

Sample Designs

In this Chapter we consider the various sample designs in common use, the reasons for using each one, and methods for selecting from suitable frames.

We'll use a variety of datasets.

5.1 Census

The simplest of all! We select every unit in the population, so $n = N$.

The probability of selection is

$$\pi = \frac{n}{N} = \frac{N}{N} = 1$$

i.e. every unit is selected with certainty, and the sample weights are all

$$w_i = \frac{1}{\pi_i} = \frac{1}{1} = 1$$

so each unit in the sample represents itself alone.

Design considerations

None: we're sampling all N members of the population.

5.2 Simple Random Sampling Without Replacement (SRSWOR)

The Simple Random Sample WithOut Replacement is the most fundamental of all sample designs, and it appears nested inside many more complex designs as we'll see later.

We have a population of size N which is fully enumerated (i.e. fully listed), and we want to select a sample of size n .

The defining property of a SRSWOR is that all samples of size n have an equal chance of selection. A **consequence** of this definition is that all population units have the same chance of selection.

There are

$$\binom{N}{n} = {}^N C_n = \frac{N!}{n!(N-n)!}$$



possible samples. The R function `choose(n,k)` can be used for evaluating the number of possible samples. e.g. the number of ways to select $n = 20$ units from $N = 87$ is

```
choose(87,20)
```

```
## [1] 2.375456e+19
```

The probability of selection of each unit is

$$\pi_i = \frac{n}{N}$$

leading to sample weights

$$w_i = \frac{N}{n}$$

The SRSWOR is used when the **costs** of different samples are not known to differ (e.g. there are no units that are more costly to visit than others), and where there is no informative auxiliary information.

If we have a sample frame stored as an R data frame `sframe` with a column `ID` uniquely identifying the units, then we can select an SRSWOR as follows:

```
bigN <- nrow(sframe)
samp <- sframe[sample(bigN,n),]
samp$weight <- bigN/n
```

or

```
bigN <- nrow(sframe)
IDset <- sample(sframe$ID,n)
samp <- sframe[sframe$ID %in% IDset,]
samp$weight <- bigN/n
```

or as a function

```
select.srswor <- function(df, n) {
  # SRSWOR
  bigN <- nrow(df)
  if(n <= bigN) {
    samp <- df[sample(bigN,n,replace=FALSE),]
    samp$bigN <- bigN
    samp$weight <- bigN/n
  } else {
    stop("Cannot have a SRSWOR with n>N")
  }
  return(samp)
}
```

Design considerations

- We only need to decide on the sample size n (see later).

Example

Select $n = 5$ units from a frame listing of airports with IATA codes sourced from Kaggle <https://www.kaggle.com/datasets/syedasimalishah/airports/data>

The top ten rows of the dataset of $N = 6072$ are shown here:

Sir



ID	IATA	Name	City	Country	Latitude	Longitude
1	AAA	Anaa Airport	Anaa	French Polynesia	-17.352600	-145.509995
2	AAC	El Arish International Airport	El Arish	Egypt	31.073299	33.835800
3	AAE	Rabah Bitat Airport	Annaba	Algeria	36.822201	7.809174
4	AAF	Apalachicola Regional Airport	Apalachicola	United States	29.727501	-85.027496
5	AAH	Aachen-Merzbrück Airport	Aachen	Germany	50.823055	6.186389
6	AAK	Buariki Airport	Buariki	Kiribati	0.185278	173.636993
7	AAL	Aalborg Airport	Aalborg	Denmark	57.092759	9.849243
8	AAM	Malamala Airport	Malamala	South Africa	-24.818100	31.544600
9	AAN	Al Ain International Airport	Al Ain	United Arab Emirates	24.261700	55.609200
10	AAO	Anaco Airport	Anaco	Venezuela	9.430225	-64.470726

Table 5.1: Sample of airports (SRSWOR)

ID	IATA	Name	City	Country	Latitude
3593	NTD	Point Mugu Naval Air Station (Naval Base Ventura Co)	Point Mugu	United States	34.12030
1472	EUM	Neumünster Airport	Neumuenster	Germany	54.07944
2512	KHZ	Kauehi Airport	Kauehi	French Polynesia	-15.78080
197	ANC	Ted Stevens Anchorage International Airport	Anchorage	United States	61.17440
4710	SRN	Strahan Airport	Strahan	Australia	-42.15500

And here is a SRSWOR of $n = 5$ airports

```
samp <- airports |> select.srswor(n=5)
```

5.3 Simple Random Sampling With Replacement (SRSWR)

This design is rarely used, but is the simplest form of Probability Proportional to Size With Replacement (PPSWR) sampling, but with the sizes all equal.

Sampling proceeds in principle by selecting one unit at a time where each unit has a $1/N$ chance of being selected each time. After n draws the probability a unit has been selected at least once is

$$p_i = 1 - \left(1 - \frac{1}{N}\right)^n$$

(This is 1 minus the chance that a unit is not selected on any of the n sampling occasions.)

If n is tiny compared to N , p_i is approximately n/N which is the SRSWOR selection probability. In practice we assign a weight of N/n to each unit as it is selected, so if unit i is selected m_i times then it ends up with a weight of

$$w_i = m_i \frac{N}{n}$$

If we use the approximate weights calculated in this way we have what is known as the Hansen-Hurwitz estimator, rather than the Horvitz-Thompson estimator. However we think of the weights in the same way and this distinction need not concern us.

Design considerations

As in a SRSWOR, we only need to decide on the sample size n .

Table 5.2: Sample of airports (SRSWOR)

ID	IATA	Name	City	Country	Latitude	Longitude	nsampled
805	CCB	Cable Airport	Upland	United States	34.11160	-117.6880	1
1063	CUD	Caloundra Airport	Caloundra	Australia	-26.80000	153.1000	1
1077	CVE	Coveñas Airport	Coveñas	Colombia	9.40092	-75.6913	1
2738	LAY	Ladysmith Airport	Ladysmith	South Africa	-28.58170	29.7497	1
4804	SXN	Sua Pan Airport	Sowa	Botswana	-20.55340	26.1158	1

Example

Here is a SRSWR of $n = 5$ airports: note that we have aggregated across any instances of multiple selections, so that in the output sample list each unit appears only once no matter how often it was selected.

```
bigN <- nrow(airports)
n <- 5
samp <- airports[sample(bigN,n,replace=TRUE),]
samp$weight <- bigN/n
samp <- samp |>
  group_by(ID) |>
  summarise(across(-weight, first), nsampled=n(), weight=sum(weight))
```

or as a function

```
select.srswr <- function(df, n) {
  # SRSWR
  bigN <- nrow(df)
  df$rownames <- 1:bigN
  samp <- df[sample(bigN,n,replace=TRUE),]
  samp$bigN <- bigN
  samp$weight <- bigN/n
  samp <- samp |>
    group_by(rownames) |>
    summarise(across(-weight, first), nsampled=n(), weight=sum(weight)) |>
    as.data.frame() |>
    ungroup() |>
    select(-rownames)
  return(samp)
}

samp <- airports |> select.srswr(n=5)
```

5.4 Probability proportion to size Sampling (PPSWR)

This sample design generalises the SRSWR design to unequal probabilities of selection. This design recognises that some units may have more influence on our estimates because of some **size** property that is known (it must be on the frame). We therefore give these units higher probabilities of being selected, and therefore lower weight.

We draw n units, one at a time, but with probabilities which are proportional to the size measure X_i : thus the probability of selection of unit i on each of the n draws from the

Country	Code	Year	Population	Prob
Afghanistan	AFG	2023	41454762	0.0051249
Albania	ALB	2023	2811660	0.0003476
Algeria	DZA	2023	46164222	0.0057071
American Samoa	ASM	2023	47544	0.0000059
Andorra	AND	2023	80869	0.0000100
Angola	AGO	2023	36749909	0.0045432
Anguilla	AIA	2023	14435	0.0000018
Antigua and Barbuda	ATG	2023	93331	0.0000115
Argentina	ARG	2023	45538402	0.0056297
Armenia	ARM	2023	2943398	0.0003639

population is

$$p_i = \frac{X_i}{\sum_{j=1}^N X_j}$$

Approximating the probability of selection π after n draws as np_i leads to the sample weights

$$w_i = \frac{1}{np_i}$$

In a PPSWR sample these weights are not guaranteed to add up to the population size.

Design considerations

We need to decide on:

- the size variable X_i (which must be known for every unit in the population), and
- the sample size n .

The size variable X_i should correlate in some way with the population characteristics we want to estimate. If we're interested in estimating the total economic output of businesses, then the amount of tax each business paid last year would be a good size variable. But it wouldn't be helpful if we were interested in estimating the proportion of businesses with family friendly employment practices.

With replacement designs lose some efficiency because their estimators do not incorporate a finite population correction in their standard errors. For small sampling fractions (i.e. small n/N) this is not a significant drawback.

Example

We have a list of $N = 236$ countries and their 2023 populations: the first 10 are listed below

Select a PPSWR sample of $n = 10$ countries, with probability proportional to their population sizes.

```
set.seed(111)
bigN <- nrow(pop)
n <- 10
samp <- pop[sample(bigN,n,prob=pop$Prob,replace=TRUE),]
samp$weight <- 1/(n*samp$Prob)
samp <- samp |>
  group_by(Country) |>
  summarise(across(-weight, first), nsampled=n(), weight=sum(weight))
```

Country	Code	Year	Population	Prob	bigN	nsampled	weight
Bangladesh	BGD	2023	171466986	0.0211977	236	1	23.6
Brazil	BRA	2023	211140731	0.0261023	236	1	23.6
India	IND	2023	1438069597	0.1777819	236	2	47.2
Indonesia	IDN	2023	281190068	0.0347622	236	1	23.6
Japan	JPN	2023	124370947	0.0153754	236	1	23.6
Pakistan	PAK	2023	247504505	0.0305978	236	1	23.6
South Africa	ZAF	2023	63212386	0.0078147	236	1	23.6
United States	USA	2023	343477330	0.0424625	236	2	47.2

or as a function

```
select.ppswr <- function(df, n, sizevar) {
  # PPSWR
  bigN <- nrow(df)
  df$rownames <- 1:bigN
  samp <- df[sample(bigN,n,prob=df[,sizevar],replace=TRUE),]
  samp$bigN <- bigN
  samp$weight <- bigN/n
  samp <- samp |>
    group_by(rownames) |>
    summarise(across(-weight, first), nsampled=n(), weight=sum(weight)) |>
    as.data.frame() |>
    ungroup() |>
    select(-rownames)
  return(samp)
}

samp <- pop |> select.ppswr(n=10, sizevar="Population")
```

5.5 Linear Systematic Random Sample (LSRS)

This design is often used in cases where the population is not fully enumerated, but is presented to us in the form of a queue. For example we might want to select one in ten passengers cars checking in for a ferry crossing. Another example is where there is information in the ordering of a list: for example a list might be sorted from youngest to oldest, or sorted from north to south. We can space out our sample down the list, avoiding the chance that our sample doesn't have a good spread of ages, or a good spread of latitudes.

We select n units from N by first calculating the **skip** k , which is the nearest integer at or below N/n . We then choose at random one unit from the first k of our list or queue, say that is unit number j . We then select units $j, j+k, j+2k, j+3k$, etc.

Selecting one unit at random in the first k (rather than just selecting the **first** unit, as we might naively think of doing) means that every unit in the list has a chance of selection.

The selection probabilities are

$$\pi_i = \frac{n}{N}$$

and weights are

$$w_i = \frac{N}{n}$$

Data from this design are often analysed as if they were drawn as a SRSWOR, but in fact this is a special case of a one-stage cluster sample (1SC, see below).

Linear Syst



Note that if there are N units but N is not an exact multiple of the skip k , then the sample size will differ depending on which of the first k units we select to initiate sampling. For example if we wanted to select $n = 3$ units by LSRS from a list of $N = 11$ units, then the skip size is $k = \text{floor}(11/3) = \text{floor}(3.67) = 3$. If we start the process by selecting the first unit, we'd end up with the sample of 4 units, $\{1, 4, 7, 10\}$, but if we start by selecting the 4th unit we'd end up with only 3 units, $\{4, 7, 10\}$.

Another common reason to take an LSRS is in household surveys: we're selecting dwellings from a small area, and for reasons of confidentiality (we don't want members of different selected households to be aware that they are both in the same survey), or to prevent collusion between neighbouring selected households, we space the sample out along the street using a LSRS to separate the selected dwellings as much as possible.

An LSRS is also a convenient way to sample if you have a printed list. After choosing the random start location among the first k units you can skip down the list at fixed intervals - something that is convenient and simple to do on a printed list.

Design considerations

- Choose a sample size n
- Consider the ordering of the list and whether units differ systematically through the list

An LSRS has one risk, and that is a periodicity in the frame that you are unaware of. For example if we are selecting apartments from a building to inspect for structural weaknesses, and select every fourth one (after a random start), we had better be sure that there aren't exactly four apartments per floor: otherwise we'll have inadvertently selected a column of apartments stacked one above the other, e.g. all in the north-east corner. We'll miss any systematic issues that occur with the north-west apartments in that case.

Such unknown periodicities are unlikely in most cases, but we should be aware that they can occur.

More importantly if the list is ordered in some way, e.g. by size of location, then a LSRS guarantees a spread across the variable that defines the ordering. This can be highly desirable since we are sure to be sampling across the full diversity of the ordering variable. A good example is sampling passengers seated in an aircraft - if we take an LSRS using a list of passengers ordered by row and seat number we'll be guaranteed to get a spread of business class, premium economy and economy class passengers in our sample.

Example

To select a LSRS of airports, let's order the airports by latitude, and then take a LSRS of $n = 5$ units.

```
bigN <- nrow(airports)
airports <- airports[order(airports$Latitude),]
n <- 5
k <- floor(bigN/n)
cat("Skip size is k =", k, "\n")
```

```
## Skip size is k = 1214
```

```
i0 <- sample(k, 1) # first selected unit
cat("Starting sampling at unit", i0, "\n")
```

```
## Starting sampling at unit 305
```

Table 5.3: Sample of airports (SRSWOR)

ID	IATA	Name	City	Country	Latitude	Longitude	weight
5304	VAO	Suavanao Airport	Suavanao	Solomon Islands	-7.58556	158.73100	1214.4
4746	STV	Surat Airport	Surat	India	21.11410	72.74180	1214.4
1714	GIB	Gibraltar Airport	Gibraltar	Gibraltar	36.15120	-5.34966	1214.4
1638	FYN	Fuyun Koktokay Airport	Fuyun	China	46.80417	89.51201	1214.4
5277	UTO	Indian Mountain LRRS Airport	Indian Mountains	United States	65.99280	-153.70399	1214.4

```

n <- (bigN-i0)%/%k + 1 # recalculate the sample size
cat("Actual sample size is",n,"\n")

## Actual sample size is 5

samp <- airports[i0 + (0:(n-1))*k,]
samp$weight <- bigN/n

select.lsr <- function(df, n) {
  # LSR
  bigN <- nrow(df)
  k <- floor(bigN/n)
  i0 <- sample(k,1) # first selected unit
  n <- (bigN-i0)%/%k + 1 # recalculate the sample size
  samp <- df[i0 + (0:(n-1))*k,]
  samp$bigN <- bigN
  samp$weight <- bigN/n
  return(samp)
}

samp <- airports |> select.lsr(5)

```

5.6 Stratified Simple Random Sample (STSRs)

In this design we take our population of N units and break it into H subgroups using our auxiliary data from the frame. We then draw separate samples from **each of the subgroups**.

To draw a stratified sample each population unit needs to be assigned in advance to one, and only one, of these subgroups - which are known as **strata** (singular: **stratum**).

There are N_h units in stratum h , and all of these stratum-level population sizes add to the total population size:

$$N = \sum_{h=1}^H N_h$$

Dividing both sides of this equation by N we have the stratum fractions $F_h = N_h/N$: these are the proportions of the **population** that belong to the various strata, and they add up to 1:

$$1 = \sum_{h=1}^H \frac{N_h}{N} = \sum_{h=1}^H F_h$$

For example, we can assign all residential addresses in New Zealand to exactly one of a defined set of non-overlapping regions, such as regional council areas. Or we might just have

two strata: North Island, South Island (neglecting any addresses on offshore islands). There are $N = 2.4$ million addresses in New Zealand, of which $N_1 = 1.72$ million in the North Island ($F_1 = 0.74 = 74$) and $N_2 = 0.62$ million are in the South Island ($F_2 = 0.26 = 26$).

We then divide or **allocate** our total sample size n to the strata. Stratum h is allocated n_h units, and the total sample size is then the sum of these:

$$n = \sum_{h=1}^H n_h$$

We then carry out independent sampling in each of the strata.

The way we choose to define the strata and the way we then choose to allocate the sample across the strata depend on our goals.

Stratified sampling is done for a number of reasons

1. **Estimation Efficiency** - efficiency is the word statisticians use when considering how we can get standard errors which are as **small as possible**.

In business surveys we often stratify because businesses differ from each other dramatically, particularly in size measures such as turnover or size of workforce. That means a small number of businesses (e.g. Spark, Air New Zealand, Fonterra) have a very large effect on estimates, and they differ dramatically from each other as well. If they happen to be included or happen to be missed from our samples that means we'd get estimates which differ very much between samples. To protect ourselves against this we stratify by business size (using tax data). The largest businesses are in the top stratum, which is small, but a census is always done of those businesses (the top stratum is called a **full coverage stratum**). The smallest businesses are in the lowest stratum, and they are selected with low probability: since they resemble each other more, and one can stand in for another.

Having a strata designed so that the units within the strata resemble each other strongly (one can stand in for another easily), but where units in different strata differ strongly from each other. For the same sample size as a SRSWOR stratified designs can sometimes achieve much much smaller standard errors.

2. **Subgroup estimates** - we may stratify by region because we want to make estimates for each region separately. In that case we need to ensure that our sample is large enough in each stratum to get reliable estimates. A national SRSWOR runs the risk that one region might happen to miss out on getting much sample.
3. **Operational considerations** - we may stratify by region because that is the easiest way to manage an interview workforce - we may have interviewers based in a number of regional cities, and having them interview in their local areas is the best way to organise that workforce.

Two common methods for allocating the n sample units across the H strata are:

1. **Equal allocation**


$$n_h = \frac{n}{H}$$

Every stratum gets the same number of units allocated - independent of the stratum size. That means we get different weights in each stratum:

$$w_h = \frac{N_h}{n_h} = H \frac{N_h}{n}$$

This is sensible when we want to have high precision estimates in each stratum separately.

Stratified S



h	Education Region	N_h	F_h
1	Bay of Plenty, Waiairiki	185	0.0739
2	Canterbury, Chatham Islands	283	0.1131
3	Hawke's Bay, Tairāwhiti	174	0.0695
4	Nelson, Marlborough, West Coast	127	0.0507
5	Otago, Southland	228	0.0911
6	Tai Tokerau	151	0.0603
7	Taranaki, Whanganui, Manawatu	234	0.0935
8	Tāmaki Herenga Manawa	187	0.0747
9	Tāmaki Herenga Tāngata	197	0.0787
10	Tāmaki Herenga Waka	179	0.0715
11	Waikato	277	0.1107
12	Wellington	281	0.1123

2. Proportional allocation

$$n_h = nF_h = n \frac{N_h}{N}$$

Here large strata get a lot of units allocated to them, and small strata get fewer units. The consequence of this design is that the weights of all units are equal, no matter which stratum they come from:

$$w_h = \frac{N_h}{n_h} = \frac{N}{n}$$

This allocation is also called **self weighting**, and the weights are the same as in SRSWOR. (Though note: even though the weights are the same, the standard errors of estimates can differ **very significantly** from those from a SRSWOR.)

Design considerations

- Choose the stratification variable(s) X_i , decide the number of strata H , define the strata using the stratification variable(s), determine the stratum sizes N_h , and stratum fractions $F_h = N_h/N$.
- Choose the overall sample size n and **allocate** it across the strata.

Example

The 2026 directory of NZ schools available from <https://www.educationcounts.govt.nz/directories/list-of-nz-schools> lists 2503 schools of various types (after removing Schools with zero enrolments or with incomplete data).

Each School is associated with one of $H = 12$ Education Regions, with counts N_h and stratum fractions F_h listed below.

We can select a stratified sample using the following function which implements an independent simple random sample within each stratum with either equal or proportional allocation.

It implements PPSWR sampling within strata as well.

```
select.stsrs <- function(df, n, stratumvar,
                        allocation="equal",
                        nmin=2, # minimum number of units to select per stratum
                        method="SRSWOR", sizevar=NULL) {
  bigN <- nrow(df)
  if(!stratumvar%in%names(df)) {
```

Education Region, h	N_h	n_h	Weight, w_h	Selected Schools
Bay of Plenty, Waiariki	10	185	18.5	482;1170;1711;1758;1787;1796;1808;1904;1929;1984
Canterbury, Chatham Islands	10	283	28.3	354;3304;3329;3348;3357;3389;3450;3553;3572;3586
Hawke's Bay, Tairāwhiti	10	174	17.4	209;1668;1678;2566;2575;2578;2593;2657;2698;2705
Nelson, Marlborough, West Coast	10	127	12.7	292;294;301;3062;3071;3180;3233;3343;3690;6975
Otago, Southland	10	228	22.8	376;399;3737;3753;3824;3827;3943;3982;4017;4050
Tai Tokerau	10	151	15.1	1000;1014;1020;1039;1059;1060;1095;1104;1130;1667
Taranaki, Whanganui, Manawatu	10	234	23.4	2165;2167;2341;2365;2382;2416;2440;2451;2462;4123
Tāmaki Herenga Manawa	10	187	18.7	48;69;460;1210;1370;1376;1382;1385;1484;4133
Tāmaki Herenga Tāngata	10	197	19.7	32;571;767;1106;1223;1309;1394;1432;1492;6977
Tāmaki Herenga Waka	10	179	17.9	96;452;1201;1271;1275;1329;1339;1442;1455;4229
Waikato	10	277	27.7	1744;1753;1777;1790;1835;1844;2001;2008;2033;6941
Wellington	10	281	28.1	133;546;2742;2874;2941;2963;2967;3032;3045;3048

```

    stop("Stratum variable is not present in the data frame")
  }
  stratumproperties <- as.data.frame(table(df[,stratumvar]))
  names(stratumproperties) <- c(stratumvar,"Nh")
  if(allocation=="equal") {
    stratumproperties$nh <- pmin(stratumproperties$Nh,pmax(nmin,
      round(n/nrow(stratumproperties))))
  } else if(allocation=="proportional") {
    stratumproperties$nh <- pmin(stratumproperties$Nh,pmax(nmin,
  } else {
    stop("Invalid allocation")
  }
  df <- merge(df,stratumproperties,by=stratumvar)

  if(method=="SRSWOR") {
    samp <- do.call(rbind, by(df, df[,stratumvar],
      function(dfx) select.srswor(dfx,first(dfx$nh))))
  } else if(method=="PPSWR") {
    samp <- do.call(rbind, by(df, df[,stratumvar],
      function(dfx) select.ppswr(dfx,first(dfx$nh),sizevar=sizevar)))
  } else {
    stop("Sampling method invalid")
  }
  return(samp)
}

```

SRSWOR with Equal allocation

```
samp <- schools |> select.stsrs(n=120,
  stratumvar="Education_Region", allocation="equal")
```

SRSWOR with proportional allocation

```
samp <- schools |> select.stsrs(n=120,
  stratumvar="Education_Region",
  allocation="proportional", sizevar="Total")
```

PPSWR with Equal Allocation

Education Region, h	N_h	n_h	Weight, w_h	Selected Schools
Bay of Plenty, Waiariki	9	185	20.55556	145;545;558;707;1645;1745;1879;1921;19
Canterbury, Chatham Islands	14	283	20.21429	328;346;531;3275;3276;3285;3412;3462;3
Hawke's Bay, Tairāwhiti	8	174	21.75000	1669;2547;2574;2588;2590;2685;2702;27
Nelson, Marlborough, West Coast	6	127	21.16667	300;2811;3181;3207;3231;3530
Otago, Southland	11	228	20.72727	384;390;3716;3719;3731;3772;3812;3826;
Tai Tokerau	7	151	21.57143	1029;1030;1075;1085;1116;1666;2104
Taranaki, Whanganui, Manawatū	11	234	21.27273	454;2157;2160;2175;2265;2336;2366;244
Tāmaki Herenga Manawa	9	187	20.77778	78;1190;1224;1240;1282;1411;1416;1515;
Tāmaki Herenga Tāngata	9	197	21.88889	45;1107;1200;1254;1311;1316;1407;1500;
Tāmaki Herenga Waka	9	179	19.88889	748;1248;1277;1332;1369;1418;1472;155
Waikato	13	277	21.30769	130;135;146;158;159;671;1344;1771;1891
Wellington	13	281	21.61538	237;258;1189;1651;2669;2818;2876;2911;

Education Region, h	N_h	n_h	Selected Schools
Bay of Plenty, Waiariki	9	185	120;121;143;152;532;1699*;1730;1761;1766
Canterbury, Chatham Islands	10	283	310;312;327;333;335;348;3289;3407;3517;3570
Hawke's Bay, Tairāwhiti	10	174	216;217;227;230;2445;2564;2581;2628;2677;2688
Nelson, Marlborough, West Coast	8	127	289;296*;1148;1587;3189;3208;3229;3472
Otago, Southland	9	228	369;374;376;400;533;548;586;3967*;4054
Tai Tokerau	9	151	5;8;13*;15;21;1032;1050;1113;1648
Taranaki, Whanganui, Manawatū	10	234	202;203;961;2164;2205;2346;2389;2397;2465;2477
Tāmaki Herenga Manawa	10	187	65;1190;1319;1370;1383;1402;1462;2085;3630;6760
Tāmaki Herenga Tāngata	9	197	46;725;767;1223;1304;1341;1396*;1406;1572
Tāmaki Herenga Waka	10	179	90;91;631;949;1203;1247;1423;1435;1460;1547
Waikato	9	277	126;132*;451;577;615;1693;1946;2021;3612
Wellington	10	281	240;253;255;275;276;478;498;2832;2878;3036

```
samp <- schools |> select.stsrs(n=120,
                               stratumvar="Education_Region",
                               allocation="equal",
                               method="PPSWR", sizevar="Total")
```

(Schools selected multiple times are indicated with a *)

5.7 One stage cluster sampling (1SC)

One stage cluster sampling superficially resembles stratified sampling. As in stratified sampling we divide the population into groups, in this case called **clusters**, assign each member of the population to one and only one group. However rather than sample members from each cluster, we instead draw a sample of clusters (by SRSWOR, or sometimes by PPSWR). We then take a **census** of each selected cluster.

In the Schools example this would could define small regions across the country, select a random sample of those regions, and then carry out a census of schools in each selected region.

Unlike stratified sampling, which is often **more efficient** than simple random sampling (i.e. it has smaller standard errors for the same sample size), cluster sampling is **less efficient**: i.e. its standard errors are worse (larger) than we'd get with a SRS with the same size.

We can see why the standard errors are worse when we consider **why** we take cluster samples:

the reason is predominantly to controlling costs. If we want to do a household survey in a city like Wellington we could list all of the residential addresses, and then draw a simple random sample of those dwellings. That would mean having a sample that is scattered all over the city - and if we have to visit each dwelling to collect the data, then that means a lot of travelling for the interviewers: especially if they have to visit some houses multiple times to make contact with the residents. So it's more **convenient** to select a set of houses on the same street and survey them - we can go from one to the next without much travel time.

The trouble is that people who live near each other tend to resemble each other: they have similar incomes, jobs, family sizes etc. That means that for each additional house you survey in a street you're not really getting as much new information as if you were to survey an additional house somewhere else in the city.

For that reason cluster samples have to be larger than SRS to achieve the same accuracy of their estimates.

Note that it makes no sense to cluster if there is no saving to be made regarding cost or some other equivalent operational consideration. If we are doing a phone survey, where no travel at all is involved, then clustering doesn't make any sense.

Our notation becomes slightly more complex here: we have N clusters, labelled $i = 1, \dots, N$. Then in cluster i we have M_i units, labelled $j = 1, \dots, M_i$. There is a total of

$$M = \sum_{i=1}^N M_i$$

units in the population.

The selection probability of a unit is simply the probability that the cluster they are a member of is selected. So if we select clusters by SRSWOR then unit j in cluster i has selection probability

$$\pi_{ij} = \frac{n}{N}$$

leading to a weight of

$$w_{ij} = \frac{N}{n}$$

where again we note that n is the number of **clusters** selected, and N is the number of **clusters** in the population.

Design considerations


- We need to decide how to break our populations into clusters, and also the sizes of those clusters
- We need to choose an overall sample size.

Clusters can vary in size, so we even though we can control the number of clusters we select, we can't control the overall sample size: we might select by chance a set of very large clusters, or a set of very small clusters.

In stratified sampling we want to make the strata as different from one another as possible, and to be as internally homogeneous.

In cluster sampling it's the reverse situation: we want our clusters to be internally diverse, and for that diversity to be present as far as possible within every cluster. This is because we're only selecting a few clusters, and they have to stand in for all of the unselected clusters. So they'd better reflect the full population diversity.

One



Territorial Authority, i	M_i	Selected Schools
Hamilton City	59	129;130;131;132;135;136;138;139 ...
Kawerau District	4	651;655;661;1770
Manawatu District	26	197;199;2333;2338;2341;2354;2356;2360 ...
Ruapehu District	22	169;176;183;1617;1845;1847;1961;1982 ...
Timaru District	28	352;354;357;358;359;360;361;1611 ...
Wairoa District	14	214;1616;1668;1669;1675;1676;1677;1678 ...
Waitematā Local Board Area	21	48;50;51;53;62;921;1210;1220 ...
Western Bay of Plenty District	25	117;123;1185;1717;1758;1765;1794;1811 ...
Whau Local Board Area	24	42;78;83;84;781;941;1206;1212 ...
Ōtorohanga District	12	157;1688;1736;1771;1779;1783;1793;1853 ...

Example

There are $N = 87$ Territorial Authorities (TAs) in New Zealand (these are local government districts). Each TA has M_i schools, so the total number of schools in the population is

$$M = \sum_{i=1}^N M_i$$

We could select a SRS of $n = 10$ TAs and then do a census of the Schools in each selected TA. Here n is the sample size of clusters, not the sample size of schools. If cluster i contains M_i schools then the total number of schools selected is

$$m = \sum_{i \in \text{sample}} M_i$$

A function to do this selection in R is:

```
select.1sc <- function(df, n, clustervar,
                        method="SRSWOR", sizevar=NULL) {
  bigM <- nrow(df)
  if(!clustervar%in%names(df)) {
    stop("Cluster ID variable is not present in the data frame")
  }
  clusterproperties <- as.data.frame(table(df[,clustervar]))
  names(clusterproperties) <- c(clustervar, "Mi")
  if(method=="SRSWOR") {
    sampcluster <- clusterproperties |>
      select.srswor(n=n)
  } else if(method=="PPSWR") {
    sampcluster <- clusterproperties |>
      select.ppswr(n=n, sizevar="Mi")
  } else {
    stop("Sampling method invalid")
  }
  samp <- merge(df, sampcluster, by=clustervar)

  return(samp)
}

samp <- schools |> select.1sc(n=10, clustervar="Territorial_Authority")
```

5.8 Two stage cluster sampling (2SC)

The obvious drawback of the One Stage cluster design is that we are getting a lot of information from clusters which may be very homogeneous. This drives the inefficiency of the design: we're wasting a lot of our sampling effort.

An obvious solution is to select a set of clusters, and then only take a sample of units within those clusters, rather than sampling all of them.

This breaks the sample design into two stages

- **Stage 1** - select a set of clusters, which we refer to as **Primary Sampling Units** (PSUs);
- **Stage 2** - select a set of units within the selected clusters. These units are called **Secondary Sampling Units** (SSUs).

Sampling at both stages can be by any method, though SRSWOR or PPSWR are the most common.

We use the same notation as in one stage clustering: there are N PSUs of which select n , and in PSU i there are M_i SSUs, of which we select m_i . The total number of SSUs in all of the PSUs in the population is

$$M = \sum_{i=1}^N M_i$$

whereas the total sample size of SSUs is

$$m = \sum_{i \in \text{Sample}} m_i$$

where we just sum over the selected PSUs.

If we use SRSWOR and select n PSUs from the N available then the probability of selection of a PSU is

$$\pi_i^{(1)} = \Pr(\text{PSU } i \text{ is selected}) = \frac{n}{N}$$

leading to the first stage weight of

$$w_i^{(1)} = \frac{1}{\pi_i^{(1)}} = \frac{N}{n}$$

If PSU i is selected we then select m_i units (SSUs) out of the M_i available. The probability that SSU j is selected given that PSU i has been selected is

$$\pi_{ij}^{(2)} = \Pr(\text{SSU } j \text{ is selected given that PSU } i \text{ was selected}) = \frac{m_i}{M_i}$$

leading to the second stage weight of

$$w_{ij}^{(2)} = \frac{1}{\pi_{ij}^{(2)}} = \frac{M_i}{m_i}$$


The overall probability that SSU j in PSU i is selected is the product of the first and second stage selection probabilities:

$$\pi_{ij} = \pi_i^{(1)} \pi_{ij}^{(2)} = \frac{n}{N} \times \frac{m_i}{M_i}$$

and its overall weight is the product of the first and second stage weights

$$w_{ij} = w_i^{(1)} w_{ij}^{(2)} = \frac{N}{n} \times \frac{M_i}{m_i}$$

Two Sta



For example I might have $N = 100$ small areas covering Wellington, and I select $n = 10$ of them. Then in one of those selected areas there might be $M_i = 120$ houses, of which I select $m_i = 15$. The first and second stage probabilities are

$$\pi_i^{(1)} = \frac{10}{100} = 0.100 \quad \text{and} \quad \pi_{ij}^{(2)} = \frac{15}{120} = 0.125$$

So the chances that a house in PSU i is selected is

$$\pi_{ij} = (0.100)(0.125) = 0.0125$$

with total weight

$$w_{ij} = \frac{100}{10} \times \frac{120}{15} = 10 \times 8 = 80$$

There are two special cases of two stage clustering which reduce to simpler designs

- A first stage census: $n = N$: we sample from every PSU. This is equivalent to **stratified sampling**, where we sample from every stratum.
- A second stage census: $m_i = M_i$: we take a census of all SSUs in each selected cluster. This is equivalent to **one stage clustering**.

Design considerations

We have a lot of decisions to make:

- How many PSUs (clusters) N , and how we define them
- The number of PSUs to select n
- The number of SSUs to select within each selected PSU. Rather than state the sample size m_i for each cluster, this is often specified as the **proportion** of the cluster size we'll select. e.g. we might select 1 in 10 of the SSUs.

Example

Let's return to the schools example. We'll select $n = 10$ TAs, but then select 20% of the schools in each selected TA - although we'll also ensure we have at least 2 schools in each selected cluster.

A function to do this selection in R is:

```
select.2sc <- function(df, n, # number of first stage clusters to select
                        cfraction, clustervar,
                        mmin=2, # minimum number of 2nd stage clusters to select
                        method1="SRSWOR", sizevar1=NULL,
                        method2="SRSWOR", sizevar2=NULL) {
  bigM <- nrow(df)
  if(!clustervar%in%names(df)) {
    stop("Cluster ID variable is not present in the data frame")
  }
  clusterproperties <- as.data.frame(table(df[,clustervar]))
  bigN <- nrow(clusterproperties)
  names(clusterproperties) <- c(clustervar,"Mi")
  if(method1=="SRSWOR") {
    sampcluster <- clusterproperties |>
      select.srswor(n=n)
  } else if(method1=="PPSWR") {
    sampcluster <- clusterproperties |>
```

Territorial Authority, i	N	n	M_i	m_i	$w_i^{(1)}$	$w_{ij}^{(2)}$	Weight, w_{ij}	Selected Schools
Masterton District	87	10	20	4	8.70	5.00	43.50	244;1660;2843;2909
Nelson City	87	10	22	4	8.70	5.50	47.85	293;294;3221;4121
Queenstown-Lakes District	87	10	16	3	8.70	5.33	46.40	746;747;958
Tasman District	87	10	35	7	8.70	5.00	43.50	292;296;3180;3203 ...
Timaru District	87	10	28	6	8.70	4.67	40.60	2109;3498;3528;3533 ...
Waimate District	87	10	9	2	8.70	4.50	39.15	3519;3574
Waitematā Local Board Area	87	10	21	4	8.70	5.25	45.67	48;921;1392;2085
Waitomo District	87	10	17	3	8.70	5.67	49.30	1778;2073;2267
Whanganui District	87	10	40	8	8.70	5.00	43.50	422;1585;2361;2412 ...
Whangarei District	87	10	54	11	8.70	4.91	42.71	13;620;1015;1017 ...

```

        select.ppswr(n=n, sizevar="Mi")
  } else {
    stop("First stage sampling method invalid")
  }
  sampcluster$mi <- pmin(sampcluster$Mi, pmax(mmin, round(cfraction*sampcluster$Mi)))
  samp1 <- merge(df, sampcluster, by=clustervar)
  samp1 <- samp1 |> rename(bigN1=bigN, weight1=weight)
  samp1$n1 <- n

  # Second stage selection
  if(method2=="SRSWOR") {
    samp2 <- do.call(rbind,
                     by(samp1, samp1[,clustervar],
                        function(ds) select.srswor(ds, n=first(ds$mi))
                     ))
  } else if(method2=="PPSWR") {
    samp2 <- do.call(rbind,
                     by(samp2, samp2[,clustervar],
                        function(ds) select.ppswr(ds, n=first(ds$mi), sizevar=sizevar2)
                     ))
  } else {
    stop("Second stage sampling method invalid")
  }
  samp2 <- samp2 |> rename(bigN2=bigN, weight2=weight)
  samp2$n2 <- samp2$mi
  samp2$weight <- samp2$weight1*samp2$weight2

  return(samp2)
}

samp <- schools |> select.2sc(n=10, cfraction=0.20, mmin=2,
                           clustervar="Territorial_Authority")

```

5.9 Complex multistage designs

The ideas in two stage cluster sampling can be extended to a third stage of sampling (i.e. we may wish to sample within SSUs), and we can also combine cluster sampling with stratified sampling.

For example, the Statistics NZ Household sampling frame is a multistage cluster design, where the sample is selected as follows:

1. NZ is stratified into $H = 12$ regions (e.g. Northland, Auckland, Hawke's Bay etc.);
2. Each stratum (region) is broken into N_h PSUs - small areas that contain (on average) about 100 households.
A PPSWR of n_h PSUs is taken within each stratum h ; (First stage sample.)
3. Within the k^{th} selected PSU, which is of size M_{hk} , a linear systematic random sample is taken of m_{hk} households; (Second stage sample.)
4. Depending on the survey, **all** A_{hkl} eligible members of the ℓ^{th} selected household will be surveyed, or a SRSWOR of a_{hkl} members may be taken. (Third stage sample.)

The sample size is therefore

$$\text{Sample Size} = \sum_{h=1}^H \sum_{k \in s_I} \sum_{\ell \in s_k} a_{hkl}$$

and the sample weights of individual j in household ℓ of PSU k in stratum h are

$$w_{hklj} = \frac{N_h}{n_h} \frac{M_{hk}}{m_{hk}} \frac{A_{hkl}}{a_{hkl}}$$

i.e. these are just the product of the weights at each stage of selection.

Function to draw a stratified two stage cluster sample

```
select.st2sc <- function(df,
                          n,          # total number of 1st stage clusters to select over a
                          stratumvar, allocation="equal",
                          nmin=2, # minimum number of 1st stage clusters to select per
                          cfraction, clustervar,
                          mmin=2, # minimum number of 2nd stage clusters to select
                          method1="SRSWOR", sizevar1=NULL,
                          method2="SRSWOR", sizevar2=NULL) {

  # stratum selection
  bigN <- nrow(df)
  if(!stratumvar%in%names(df)) {
    stop("Stratum variable is not present in the data frame")
  }
  if(!clustervar%in%names(df)) {
    stop("Clustering variable is not present in the data frame")
  }
  xx <- tapply(df[,clustervar], df[,stratumvar], function(x) length(unique(x)))
  stratumproperties <- data.frame(Var1=names(xx), Freq=xx)
  names(stratumproperties) <- c(stratumvar, "Nh")
  rownames(stratumproperties) <- NULL
  if(allocation=="equal") {
    stratumproperties$nh <- pmin(stratumproperties$Nh, pmax(nmin,
                                                            round(n/nrow(stratumproperties))))
  } else if(allocation=="proportional") {
    stratumproperties$nh <- pmin(stratumproperties$Nh, pmax(nmin,
                                                            round(n*stratumproperties$Nh/sum(stratumproperties$Nh))))
  } else {
    stop("Invalid allocation")
  }
}
```

Region, h	Territorial Authority, i	N_h	n_h	M_{hi}	m_{hi}	$w_i^{(1)}$	$w_{ij}^{(2)}$	Weight, w_{ij}	
Auckland Region	Aotea/Great Barrier Local Board Area	21	7	3	2	3.00	1.50	4.50	
Auckland Region	Devonport-Takapuna Local Board Area	21	7	24	5	3.00	4.80	14.40	
Auckland Region	Kaipātiki Local Board Area	21	7	25	5	3.00	5.00	15.00	
Auckland Region	Maungakiekie-Tāmaki Local Board Area	21	7	29	6	3.00	4.83	14.50	
Auckland Region	Māngere-Ōtāhuhu Local Board Area	21	7	35	7	3.00	5.00	15.00	
Auckland Region	Puketāpapa Local Board Area	21	7	21	4	3.00	5.25	15.75	
Auckland Region	Waitākere Ranges Local Board Area	21	7	14	3	3.00	4.67	14.00	
Bay of Plenty Region	Whakatane District	6	2	32	6	3.00	5.33	16.00	
Bay of Plenty Region	Ōpōtiki District	6	2	13	3	3.00	4.33	13.00	
Canterbury Region	Ashburton District	9	3	23	5	3.00	4.60	13.80	

```

}
df <- merge(df, stratumproperties, by=stratumvar)

samp <- do.call(rbind, by(df, df[,stratumvar],
  function(dfx) {
    select.2sc(dfx, n=first(dfx$nh),
      cfraction=cfraction,
      clustervar=clustervar,
      mmin=mmin,
      method1=method1, sizevar1=sizevar1,
      method2=method2, sizevar2=sizevar2)
  }))

return(samp)
}

samp <- schools |> select.st2sc(n=30, stratumvar="Region", allocation="proportional",
  cfraction=0.20, clustervar="Territorial_Authority")

```

First 10 of the selected first stage clusters:

```

## `summarise()` has regrouped the output.
## i Summaries were computed grouped by Region and Territorial_Authority.
## i Output is grouped by Region.
## i Use `summarise(.groups = "drop_last")` to silence this message.
## i Use `summarise(.by = c(Region, Territorial_Authority))` for per-operation
## grouping (`?dplyr::dplyr_by`) instead.

```


Chapter 6

Sample Designs in R

The method of analysis of survey data depends crucially on its sample design. Simple methods of analysis which neglect the sample design can lead to estimates which are severely biased, and standard errors that are incorrect.

In this and the following Chapter we'll use the `survey` package in R to carry out analyses. We'll use the `schools` data set as our primary example, but also use data on populations and Gross Domestic Product for each country in 2023 from the 'Our World in Data' website.

```
# Census
schools.census <- schools |> select.srswor(n=nrow(schools))
# SRSWOR
schools.srswor <- schools |> select.srswor(n=20)
# PPSWR
schools.ppswr <- schools |> select.ppswr(n=20, sizevar="Total")
# STSRS
schools.stsrs <- schools |> select.stsrs(n=60, stratumvar="Education_Region")
# 1SC
schools.1sc <- schools |> select.1sc(n=10, clustervar="Territorial_Authority")
# 2SC
schools.2sc <- schools |> select.2sc(n=10, cfraction=0.20, mmin=2,
                                   clustervar="Territorial_Authority")
# ST2SC
schools.st2sc <- schools |> select.st2sc(n=30, stratumvar="Region", allocation="proportional",
                                       cfraction=0.20, clustervar="Territorial_Authority")
```

In order to make estimates from sampled data using the `survey` package we need to specify the sample design we're using. We'll do this for the various sample designs covered in the previous chapter.

Note that in the `survey` package we identify variables with their names without quotes and with a `~` symbol in front. (If you're interested you can use the `survyr` package which is a `dplyr`-like wrapper for `survey` with functions defined so that the `~` isn't necessary.)

We'll define each design, and then use it to estimate the mean size of school rolls in the population (stored in the variable `Total` in the `schools` data frame. (One exception will be the PPSWR example where we'll select a sample of countries using their population size as the size variable.)

The true mean school size is

```
ybar <- mean(schools$Total)
ybar
```

```
## [1] 334.481
```

6.1 Census

The analysis of a census is trivial (at least, if there is no non-response). We can create a census dataset by taking an SRSWOR of size $n = N$.

```
schools.census <- schools |> select.srswor(n=nrow(schools))
```

This function adds the column `bigN` to the dataset, which contains population size N repeated on every row.

Now we can specify the design.

```
schools.census.des <- svydesign(id=~1, data=schools.census, fpc=~bigN)
```

Note that we tell R which column contains the population size using the `fpc` argument. Our data aren't clustered, so the `id` variable (which identifies clusters) is set to 1.

```
svymean(~Total, design=schools.census.des, deff=TRUE)
```

```
##          mean      SE DEff
## Total 334.48    0.00  NaN
```

The estimated mean is 334.48, which is of course exactly equal to the population mean - and note that the standard error is zero: censuses have no sampling error. We've measured everything there is to measure so we have no uncertainty.

We've requested `Deff=TRUE`, and have got a `Deff` value of `NaN`, which means not able to be calculated. We'll discuss what this means later on.

If we had left out the `fpc` command

```
schools.census.des <- svydesign(id=~1, data=schools.census)
```

```
## Warning in svydesign.default(id = ~1, data = schools.census): No weights or
## probabilities supplied, assuming equal probability
```

we get a warning that we can interpret as R assuming an SRSWR from an infinite population.

The estimated mean is the same, but the standard error are now what'd we'd have learned from first year statistics.

```
svymean(~Total, design=schools.census.des, deff=TRUE)
```

```
##          mean      SE DEff
## Total 334.4810    8.4582  Inf
sy <- sd(schools.census$Total)
n <- nrow(schools.census)
se <- sy/sqrt(n)
```

i.e. the standard error is

$$SE[\hat{Y}] = \frac{s_y}{\sqrt{n}} = \frac{423.162}{\sqrt{2503}} = 8.458$$

Also notice that the `Deff` value has changed to `Inf` (meaning infinity) - again we'll come back to what this means this shortly.

6.2 Simple Random Sample WithOut Replacement (SR-SWOR)

Let's draw a sample of $n = 60$ schools from the population by SRSWOR.

```
schools.srswor <- schools |> select.srswor(n=60)
```

To specify the design we use the same syntax as above:

```
schools.srswor.des <- svydesign(id=~1, data=schools.srswor, fpc=~bigN)
```

Estimating the mean school size is also just the same:

```
svymean(~Total, design=schools.srswor.des, deff=TRUE)
```

```
##           mean      SE DEff
## Total 347.367  49.582    1
```

This is the mean of a genuine sample now, so we have an estimated mean value that differs from the truth, and the standard error is gauge of its accuracy.

The standard error comes from the finite population corrected formula:

```
bigN <- first(schools.srswor$bigN) # take the entry in the first row: they're all the same
sy <- sd(schools.srswor$Total)
n <- nrow(schools.srswor)
se <- sqrt(1-n/bigN)*sy/sqrt(n)
```

$$SE[\widehat{Y}] = \sqrt{1 - \frac{n}{N}} \frac{s_y}{\sqrt{n}} = \sqrt{1 - \frac{60}{2503}} \frac{388.745}{\sqrt{60}} = 49.582$$

Compare this to the simple SE, which is just slightly bigger

```
se <- sy/sqrt(n)
```

$$SE[\widehat{Y}] = \frac{s_y}{\sqrt{n}} = \frac{388.745}{\sqrt{60}} = 50.187$$

This is what we get if we omit the `fpc`

```
schools.srswor.des <- svydesign(id=~1, data=schools.srswor)
```

```
## Warning in svydesign.default(id = ~1, data = schools.srswor): No weights or
## probabilities supplied, assuming equal probability
```

```
svymean(~Total, design=schools.srswor.des, deff=TRUE)
```

```
##           mean      SE DEff
## Total 347.367  50.187  Inf
```

Country	Code	GovtSpend	GDP	Region	Population	GDPPerCapita
Albania	ALB	1863822443	1.496298e+10	Europe	2811660	5321.759
Algeria	DZA	44321604142	2.151440e+11	Africa	46164222	4660.405
Angola	AGO	15741572839	1.032363e+11	Africa	36749909	2809.157
Argentina	ARG	111098899281	5.916847e+11	South America	45538402	12993.093
Armenia	ARM	1598415044	1.540616e+10	Asia	2943398	5234.142
Australia	AUS	385802798739	1.649425e+12	Oceania	26451126	62357.456
Austria	AUT	84758230499	4.246078e+11	Europe	9130434	46504.669
Azerbaijan	AZE	NA	5.755883e+10	Asia	10318209	5578.374
Bahrain	BHR	6908975351	3.996681e+10	Asia	1569674	25461.855
Bangladesh	BGD	18377804014	3.232800e+11	Asia	171466986	1885.377

6.3 Probability Proportional to Size With Replacement (PPSWR)

We have a dataset from ‘Our World in Data’ giving population sizes, Gross Domestic Product and military expenditure in 2023 (all in US dollars). The first 10 rows are below:

We can select a PPSWR using population as the size variable

```
mil.ppswr <- mil |> select.ppswr(n=30, sizevar="Population")
```

The **weight** column here combines the weight over the number of times a unit is selected

```
mil.ppswr.des <- svydesign(id=~1, data=mil.ppswr, weight=~weight)
```

The mean military expenditure per country is

```
svymean(~MilSpend, design=mil.ppswr.des, deff=TRUE)
```

```
##              mean          SE   DEff
## MilSpend 7.4657e+10 3.4650e+10 1.0362
```

though more interesting is the **total** of this variable, to see how much the world is spending.

```
svytotal(~MilSpend, design=mil.ppswr.des, deff=TRUE)
```

```
##              total          SE   DEff
## MilSpend 1.1199e+13 5.3452e+12 1.096
```

The estimation of the **total** of a variable is only something that makes sense in a finite population, and is a type of estimate that is rarely found anywhere outside sample survey analysis.

It is frequently of interest in Official (i.e. government) statistics that we want to estimate the total number of people with a certain characteristic, or the total amount of a commodity being exported, or total export earnings.

Note again that in a with replacement design we can’t apply the finite population correction.

6.4 Stratified Sample

Recall that in this design we specify a set of groups called strata, and each population unit belongs to one of these strata.

Selecting a sample of $n = 60$ schools by equal allocation across the 12 Education Regions:

```
schools.stsrs <- schools |> select.stsrs(n=60, stratumvar="Education_Region")
```

Specify this design

```
schools.stsrs.des <- svydesign(id=~1, strata=~Education_Region, fpc=~Nh,
                             data=schools.stsrs,
                             survey.lonely.psu="remove")
```

Note that we've included the option `survey.lonely.psu="remove"` to mean that the estimation of standard errors doesn't fail if a stratum only has a single PSU.

We can now return to estimating the mean school size

```
svymean(~Total, design=schools.stsrs.des, deff=TRUE)
```

```
##           mean      SE   DEff
## Total 337.235  39.972 0.6816
```

Now it's finally time to talk about the **Design Effect**, or the Deff for short.

The Deff of an estimator $\hat{T}$ of a population characteristic T in a particular sample design is defined to be

$$\text{Deff}[\hat{T}_{\text{design}}] = \frac{\text{Var}[\hat{T}_{\text{design}}]}{\text{Var}[\hat{T}_{\text{SRSWOR}}]} = \left(\frac{\text{SE}[\hat{T}_{\text{design}}]}{\text{SE}[\hat{T}_{\text{SRSWOR}}]} \right)^2$$

where both estimators are based on samples of the same size.

It tells us how the standard error differs between our design of interest and a SRSWOR of the same sample size.

Note that

- We have a different Deff for every estimate we make from a survey.
- A $\text{Deff} < 1$ means that the Standard Errors are smaller than what we'd get in SR-SWOR. A small Deff indicates a **more efficient** design.
- A Deff of 1 means the design is no better than SRSWOR - and of course (look back above) the Deffs of all estimates in SRSWOR are 1 by definition.
- If we're dealing with a census, then the Standard Error is zero, and in any case there's no difference between designs, so the Deff isn't defined.
- **Stratified designs** can achieve Deffs much much less than 1 if their strata are well defined to be informative about the estimates of interest. e.g. in business surveys where the strata are defined by business size we can get very good Deffs.
- **Cluster designs** often have Deffs greater than 1, which show them to be significantly less efficient than SRSWOR.

The Deff values we get from R are approximations, and aren't very reliable when sample sizes are small. But do notice that we have a $\text{Deff} < 1$ here for the stratified sample of $n = 60$ units, and the SE is lower than that achieved in the SRSWOR of $n = 60$ units above.

6.5 One stage cluster design (1SC)

A one stage cluster design takes $n = 10$ Territorial Authority areas (TAs) at random and selects all schools in each selected TA.

```
schools.1sc <- schools |> select.1sc(n=10, clustervar="Territorial_Authority")
```

We need a column `bigN` which gives the number of TAs in the population.

```
schools.1sc.des <- svydesign(id=~Territorial_Authority, data=schools.1sc,
                           fpc=~bigN)
```

```
svymean(~Total, design=schools.1sc.des, deff=TRUE)
```

```
##          mean      SE  DEff
## Total 287.197  43.829 5.7634
```

The Deff of this one stage cluster sample is bad (much larger than 1) - this is due to strong differences between TAs in the mean school size. The more the clusters differ from each other on the characteristic of interest, the worse the Deff gets.

6.6 Two stage cluster design (2SC)

A two stage cluster design takes $n = 10$ Territorial Authority areas (TAs) at random and selects 20% of the schools in each selected TA.

```
schools.2sc <- schools |> select.2sc(n=10, cfraction=0.20, mmin=2,
                                   clustervar="Territorial_Authority")
```

We need a column `bigN1` which gives the number of TAs in the population, and `bigN2` which is the number of schools in each TA. We also need to specify names of unique identifiers of the first stage units (the PSUs: these are Territorial Authorities) and the second stage units (the SSUs: these are Schools).

```
schools.2sc.des <- svydesign(id=~Territorial_Authority+School_Id, data=schools.2sc,
                           fpc=~bigN1+bigN2,
                           survey.lonely.psu="remove")
```

```
svymean(~Total, design=schools.2sc.des, deff=TRUE)
```

```
##          mean      SE  DEff
## Total 486.981  84.261 2.1739
```

6.7 Stratified two stage cluster design

Finally here is a stratified two stage sample of schools within TA's with Regions. The Regions are strata, the TAs are clusters within the Regions, and then 20% of schools are sampled in the selected TAs.

```
schools.st2sc <- schools |> select.st2sc(n=30, stratumvar="Region", allocation="proportional",
                                       cfraction=0.20, clustervar="Territorial_Authority")
```

The design is the same as the two stage cluster sample, except we add the stratum specification.

```
schools.st2sc.des <- svydesign(id=~Territorial_Authority+School_Id, strata=~Region, data=schools.st2sc,
                              fpc=~bigN1+bigN2,
                              survey.lonely.psu="remove")
```

```
svymean(~Total, design=schools.st2sc.des, deff=TRUE)
```

```
##          mean      SE  DEff
## Total 295.326  25.898 1.248
```

Chapter 7

Survey Analysis

With the ability to load a survey dataset and specify its survey design we're now ready to do some analysis.

We'll use the `api` dataset which comes with the `survey` package: it's a set of observations of the Academic Performance Index (API) on all California Schools with at least 100 students in the years 1999 and 2000.

```
data(api)
```

In particular we'll use the `apiclus2` dataset, which is a two stage cluster sample of 126 schools drawn from the population of 6194. The top 10 rows of the data set and some of the key variables are shown below.

```
apiclus2$type <- factor(apiclus2$type, levels=c("E","M","H"))
levels(apiclus2$type) <- c("Elementary","Middle","High")

apiclus2 |>
  slice_head(n=10) |>
  select(cds, stype, name, dnum, cnum, api99, api00) |>
  kbl() |>
  kable_styling()
```

The full set of variables in the dataset is:

The `apiclus2` dataset was drawn by taking a sample of districts (labelled by district number

cds	stype	name	dnum	cnum	api99	api00
31667796031017	Elementary	Alta-Dutch Flat	15	30	785	821
55751846054837	Elementary	Tenaya Elementa	63	54	718	773
41688746043517	Elementary	Panorama Elemen	83	40	632	600
41688746043509	Middle	Lipman Middle	83	40	740	740
41688746043491	Elementary	Brisbane Elemen	83	40	711	716
40687266042998	Elementary	Cayucos Element	117	39	779	811
20651936023881	Elementary	Fuller (Merle L	132	19	432	472
20651936023931	Elementary	Fairmead Elemen	132	19	494	520
20651936023907	Middle	Wilson Elementa	132	19	589	568
06615980631259	High	Colusa High	152	5	585	591

Variable	Description
cds	Unique identifier
stype	Elementary/Middle/High School
name	School name (15 characters)
sname	School name (40 characters)
snum	School number
dname	District name
dnum	District number
cname	County name
cnum	County number
flag	reason for missing data
pcttest	percentage of students tested
api00	API in 2000
api99	API in 1999
target	target for change in API
growth	Change in API
sch.wide	Met school-wide growth target?
comp.imp	Met Comparable Improvement target
both	Met both targets
awards	Eligible for awards program
meals	Percentage of students eligible for subsidized meals
ell	‘English Language Learners’ (percent)
yr.rnd	Year-round school
mobility	percentage of students for whom this is the first year at the school
acs.k3	average class size years K-3
acs.46	average class size years 4-6
acs.core	Number of core academic courses
pct.resp	percent where parental education level is known
not.hsg	percent parents not high-school graduates
hsg	percent parents who are high-school graduates
some.col	percent parents with some college
col.grad	percent parents with college degree
grad.sch	percent parents with postgraduate education
avg.ed	average parental education level
full	percent fully qualified teachers
emer	percent teachers with emergency qualifications
enroll	number of students enrolled
api.stu	number of students tested.

dnum), and then a sample of schools within districts (schools are labelled by their school number snum).

The column fpc1 contains the number of districts (757), and the column fpc2 contains the number of schools in each district (which varies by district).

We can specify this two stage cluster design using

```
apiclus2.des <- svydesign(ids=~dnum+snum, fpc=~fpc1+fpc2, data=apiclus2)
```

The summary() function tells us the basics of the design we've just specified:

```
summary(apiclus2.des)

## 2 - level Cluster Sampling design
## With (40, 126) clusters.
## svydesign(ids = ~dnum + snum, fpc = ~fpc1 + fpc2, data = apiclus2)
## Probabilities:
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
## 0.003669 0.037743 0.052840 0.042390 0.052840 0.052840
## Population size (PSUs): 757
## Data variables:
## [1] "cds"      "stype"    "name"     "sname"    "snum"     "dname"
## [7] "dnum"     "cname"    "cnum"     "flag"     "pcttest"  "api00"
## [13] "api99"    "target"   "growth"   "sch.wide" "comp.imp" "both"
## [19] "awards"   "meals"    "ell"      "yr.rnd"   "mobility" "acs.k3"
## [25] "acs.46"   "acs.core" "pct.resp" "not.hsg"  "hsg"      "some.col"
## [31] "col.grad" "grad.sch" "avg.ed"   "full"     "emer"     "enroll"
## [37] "api.stu"  "pw"       "fpc1"     "fpc2"
```

We can see that from the summary that $n = 40$ districts were sampled (out of $N = 757$). Let's count the number of schools sampled per district. We'll just display the top 15 largest districts:

```
apiclus2 |>
  group_by(dnum,dname) |>
  summarise(Mi=first(fpc2),mi=n()) |>
  ungroup() |>
  arrange(desc(Mi)) |>
  slice_head(n=15) |>
  kbl(col.names=c("dnum","District Name","$M_i$","$m_i$"),escape=FALSE) |>
  kable_styling(full_width=FALSE)
```

After the 10 largest districts a census of schools is taken in each selected district $m_i = M_i$

7.1 Continuous variables

Consider the key outcome variable 'api00', the API score in the year 2000 for each school.

Display

The basic workhorse for the display of numerical variables is the histogram

```
svyhist(~api00, design=apiclus2.des, prob=FALSE,
        xlab="API Score (Year 2000)",
        main="Weighted distribution of API (2000) scores")
```

dnum	District Name	M_i	m_i
620	Sacramento City Unified	72	5
570	Pomona Unified	36	5
575	Poway Unified	28	5
638	San Lorenzo Unified	14	5
200	El Centro Elementary	11	5
731	Sylvan Union Elementary	9	5
596	Rincon Valley Union Elem	7	5
639	San Lorenzo Valley Unif	7	5
781	Walnut Creek Elementary	6	5
295	Healdsburg Unified	5	5
403	Los Gatos Union Elem	5	5
480	Natomas Unified	5	5
534	Palo Verde Unified	5	5
549	Piedmont City Unified	5	5
173	Del Mar Union Elementary	4	4

Weighted distribution of API (2000) scores




Figure 7.1: Weighted histogram of API (2000) scores

Note that this histogram differs strongly from the simple histogram of the sample values:

```
hist(apiclus2$api00,
     xlab="API Score (Year 2000)",
     main="Unweighted distribution of API (2000) scores")
```




Figure 7.2: Unweighted histogram of API (2000) sample data

a warning to us that with survey data we can go badly astray if we ignore the sample design.

Boxplots are an alternative to histograms:

```
svyboxplot(api00~1, design=apiclus2.des,
           ylab="API (2000)", main="API (2000) Scores")
```

... especially if we want to display variables divided by a category:

```
svyboxplot(api00~stype, design=apiclus2.des,
           xlab="School Type",
           ylab="API (2000)", main="API (2000) Scores")
```

If we have two continuous variables we can plot them as a scatterplot, showing the weights of each observation:

```
svyplot(api00~api99, design=apiclus2.des, style="bubble",
        xlab="API (1999)", ylab="API (2000)")
```




Figure 7.3: Distribution of API (2000) sample data




Figure 7.4: Distribution of API (2000) sample data




Figure 7.5: Distribution of API scores in 1999 and 2000

Estimation

We've already seen that we can estimate means using the `svymean()` function.

```
svymean(~api00, design=apiclus2.des)
```

```
##           mean      SE
## api00 670.81 30.099
```

The `jtools` package allows us to estimate the population standard deviation too

```
library(jtools)
svysd(~api00, design=apiclus2.des)
```

```
##           std. dev.
## api00      136.83
```

(Pause a minute and remind yourself of the difference between a standard deviation and a standard error.)

We can estimate means for a set of separate variables, using a perhaps unintuitive notation:

```
svymean(~api99+api00, design=apiclus2.des)
```

```
##           mean      SE
## api99 645.03 29.711
## api00 670.81 30.099
```

or for a variable split by categories by wrapping `svymean()` up inside a call to `svyby()`

```
svyby(~api00, by=~stype, svymean, design=apiclus2.des)
```

```
##           stype      api00      se
## Elementary Elementary 692.8104 29.92660
## Middle      Middle 642.3520 45.09132
## High        High 598.3407 17.69417
```

A **dotchart** is a nice way of showing the estimated means and their confidence intervals:

```
ss <- svyby(~api00, by=~stype, svymean, design=apiclus2.des)
sc <- confint(ss)
ns <- nrow(sc)
dotchart(ss, xlim=range(sc), xlab="Mean API (2000) Score")
arrows(sc[,1], 1:ns, sc[,2], 1:ns, angle=90, code=3, length=0.1)
```

If we want to test statistical hypotheses about differences between groups one way is to use the regression framework

```
summary(svyglm(api00~stype, design=apiclus2.des))
```

```
##
## Call:
## svyglm(formula = api00 ~ stype, design = apiclus2.des)
##
## Survey design:
## svydesign(ids = ~dnum + snum, fpc = ~fpc1 + fpc2, data = apiclus2)
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)   692.81      29.93  23.150 < 2e-16 ***
```




Figure 7.6: Mean API (2000) scores by School Type

```
## stypeMiddle   -50.46      25.30  -1.994  0.05354 .
## stypeHigh     -94.47      28.21  -3.348  0.00188 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 17528.44)
##
## Number of Fisher Scoring iterations: 2
```

Here we can see that with Elementary as the reference level we get the same estimate of the mean API reported as the (`Intercept`). We further see that there is no significant difference between Middle and Elementary Schools in mean API, but the mean API is significantly lower at High Schools compared to Elementary Schools.

Regressions take of course take lots of variables into account, and we can apply familiar regression techniques to investigating conditional associations between predictors and a continuous outcome.

For example we can test whether the percentage of children at a school who have free school meals (i.e. the `meals` variable) associates with the API score:

```
summary(svyglm(api00~stype+meals, design=apiclus2.des))

##
## Call:
## svyglm(formula = api00 ~ stype + meals, design = apiclus2.des)
##
## Survey design:
## svydesign(ids = ~dnum + snum, fpc = ~fpc1 + fpc2, data = apiclus2)
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  874.8638    16.4677   53.126 < 2e-16 ***
## stypeMiddle  -60.9517    20.6964   -2.945  0.00563 **
## stypeHigh    -169.2128    20.2618  -8.351 6.06e-10 ***
## meals         -3.2367     0.2888 -11.209 2.70e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 5560.106)
##
## Number of Fisher Scoring iterations: 2
```

... and there is indeed a strong negative association: the higher the `meals` value the lower the API score.

We can estimate **totals** using the `svytotal()` function - for example we can estimate the total number of students enrolled at all schools in California:

```
svytotal(~enroll, design=apiclus2.des, na.rm=TRUE)
```

```
##           total      SE
## enroll 2639273 799638
```

Although note that we have had to remove some schools with unknown `enroll` values by stating `na.rm=TRUE`, and these forced removals likely mean that our estimate of the total is an underestimate.

Confidence intervals can be calculated using the `confint()` function

```
confint(svymean(~api99+api00, design=apiclus2.des), level=0.95)
```

```
##           2.5 %    97.5 %
## api99 586.8009 703.2670
## api00 611.8188 729.8048
```

```
confint(svyby(~api00, by=~stype, svymean, design=apiclus2.des))
```

```
##           2.5 %    97.5 %
## Elementary 634.1553 751.4655
## Middle     553.9746 730.7294
## High       563.6607 633.0206
```

The standard errors can be directly extracted using `SE()`

```
SE(svyby(~api00, by=~stype, svymean, design=apiclus2.des))
```

```
## [1] 29.92660 45.09132 17.69417
```

The **relative standard errors**, also known as **relative sampling errors** (RSEs) can be extracted by `cv()`:

```
cv(svyby(~api00, by=~stype, svymean, design=apiclus2.des))
```

```
## Elementary      Middle      High
## 0.04319595 0.07019721 0.02957206
```

RSEs are a useful measure of quality of estimates of quantities that are strictly positive, such as the API score.

If we learned that the SE of an estimate (e.g. the mean of `api00`) was 30.1 we wouldn't be able to assess whether it was a good estimate or not without knowing that the value of the estimate was 670.81. A good way to assess the quality of an estimate is the SE of an estimate $\hat{T}$ as a proportion of the estimate $\hat{T}$:

$$\mathbf{RSE}[\hat{T}] = \frac{\mathbf{SE}[\hat{T}]}{\hat{T}}$$

For overall mean `api00` the RSE is

$$\mathbf{RSE}[\hat{Y}] = \frac{30.1}{670.81} = 0.045$$

This is a highly precise estimate: the SE is only 4.5% of the value of the estimate.

Good RSEs are only a few percent at most. Poor RSEs are larger than 10%, and anything bigger than 20% isn't really publishable.

In general a standard deviation divided by a mean goes by the term **coefficient of variation** (CV), which is why the R function is called `cv()`.

```
cv(svymean(~api00, design=apiclus2.des))
```

```
##           api00
## api00 0.04486956
```

7.2 Categorical Variables

The workhorse display for categorical variables is the barplot, showing **proportions** in categories. Proportions are first calculated using the `svymean()` function:

```
barplot(svymean(~stype, design=apiclus2.des),
        names=levels(apiclus2$stype),
        xlab="School Type", ylab="Proportion")
```




Figure 7.7: Estimates of the proportions of Schools of each type

```
svymean(~stype, design=apiclus2.des)
```

```
##               mean      SE
## stypeElementary 0.68118 0.0556
## stypeMiddle     0.18450 0.0222
## stypeHigh       0.13432 0.0567
```

It's worth pausing to recognise that a proportion is a mean of a special type of variable called an **indicator variable**: it takes the value 1 if the observation is a member of specific category and 0 if not. Add up all the ones and zeroes and you get the number of observations in the category, and divide by the number of observations you get a proportion.

$$p = \frac{1 + 0 + 1 + 1 + 0 + 0 + \dots}{n} = \frac{1}{n} \sum_{i=1}^n Y_i$$

where Y_i is an indicator variable.

Table 7.1: (#tab:tabci.symmetric)Symmetric confidence interval estimates

	Estimate	2.5%	97.5%
stypeElementary	0.681	0.572	0.790
stypeMiddle	0.185	0.141	0.228
stypeHigh	0.134	0.023	0.246

With sample data we calculate **weighted sums** though:

$$\hat{p} = \frac{\sum_{i=1}^n w_i Y_i}{\sum_{i=1}^n w_i} = \frac{\hat{Y}}{\hat{N}}$$

where the top line of this fraction is $\hat{Y}$, our estimate of the number of individuals (in a finite population) who belong to the category of interest. The bottom line is $\hat{N}$, the estimated size of the population. $\hat{N}$ is exact in some designs (e.g. SRSWOR, or Stratified SRSWOR), but is approximate in other designs (especially cluster designs).

These total estimates within categories can be estimated using the `svytotal()` function.

```
svytotal(~stype, design=apiclus2.des)
```

```
##              total      SE
## stypeElementary 3493.56 1119.75
## stypeMiddle     946.25  311.81
## stypeHigh       688.87  289.35
```

When creating barplots of estimated proportions and/or totals we really should indicate the uncertainties in our estimates, and show confidence intervals on the bars. We can use the standard errors provided by `svymean()` and `svytotal()`: the function `confint()` will provide these:

```
confint(svytotal(~stype, design=apiclus2.des))
```

```
##              2.5 %   97.5 %
## stypeElementary 1298.8892 5688.221
## stypeMiddle     335.1077 1557.392
## stypeHigh       121.7472 1255.993
```

These confidence intervals are symmetric about the point estimates:

```
ss <- svymean(~stype, design=apiclus2.des)
ci <- confint(ss)
bp <- barplot(ss, xlab="School Type", ylab="Estimated Size",
              names=levels(apiclus2$stype),
              ylim=c(0,max(ci)))
arrows(bp, ci[,1], bp, ci[,2], angle=90, length=0.2, code=3)

names(ss) <- levels(apiclus2$stype)
cbind(as.numeric(ss),ci) |>
  kbl(col.names=c("Estimate","2.5%","97.5%"),
      digits=3,
      caption="Symmetric confidence interval estimates") |>
  kable_styling(full_width=FALSE)
```

These symmetric intervals are actually somewhat unrealistic, because `confint()` doesn't




Figure 7.8: Estimated proportions of schools by school type - with symmetric confidence intervals

Table 7.2: (`#tab:tabci.asymmetric`) Asymmetric confidence interval estimates

	Estimate	2.5%	97.5%
Elementary	0.681	0.560	0.782
Middle	0.185	0.144	0.234
High	0.134	0.055	0.294

know it’s dealing with indicator variables, and doesn’t take into account the fact that proportions and totals can’t go negative.

A better method of constructing confidence intervals is the special function `svyciprop()`:

```
svyciprop(~I(stype=="Elementary"), design=apiclus2.des)
```

```
##                                2.5% 97.5%
## I(stype == "Elementary") 0.681 0.560 0.782
```

Note that we’ve had to create an indicator variable here using the `I()` function for just one level in order to calculate its confidence interval.

To get a confidence interval proportion for every level of a factor it’s somewhat fiddly with the raw `survey` package:

```
ss <- svymean(~stype, design=apiclus2.des)
ci <- t(sapply(levels(apiclus2$stype),
  function(stypelevel) {
    form <- eval(parse(text=paste0("formula(~I(stype==\"", stypelevel, "\"))")))
    confint(svyciprop(form, design=apiclus2.des))
  }))
bp <- barplot(ss, xlab="School Type", ylab="Estimated Size",
  names=levels(apiclus2$stype),
  ylim=c(0,max(ci)))
arrows(bp, ci[,1], bp, ci[,2], angle=90, length=0.2, code=3)

names(ss) <- levels(apiclus2$stype)
cbind(as.numeric(ss),ci) |>
  kbl(col.names=c("Estimate","2.5%","97.5%"),
    digits=3,
    caption="Asymmetric confidence interval estimates") |>
  kable_styling(full_width=FALSE)
```

The biggest thing to notice is that the confidence interval for “High” is now asymmetric: it’s longer on the high side than the low side, and more correctly represents our uncertainty. The estimates for “Elementary” and “Middle” are essentially unchanged.

For **crosstabulations** we can use `svytable()`, which gives weighted counts in categories - e.g. of whether School Wide growth target had been met

```
svytable(~stype+sch.wide, design=apiclus2.des)
```

```
##          sch.wide
## stype          No      Yes
## Elementary 242.240 3251.315
## Middle    446.630  499.620
## High      586.675  102.195
```




Figure 7.9: Estimated proportions of schools by school type - with asymmetric confidence intervals

Which can be displayed with a side-by-side barplot

```
st <- svytable(~sch.wide+stype, design=apiclus2.des)
bp <- barplot(st, beside=TRUE, col=c("light blue","dark red"),
              legend=TRUE, args.legend=list(title="Met targets"),
              xlab="School Type", ylab="Frequency")
```




Figure 7.10: Counts of schools meeting targets, by school type

A test of association is given by the weighted Chi-Square test

```
svychisq(~stype+sch.wide, design=apiclus2.des)
```

```
##
## Pearson's X^2: Rao & Scott adjustment
##
## data: svychisq(~stype + sch.wide, design = apiclus2.des)
## F = 11.873, ndf = 1.3825, ddf = 53.9175, p-value = 0.0003221
```

We can carry out a logistic regression to make a similar test of association. We have `sch.wide` as the binary outcome variable and `stype` as the predictor:

```
summary(svyglm(I(sch.wide=="Yes")~stype, family=quasibinomial(link="logit"),
               design=apiclus2.des))
```

```
##
## Call:
```

```
## svyglm(formula = I(sch.wide == "Yes") ~ stype, design = apiclus2.des,
##       family = quasibinomial(link = "logit"))
##
## Survey design:
## svydesign(ids = ~dnum + snum, fpc = ~fpc1 + fpc2, data = apiclus2)
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.5969     0.5997   4.331 0.000109 ***
## stypeMiddle  -2.4848     1.1565  -2.148 0.038292 *
## stypeHigh    -4.3445     0.8757  -4.961 1.59e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for quasibinomial family taken to be 1.008)
##
## Number of Fisher Scoring iterations: 5
```

With regular `lm()` and `glm()` regression fits we can use F tests to determine whether specific terms in the model should be included. A similar test can be done to test for the inclusion of specific terms using `regTermTest`

```
sr <- svyglm(I(sch.wide=="Yes")~stype, family=quasibinomial(link="logit"),
            design=apiclus2.des)
regTermTest(sr, ~stype)
```

```
## Wald test for stype
## in svyglm(formula = I(sch.wide == "Yes") ~ stype, design = apiclus2.des,
##       family = quasibinomial(link = "logit"))
## F = 13.61363 on 2 and 37 df: p= 3.7061e-05
```

Here we can see that `stype` is contributing significantly to the fit of the model.

Chapter 8

Weight adjustments

At the end of a survey we may find ourselves in the undesirable situation that

- Purely by chance, the sample is somewhat unbalanced: e.g. by chance we may have selected many more males than females, or more large households than small ones, or
- We may have had **unit nonresponse** - i.e. there may be selected units from which we have collected no data. In the case of human subjects this may have been from a failure to contact the selected person, or the refusal of the person to answer our questions.

In both cases we are in the situation that our sample, even when weighted, is not a good representation of the population from which it is drawn.

For example consider the SURF dataset, which is a synthetic sample of 200 individuals between the ages of

8.1 Calibration

```
apipopc <- apipop[!is.na(apipop$enroll),]  
apipopc$type <- factor(as.character(apipopc$type), levels=c("E", "M", "H"))  
levels(apipopc$type) <- c("Elementary", "Middle", "High")  
apipopc$yr.rnd[is.na(apipopc$yr.rnd)] <- "Yes"
```

```
mean(apipopc$enroll)
```

```
## [1] 619.0469
```

```
tapply(apipopc$enroll, apipopc$type, mean)
```

```
## Elementary      Middle      High  
##   426.9616   912.0892  1349.9654
```

```
table(apipopc$type)
```

```
##  
## Elementary      Middle      High  
##      4397      1009      751
```

```
table(apipopc$type)/nrow(apipopc)
```

```
##
```

```
## Elementary      Middle      High
## 0.7141465 0.1638785 0.1219750
```

```
stype.benchmarks <- as.data.frame(table(apipopc$stype))
names(stype.benchmarks) <- c("stype", "Freq")
stype.benchmarks
```

```
##          stype Freq
## 1 Elementary 4397
## 2      Middle 1009
## 3          High  751
```

```
sum(stype.benchmarks$Freq) # 6157
```

```
## [1] 6157
```

```
yr.rnd.benchmarks <- as.data.frame(table(apipopc$yr.rnd))
names(yr.rnd.benchmarks) <- c("yr.rnd", "Freq")
yr.rnd.benchmarks
```

```
## yr.rnd Freq
## 1      No  874
## 2      Yes 5283
```

```
sum(yr.rnd.benchmarks$Freq) # 6157
```

```
## [1] 6157
```

```
stype.yr.rnd.benchmarks <- as.data.frame(table(apipopc$stype, apipopc$yr.rnd))
names(stype.yr.rnd.benchmarks) <- c("stype", "yr.rnd", "Freq")
stype.yr.rnd.benchmarks
```

```
##          stype yr.rnd Freq
## 1 Elementary      No   791
## 2      Middle      No    60
## 3          High      No    23
## 4 Elementary      Yes 3606
## 5      Middle      Yes   949
## 6          High      Yes   728
```

Select a two stage cluster sample - make estimates

```
set.seed(111)
api2sc <- apipopc |>
  select.2sc(n=40, cfraction=0.30, mmin=2, clustervar="dnum")
api2sc.des <- svydesign(ids=~dnum+snum, fpc=~bigN1+bigN2, data=api2sc)
table(api2sc$stype)/nrow(api2sc)
```

```
##
## Elementary      Middle      High
## 0.6224490 0.1428571 0.2346939
```

```
svytable(~stype, design=api2sc.des)
```

```
## stype
## Elementary      Middle      High
## 2972.2159 589.3841 946.0500
```

```
svytable(~stype, design=api2sc.des)/sum(api2sc$weight)
```

```
## stype
## Elementary      Middle      High
## 0.6593715 0.1307520 0.2098765
```

```
svymean(~enroll, design=api2sc.des)
```

```
##          mean      SE
## enroll 711.45 81.057
```

```
svyby(~enroll, by=~stype, api2sc.des, svymean)
```

```
##              stype      enroll      se
## Elementary Elementary 415.6678 46.44827
## Middle      Middle 834.1254 113.51448
## High              High 1564.3007 190.86274
```

Poststratify by stype and reestimate

```
api2sc.des.ps <- postStratify(api2sc.des, strata=~stype, population=stype.benchmarks)
svymean(~enroll, api2sc.des.ps)
```

```
##          mean      SE
## enroll 624.35 53.142
```

Postratify by stype x yr.rnd and reestimate

```
api2sc.des.ps2 <- postStratify(api2sc.des, strata=~stype+yr.rnd, population=stype.yr.rnd.benchmarks)
svymean(~enroll, api2sc.des.ps2)
```

```
##          mean      SE
## enroll 629.45 40.239
```

Poststratify by stype + yr.rnd (i.e. we only have the marginals) and reestimate

```
api2sc.des.rake <- rake(api2sc.des,
                        sample=list(~stype, ~yr.rnd),
                        population=list(stype.benchmarks, yr.rnd.benchmarks))
svymean(~enroll, api2sc.des.rake)
```

```
##          mean      SE
## enroll 632.04 46.705
```

```
ff <- function(x) c(as.numeric(x), SE(x))
```

Compare

```
ss <- rbind(
  c(mean(apipopc$enroll), 0),
  ff(svymean(~enroll, api2sc.des)),
  ff(svymean(~enroll, api2sc.des.ps)),
  ff(svymean(~enroll, api2sc.des.ps2)),
  ff(svymean(~enroll, api2sc.des.rake)))
dimnames(ss) <- list(c("Census", "2SC", "2SC+PS(stype)", "2SC+PS(stype x yr.rnd)", "2SC+Rake(stype+yr.rnd)"),
                     c("Estimate", "SE"))
ss
```

```
##              Estimate      SE
```

```
## Cenus                619.0469  0.00000
## 2SC                  711.4530  81.05714
## 2SC+PS(stype)        624.3485  53.14154
## 2SC+PS(stype x yr.rnd) 629.4503  40.23927
## 2SC+Rake(stype+yr.rnd) 632.0372  46.70502
```

With replacement: specify weights only - same as 2SC estimate, but it gets SE wrong

```
api2sc.des.wr <- svydesign(ids=~1, data=api2sc, weight=~weight)
svytable(~stype, design=api2sc.des.wr)
```

```
## stype
## Elementary      Middle      High
## 2972.2159    589.3841    946.0500
```

```
svytable(~stype, design=api2sc.des.wr)/sum(api2sc$weight)
```

```
## stype
## Elementary      Middle      High
## 0.6593715    0.1307520    0.2098765
```

```
svymean(~enroll, design=api2sc.des.wr)
```

```
##          mean      SE
## enroll 711.45 70.509
```

```
svyby(~enroll, ~stype, design=api2sc.des.wr, svymean)
```

```
##          stype      enroll      se
## Elementary Elementary 415.6678 25.89762
## Middle      Middle 834.1254 113.93267
## High          High 1564.3007 171.61203
```

8.2 Replicate weights

So we need replicate weights - generate these

In such circumstances we often create multiple sets of weights in addition to the adjusted weights. These **replicate weights** are what are used to estimate the standard errors, and they have the advantage that they can create standard errors for a wider range of estimators than the usual means and totals.

Jackknife weights are calculated by systematically deleting PSUs at the first stage of selection, and then scaling the weights on the remaining rows so that they still add up to the same total.

```
api2sc.jkdes <- as.svrepdesign(api2sc.des, type="JK1")
```

```
## Warning in as.svrepdesign.default(api2sc.des, type = "JK1"): Finite population
## corrections after first stage have been dropped
```

Informative warning here: no fpc can apply at second stage or deeper.

```
names(api2sc.jkdes)
```

```
## [1] "repweights"      "pweights"        "type"            "rho"
## [5] "scale"           "rscales"         "call"            "combined.weights"
## [9] "selfrep"         "mse"             "variables"       "degf"
```

```
api2sc.jkdes$repweights$weights[1:3,] # an n x n matrix of weights: each column is a replicate, each
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]      [,8]
## [1,] 0.000000 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [2,] 1.025641 0.000000 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [3,] 1.025641 1.025641 0.000000 1.025641 1.025641 1.025641 1.025641 1.025641
##           [,9]      [,10]     [,11]     [,12]     [,13]     [,14]     [,15]     [,16]
## [1,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [2,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [3,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
##           [,17]     [,18]     [,19]     [,20]     [,21]     [,22]     [,23]     [,24]
## [1,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [2,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [3,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
##           [,25]     [,26]     [,27]     [,28]     [,29]     [,30]     [,31]     [,32]
## [1,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [2,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [3,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
##           [,33]     [,34]     [,35]     [,36]     [,37]     [,38]     [,39]     [,40]
## [1,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [2,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
## [3,] 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641 1.025641
```

Replicate Weights are $n/(n-1) = 40/39 = 1.025641$ - drops one `dnum` out for each replicate. These need to be multiplied by the original sampling weight to get the actual weights.

Use JK1 for unstratified designs, JK_n for stratified designs

```
svymean(~enroll, design=api2sc.jkdes)
```

```
##           mean      SE
## enroll 711.45 82.848
```

Better estimate of the SE.

We can poststratify these weights by the way: we poststratify each replicate weight set in addition to the probability weights themselves.

```
api2sc.jkdes.ps <- postStratify(api2sc.jkdes, strata=~stype, population=stype.benchmarks)
svymean(~enroll, design=api2sc.jkdes.ps)
```

```
##           mean      SE
## enroll 624.35 60.61
```

We get immediate reduction of the standard errors

Comparison

```
ss <- rbind(
  c(mean(api2sc.jkdes$enroll), 0),
  ff(svymean(~enroll, api2sc.des)),
  ff(svymean(~enroll, api2sc.des.ps)),
  ff(svymean(~enroll, api2sc.des.ps2)),
  ff(svymean(~enroll, api2sc.des.rake)),
  ff(svymean(~enroll, api2sc.jkdes)),
  ff(svymean(~enroll, api2sc.jkdes.ps))
)
```

SS

Take a look at the replicate weights after post-stratification

##	Weights								
##	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
## 1	0.00000	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 3	0.00000	1.00000	1.000000	1.000000	1.049550	1.00	1.00000	1.067176	1.000000
## 1	1.00945	0.00000	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 1	1.00945	1.00628	0.000000	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 2	1.00000	1.00000	0.000000	1.000000	1.030303	1.02	1.00000	1.000000	1.030303
## 1	1.00945	1.00628	1.047964	0.000000	1.000000	1.00	1.00628	1.012640	1.009450
## 3	1.04955	1.00000	1.000000	1.000000	0.000000	1.00	1.00000	1.067176	1.000000
## 2	1.00000	1.00000	1.077465	1.000000	0.000000	1.02	1.00000	1.000000	1.030303
## 2	1.00000	1.00000	1.077465	1.000000	1.030303	0.00	1.00000	1.000000	1.030303
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	0.00000	1.012640	1.009450
## 3	1.04955	1.00000	1.000000	1.000000	1.049550	1.00	1.00000	0.000000	1.000000
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	0.000000	1.009450
## 2	1.00000	1.00000	1.077465	1.000000	1.030303	1.02	1.00000	1.000000	0.000000
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	0.000000
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 2	1.00000	1.00000	1.077465	1.000000	1.030303	1.02	1.00000	1.000000	1.030303
## 2	1.00000	1.00000	1.077465	1.000000	1.030303	1.02	1.00000	1.000000	1.030303
## 2	1.00000	1.00000	1.077465	1.000000	1.030303	1.02	1.00000	1.000000	1.030303
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 3	1.04955	1.00000	1.000000	1.000000	1.049550	1.00	1.00000	1.067176	1.000000
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 3	1.04955	1.00000	1.000000	1.000000	1.049550	1.00	1.00000	1.067176	1.000000
## 2	1.00000	1.00000	1.077465	1.000000	1.030303	1.02	1.00000	1.000000	1.030303
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 3	1.04955	1.00000	1.000000	1.000000	1.049550	1.00	1.00000	1.067176	1.000000
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 2	1.00000	1.00000	1.077465	1.000000	1.030303	1.02	1.00000	1.000000	1.030303
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 3	1.04955	1.00000	1.000000	1.000000	1.049550	1.00	1.00000	1.067176	1.000000
## 1	1.00945	1.00628	1.047964	1.019081	1.000000	1.00	1.00628	1.012640	1.009450
## 3	1.04955	1.00000	1.000000	1.000000	1.049550	1.00	1.00000	1.067176	1.000000
## 1	1.00945	1.0							

[illegible]

[illegible]

[illegible]

```
## 1 0.00000 1.000000 1.000000 1.019081
## 2 1.02000 0.000000 1.085106 1.000000
## 3 1.00000 1.000000 0.000000 1.000000
## 2 1.02000 1.108696 0.000000 1.000000
## 1 1.00628 1.000000 1.000000 0.000000
##
## $index
## [1] 1 2 3
##
## attr(,"class")
## [1] "repweights_compressed" "repweights"
```

We might have been **supplied** these weights, someone else might have created them. In which case our dataset would look like this

```
fwmatt <- api2sc.jkdes.ps$repweights$weights[api2sc.jkdes.ps$repweights$index,]
colnames(fwmatt) <- paste0("repweight",sprintf("%0.3d",1:ncol(fwmatt)))
pweight <- api2sc.jkdes.ps$pweights
fwmatt <- diag(pweight)%*%fwmatt
sum(pweight)

## [1] 6157
apply(fwmatt,2,sum)
```

```
## repweight001 repweight002 repweight003 repweight004 repweight005 repweight006
##          6157          6157          6157          6157          6157          6157
## repweight007 repweight008 repweight009 repweight010 repweight011 repweight012
##          6157          6157          6157          6157          6157          6157
## repweight013 repweight014 repweight015 repweight016 repweight017 repweight018
##          6157          6157          6157          6157          6157          6157
## repweight019 repweight020 repweight021 repweight022 repweight023 repweight024
##          6157          6157          6157          6157          6157          6157
## repweight025 repweight026 repweight027 repweight028 repweight029 repweight030
##          6157          6157          6157          6157          6157          6157
## repweight031 repweight032 repweight033 repweight034 repweight035 repweight036
##          6157          6157          6157          6157          6157          6157
## repweight037 repweight038 repweight039 repweight040
##          6157          6157          6157          6157
```

```
api2sc.withrep.ps <- cbind(api2sc, pweight, fwmatt)
```

We create the design using the probability weights `pweight` and the set of replicate weights `repweight001-repweight040`.

```
pweightcol <- match("pweight",names(api2sc.withrep.ps)) # which column name is exactly
repweightcols <- grep("^repweight",names(api2sc.withrep.ps)) # which columnnames start
varcols <- (1:ncol(api2sc.withrep.ps))
varcols <- varcols[!(varcols%in%c(pweightcol,repweightcols))]
n <- length(unique(api2sc$dnum))
api2sc.repdes.ps <- svrepdesign(variables=api2sc.withrep.ps[,varcols],
                              weights=api2sc.withrep.ps[,pweightcol],
                              repweights=api2sc.withrep.ps[,repweightcols],
                              combined.weights=TRUE,
                              scale=(n-1)/n,
                              type="JK1")
```

```
svymean(~enroll, api2sc.repdes.ps)
```

```
##           mean      SE
## enroll 624.35 62.313
```

```
ss <- rbind(
  c(mean(api2sc$enroll), 0),
  ff(svymean(~enroll, api2sc.des)),
  ff(svymean(~enroll, api2sc.des.ps)),
  ff(svymean(~enroll, api2sc.des.ps2)),
  ff(svymean(~enroll, api2sc.des.rake)),
  ff(svymean(~enroll, api2sc.jkdes)),
  ff(svymean(~enroll, api2sc.jkdes.ps)),
  ff(svymean(~enroll, api2sc.repdes.ps))
)
dimnames(ss) <- list(c("Census", "2SC", "2SC+PS(stype)", "2SC+PS(stype x yr.rnd)", "2SC+Rake(stype+yr.rnd)",
  "JK1", "JK1+PS(stype)", "JK1+PS(stype)"),
  c("Estimate", "SE"))
ss
```

```
##           Estimate      SE
## Census          619.0469 0.00000
## 2SC              711.4530 81.05714
## 2SC+PS(stype)    624.3485 53.14154
## 2SC+PS(stype x yr.rnd) 629.4503 40.23927
## 2SC+Rake(stype+yr.rnd) 632.0372 46.70502
## JK1             711.4530 82.84780
## JK1+PS(stype)   624.3485 60.60973
## JK1+PS(stype)   624.3485 62.31259
```

JKn is used with stratified designs

E.g. draw a two stage cluster sample within two strata: schools which are or are not eligible for an awards programme:

```
set.seed(999)
apist2sc <- apipopc |>
  select.st2sc(n=40, stratumvar="awards", cfraction=0.20, clustervar="dnum")
apist2sc.des <- svydesign(id=~dnum+snum, strata=~awards, fpc=~bigN1+bigN2,
  data=apist2sc, nest=TRUE)
apist2sc.des.ps <- postStratify(apist2sc.des, strata=~stype, population=stype.benchmarks)
```

Note: we need `nest=TRUE` here because districts `dnum` do not nest within strata. I.e. it is the **combination** of `awards` and `dnum` that actually defines the clusters here.

```
svymean(~enroll, design=apist2sc.des)
```

```
##           mean      SE
## enroll 497.75 64.645
```

```
svymean(~enroll, design=apist2sc.des.ps)
```

```
##           mean      SE
## enroll 523.97 63.001
```

```
apist2sc.jkdes <- as.svrepdesign(apist2sc.des, type="JKn")
```

```
## Warning in as.svrepdesign.default(apist2sc.des, type = "JKn"): Finite
## population corrections after first stage have been dropped
```

```
svymean(~enroll, design=apist2sc.jkdes)
```

```
##           mean      SE
## enroll 497.75 84.997
```


```
apist2sc.jkdes.ps <- postStratify(apist2sc.jkdes, strata=~stype, population=stype.bench)
svymean(~enroll, design=apist2sc.jkdes.ps)
```

```
##           mean      SE
## enroll 523.97 82.125
```


There are alternative methods to Jackknife generation of replicate weights. You may hear of Bootstrap weights, and Balanced Repeated Replication. Some of these have multiple variants. We don't need to worry about these.

```
set.seed(222)
R <- 100
sim <- t(sapply(1:R,
  function(i) {
    ss <- apipopc |>
      select.2sc(n=100, cfraction=0.30, mmin=2, clustervar="dnum")
    d1 <- svydesign(ids=~dnum+snum, fpc=~bigN1+bigN2, data=ss)
    d2 <- postStratify(d1, strata=~stype, population=stype.benchmarks)
    e1 <- ff(svymean(~api00, design=d1))
    e2 <- ff(svymean(~api00, design=d2))
    return(c(e1,e2))
  })))

par(mfrow=c(1,2))
hist(sim[,1],main="Simple estimates",xlim=c(0.9,1.1)*range(sim[,c(1,3)]))
abline(v=mean(apipopc$api00),col="red",lwd=2)
hist(sim[,3],main="Postratified estimates",xlim=c(0.9,1.1)*range(sim[,c(1,3)]))
abline(v=mean(apipopc$api00),col="red",lwd=2)
```



```
par(mfrow=c(1,2))
hist(sim[,2],main="SE of Simple estimates",xlim=c(0.8,1.2)*range(sim[,c(2,4)]))
abline(v=sd(sim[,1]),col="red",lwd=2)
hist(sim[,4],main="SE of Postratified estimates",xlim=c(0.8,1.2)*range(sim[,c(2,4)]))
abline(v=sd(sim[,3]),col="red",lwd=2)
```



Chapter 9

Imputation

Dealing with item non-response

Can do complete case analysis, but that risks both bias and loss of precision due to reduced sample size.

Another option is imputation: fill in the data using a model. All such models involve untestable assumptions.

```
surf <- read.csv("data/surf.csv")[,1:8]
```

Remove some data: marital status and income

```
set.seed(111)
msurf <- surf
msurf$Marital[sample(nrow(surf),10)] <- NA
msurf$Income[sample(nrow(surf),10)] <- NA
table(msurf$Marital,exclude=NULL)
```

```
##
##   married      never      other previously      <NA>
##       68         82         20         20         10
```

```
sum(is.na(msurf$Income))
```

```
## [1] 10
```

```
library(mice)
library(lattice)
```

Package mice provides for multiple imputation methods, but gives some datasets and some visualisation tools. It's a bit tricky to use, so we'll use a simpler function `impute()`.

Example dataset supplied with mice is `nhanes` - note all variables are numeric

```
mnhanes <- mice::nhanes
mnhanes$age <- factor(mnhanes$age)
mnhanes <- data.frame(id=factor(1:nrow(mnhanes)),mnhanes)
dim(mnhanes) # 25 x 5
```

```
## [1] 25 5
```

```
summary(mnhanes)
```

```
##          id      age      bmi      hyp      chl
## 1      : 1    1:12  Min.    :20.40  Min.    :1.000  Min.    :113.0
## 2      : 1    2: 7  1st Qu.:22.65  1st Qu.:1.000  1st Qu.:185.0
## 3      : 1    3: 6  Median   :26.75  Median   :1.000  Median   :187.0
## 4      : 1      Mean    :26.56  Mean     :1.235  Mean     :191.4
## 5      : 1      3rd Qu.:28.93  3rd Qu.:1.000  3rd Qu.:212.0
## 6      : 1      Max.    :35.30  Max.     :2.000  Max.     :284.0
## (Other):19      NA's    :9      NA's     :8      NA's     :10
```


Surf data is a mixture of numeric and categorical variables

```
summary(msurf)
```

```
##      Personid      Gender      Qualification      Age
## Min.   : 1.00  Length:200  Length:200  Min.   :15.00
## 1st Qu.: 50.75  Class :character  Class :character  1st Qu.:22.75
## Median :100.50  Mode  :character  Mode  :character  Median :32.00
## Mean   :100.50      Mean   :30.69
## 3rd Qu.:150.25      3rd Qu.:38.00
## Max.   :200.00      Max.   :45.00
##
##      Hours      Income      Marital      Ethnicity
## Min.   : 2.00  Min.   : 11.0  Length:200  Length:200
## 1st Qu.:20.00  1st Qu.: 355.2  Class :character  Class :character
## Median :40.00  Median : 532.5  Mode  :character  Mode  :character
## Mean   :33.71  Mean   : 586.9
## 3rd Qu.:45.00  3rd Qu.: 819.8
## Max.   :70.00  Max.   :1789.0
##              NA's    :10
```


Inspect missing pattern

```
md.pattern(mnhanes)
```



```
##      id age hyp bmi chl
## 13  1  1  1  1  1  0
## 3   1  1  1  1  0  1
## 1   1  1  1  0  1  1
## 1   1  1  0  0  1  2
## 7   1  1  0  0  0  3
##    0  0  8  9 10 27
```

```
md.pattern(msurf)
```



```
##      Personid Gender Qualification Age Hours Ethnicity Income Marital
## 181         1      1             1  1    1           1      1        0
## 9          1      1             1  1    1           1      1        1
## 9          1      1             1  1    1           1      0        1
## 1          1      1             1  1    1           1      0        2
##          0      0             0  0    0           0     10       10 20
```

Mean, median or mode imputation

```
dd <- impute(mnhanes, varname="bmi", method="sample")
table(mnhanes$bmi, exclude=NULL)
```

```
##
## 20.4 21.7  22 22.5 22.7 24.9 25.5 26.3 27.2 27.4 27.5 28.7 29.6 30.1 33.2 35.3
##    1    1    1    1    1    1    1    1    1    1    1    1    1    1    1    1
## <NA>
##      9
```

```
table(dd$bmi, exclude=NULL)
```

```
##
## 20.4 21.7  22 22.5 22.7 24.9 25.5 26.3 27.2 27.4 27.5 28.7 29.6 30.1 33.2 35.3
##    1    1    2    2    2    3    1    1    3    1    1    2    2    1    1    1
```

```
dx <- mnhanes
dx$bmi[dx$age=="1"] <- NA
dd <- impute(dx, varname="bmi", method="sample", formula=~age)
```

```
## Warning in FUN(data[x, , drop = FALSE], ...): Cell with missing data but no
## donors: data not imputed
```

```
table(mnhanes$bmi,exclude=NULL)
```

```
##
## 20.4 21.7 22 22.5 22.7 24.9 25.5 26.3 27.2 27.4 27.5 28.7 29.6 30.1 33.2 35.3
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## <NA>
## 9
```

```
table(dd$bmi,exclude=NULL)
```

```
##
## 21.7 22 22.7 24.9 25.5 26.3 27.2 27.4 28.7 <NA>
## 1 1 2 3 1 2 1 1 1 12
```

```
varname <- "bmi"; tapply(mnhanes[,varname], mnhanes[, "age"], mean, na.rm=TRUE)
```

```
## 1 2 3
## 28.37143 25.42000 24.82500
```

Overall mean and median and mode

```
impute(mnhanes, varname="bmi", method="mean")
```

```
## id age bmi hyp chl bmi.imp
## 1 1 1 26.5625 NA NA TRUE
## 2 2 2 22.7000 1 187 FALSE
## 3 3 1 26.5625 1 187 TRUE
## 4 4 3 26.5625 NA NA TRUE
## 5 5 1 20.4000 1 113 FALSE
## 6 6 3 26.5625 NA 184 TRUE
## 7 7 1 22.5000 1 118 FALSE
## 8 8 1 30.1000 1 187 FALSE
## 9 9 2 22.0000 1 238 FALSE
## 10 10 2 26.5625 NA NA TRUE
## 11 11 1 26.5625 NA NA TRUE
## 12 12 2 26.5625 NA NA TRUE
## 13 13 3 21.7000 1 206 FALSE
## 14 14 2 28.7000 2 204 FALSE
## 15 15 1 29.6000 1 NA FALSE
## 16 16 1 26.5625 NA NA TRUE
## 17 17 3 27.2000 2 284 FALSE
## 18 18 2 26.3000 2 199 FALSE
## 19 19 1 35.3000 1 218 FALSE
## 20 20 3 25.5000 2 NA FALSE
## 21 21 1 26.5625 NA NA TRUE
## 22 22 1 33.2000 1 229 FALSE
## 23 23 1 27.5000 1 131 FALSE
## 24 24 3 24.9000 1 NA FALSE
## 25 25 2 27.4000 1 186 FALSE
```

```
impute(mnhanes, varname="bmi", method="median")
```

```
## id age bmi hyp chl bmi.imp
## 1 1 1 26.75 NA NA TRUE
## 2 2 2 22.70 1 187 FALSE
## 3 3 1 26.75 1 187 TRUE
```

```
## 4 4 3 26.75 NA NA TRUE
## 5 5 1 20.40 1 113 FALSE
## 6 6 3 26.75 NA 184 TRUE
## 7 7 1 22.50 1 118 FALSE
## 8 8 1 30.10 1 187 FALSE
## 9 9 2 22.00 1 238 FALSE
## 10 10 2 26.75 NA NA TRUE
## 11 11 1 26.75 NA NA TRUE
## 12 12 2 26.75 NA NA TRUE
## 13 13 3 21.70 1 206 FALSE
## 14 14 2 28.70 2 204 FALSE
## 15 15 1 29.60 1 NA FALSE
## 16 16 1 26.75 NA NA TRUE
## 17 17 3 27.20 2 284 FALSE
## 18 18 2 26.30 2 199 FALSE
## 19 19 1 35.30 1 218 FALSE
## 20 20 3 25.50 2 NA FALSE
## 21 21 1 26.75 NA NA TRUE
## 22 22 1 33.20 1 229 FALSE
## 23 23 1 27.50 1 131 FALSE
## 24 24 3 24.90 1 NA FALSE
## 25 25 2 27.40 1 186 FALSE
```

```
impute(msurf, varname="Marital", method="mode")
```

```
##      Personid Gender Qualification Age Hours Income      Marital Ethnicity
## 1           1 female      school  15    4    87      never European
## 2           2 female vocational  40   42   596    married European
## 3           3 male      none    38   40   497    married   Maori
## 4           4 female vocational  34    8   299      never European
## 5           5 female      school  45   16   301    married European
## 6           6 male      degree  45   50  1614    married European
## 7           7 female      none   36   12   201      other European
## 8           8 male      degree  35   45   934 previously European
## 9           9 female vocational  38   26   624    married European
## 10          10 male      school  37   50   533 previously European
## 11          11 male      none   44   45   609 previously European
## 12          12 female      school  35   40   620    married European
## 13          13 female      school  34   24   371    married European
## 14          14 female vocational  36   20   404    married European
## 15          15 female      school  19   40   623    married European
## 16          16 male      school  16   54   616      never   other
## 17          17 male      degree  41   54   856 previously European
## 18          18 male vocational  25   40   708      never European
## 19          19 male      degree  42   40   743    married European
## 20          20 female      none   34   25   386      other European
## 21          21 male vocational  24   47   999      never European
## 22          22 female vocational  42    4    74    married Pacific
## 23          23 female      school  16   20   387      never European
## 24          24 male      school  18    5    44      never   Maori
## 25          25 female      none   45   38   NA      never   Maori
## 26          26 male vocational  36   40   260      never   other
## 27          27 male      school  32   48   895    married European
```

## 28	28 female	vocational	33	14	368	never	European
## 29	29 female	none	43	40	559	never	European
## 30	30 male	vocational	31	40	526	other	other
## 31	31 male	vocational	29	60	1084	other	other
## 32	32 female	school	18	45	628	never	Maori
## 33	33 male	none	21	35	383	never	European
## 34	34 female	vocational	38	40	713	never	other
## 35	35 male	vocational	37	40	599	married	European
## 36	36 female	vocational	30	18	450	married	European
## 37	37 male	none	34	30	395	never	European
## 38	38 male	school	34	40	532	married	European
## 39	39 female	vocational	42	35	447	previously	European
## 40	40 male	vocational	29	37	795	married	European
## 41	41 male	school	35	40	809	married	European
## 42	42 female	school	22	15	309	never	European
## 43	43 male	degree	38	50	1217	previously	Maori
## 44	44 male	none	36	35	426	married	European
## 45	45 male	none	17	15	222	never	European
## 46	46 female	school	45	44	508	married	European
## 47	47 female	school	39	17	351	married	Maori
## 48	48 female	none	43	6	206	married	European
## 49	49 male	none	34	54	826	never	European
## 50	50 male	degree	31	70	1724	married	European
## 51	51 female	school	41	20	264	previously	Maori
## 52	52 male	none	37	48	949	married	European
## 53	53 female	school	32	24	143	other	European
## 54	54 female	none	27	40	474	never	European
## 55	55 female	school	45	40	708	previously	European
## 56	56 female	vocational	27	5	112	never	European
## 57	57 male	vocational	40	39	525	previously	Maori
## 58	58 male	vocational	34	55	849	married	European
## 59	59 female	school	22	6	NA	never	European
## 60	60 male	school	22	12	85	never	European
## 61	61 female	school	19	15	231	married	European
## 62	62 female	degree	28	50	544	married	European
## 63	63 female	vocational	38	32	462	married	European
## 64	64 male	none	18	14	239	never	European
## 65	65 female	degree	29	60	1543	never	European
## 66	66 male	school	35	40	548	married	European
## 67	67 female	vocational	41	17	18	married	European
## 68	68 female	vocational	20	25	552	never	Maori
## 69	69 female	none	45	40	556	never	Maori
## 70	70 female	none	43	25	266	previously	European
## 71	71 female	none	25	6	286	married	Pacific
## 72	72 male	vocational	21	70	1157	never	Maori
## 73	73 male	none	28	54	480	married	European
## 74	74 female	degree	36	50	979	never	European
## 75	75 male	school	22	44	475	never	European
## 76	76 male	none	18	35	507	never	Pacific
## 77	77 male	school	21	10	107	never	European
## 78	78 male	school	44	36	NA	never	European
## 79	79 female	school	24	16	240	married	European
## 80	80 female	school	34	38	567	married	European

## 81	81	male	none	29	40	477	other	other
## 82	82	male	none	39	35	662	other	European
## 83	83	male	none	43	70	1005	married	European
## 84	84	male	school	17	40	642	never	European
## 85	85	male	degree	31	40	985	never	European
## 86	86	female	vocational	35	40	501	previously	Maori
## 87	87	male	vocational	23	25	460	never	European
## 88	88	female	school	18	6	86	other	European
## 89	89	female	school	40	15	287	married	European
## 90	90	female	school	18	9	NA	other	Pacific
## 91	91	female	vocational	24	40	630	never	European
## 92	92	female	none	42	40	402	married	European
## 93	93	female	school	27	20	406	never	Maori
## 94	94	male	vocational	23	40	1062	never	European
## 95	95	female	none	16	3	108	never	European
## 96	96	female	none	43	9	220	married	European
## 97	97	female	vocational	34	40	615	never	European
## 98	98	female	school	30	40	820	never	Maori
## 99	99	female	degree	24	39	626	never	European
## 100	100	male	vocational	34	60	1373	married	European
## 101	101	female	school	21	35	874	other	European
## 102	102	male	vocational	41	46	972	married	European
## 103	103	female	vocational	30	34	609	never	European
## 104	104	male	vocational	27	60	710	other	European
## 105	105	male	school	35	70	1131	married	European
## 106	106	male	vocational	40	50	868	other	other
## 107	107	female	degree	40	12	392	other	European
## 108	108	female	none	33	40	517	married	European
## 109	109	female	vocational	23	38	517	never	Maori
## 110	110	female	vocational	34	38	532	previously	European
## 111	111	female	school	40	21	NA	married	European
## 112	112	male	school	19	16	110	never	European
## 113	113	male	school	16	16	NA	never	European
## 114	114	female	school	21	2	11	never	Maori
## 115	115	male	none	38	50	835	married	European
## 116	116	female	degree	36	40	925	previously	European
## 117	117	male	degree	43	45	950	married	European
## 118	118	male	degree	24	40	860	never	European
## 119	119	female	degree	40	13	425	never	European
## 120	120	female	school	33	17	244	never	Pacific
## 121	121	female	vocational	21	40	830	other	European
## 122	122	male	vocational	32	38	562	married	European
## 123	123	male	degree	41	45	1145	married	other
## 124	124	female	vocational	37	6	193	married	European
## 125	125	male	school	23	45	515	never	European
## 126	126	female	none	17	17	259	never	European
## 127	127	male	none	43	10	99	never	Maori
## 128	128	female	vocational	20	50	757	never	other
## 129	129	male	vocational	27	40	952	never	European
## 130	130	female	vocational	38	40	1097	never	European
## 131	131	female	none	20	2	105	never	European
## 132	132	female	vocational	44	48	743	married	European
## 133	133	female	school	31	10	218	previously	European

## 134	134 female	school	22	8	94	never	Pacific
## 135	135 male	degree	33	38	NA	married	European
## 136	136 male	degree	26	38	NA	never	European
## 137	137 male	vocational	45	44	902	married	European
## 138	138 female	none	32	40	406	married	European
## 139	139 male	degree	41	45	1034	never	Maori
## 140	140 male	vocational	24	54	868	never	European
## 141	141 female	degree	23	45	658	never	other
## 142	142 female	vocational	38	32	521	married	European
## 143	143 male	vocational	34	45	1099	previously	European
## 144	144 male	degree	33	60	540	previously	European
## 145	145 female	vocational	34	40	967	married	European
## 146	146 male	vocational	37	40	1163	married	European
## 147	147 female	none	32	10	211	married	European
## 148	148 male	school	16	40	539	never	European
## 149	149 male	school	27	45	741	never	Maori
## 150	150 male	vocational	40	45	768	married	European
## 151	151 male	vocational	18	40	1099	other	European
## 152	152 female	school	18	17	NA	never	European
## 153	153 male	none	33	40	496	married	Maori
## 154	154 female	vocational	27	39	724	other	European
## 155	155 female	vocational	26	42	431	married	European
## 156	156 male	vocational	37	40	819	married	European
## 157	157 male	degree	36	38	849	married	European
## 158	158 female	degree	39	55	1152	other	European
## 159	159 female	vocational	42	27	525	married	other
## 160	160 male	vocational	30	24	463	never	European
## 161	161 male	vocational	22	65	884	never	European
## 162	162 male	vocational	35	45	700	married	European
## 163	163 female	vocational	25	25	409	never	Maori
## 164	164 female	vocational	42	24	437	previously	European
## 165	165 female	none	45	40	439	never	European
## 166	166 female	school	18	4	45	never	Maori
## 167	167 female	school	16	9	232	never	European
## 168	168 female	vocational	16	46	781	married	European
## 169	169 male	school	25	40	506	never	European
## 170	170 female	vocational	22	40	442	never	European
## 171	171 male	vocational	20	58	1147	other	other
## 172	172 male	school	19	32	156	never	European
## 173	173 female	school	20	10	62	never	European
## 174	174 female	vocational	33	48	801	married	European
## 175	175 male	school	21	50	867	never	Maori
## 176	176 female	school	30	19	NA	never	European
## 177	177 female	school	37	20	425	married	European
## 178	178 male	vocational	45	40	564	never	European
## 179	179 female	none	15	40	323	other	Maori
## 180	180 male	school	18	60	759	never	other
## 181	181 male	school	28	20	161	never	European
## 182	182 male	vocational	34	40	863	previously	European
## 183	183 male	school	22	40	375	never	European
## 184	184 male	school	18	45	455	never	European
## 185	185 male	school	23	40	706	never	European
## 186	186 female	school	16	20	264	never	European

## 187	187 female	none	42	40	614	never	European
## 188	188 female	vocational	25	8	132	never	European
## 189	189 female	school	39	15	288	previously	European
## 190	190 female	degree	32	40	581	married	European
## 191	191 female	none	15	8	172	never	European
## 192	192 female	school	28	8	227	never	Pacific
## 193	193 male	vocational	36	40	927	other	European
## 194	194 female	degree	32	20	746	previously	European
## 195	195 female	school	21	35	485	never	European
## 196	196 female	school	45	24	386	married	European
## 197	197 male	degree	26	50	1789	never	European
## 198	198 male	degree	25	45	1558	never	European
## 199	199 female	none	38	20	230	married	European
## 200	200 male	school	31	50	954	never	European
##	Marital.imp						
## 1	FALSE						
## 2	FALSE						
## 3	FALSE						
## 4	FALSE						
## 5	FALSE						
## 6	FALSE						
## 7	FALSE						
## 8	FALSE						
## 9	FALSE						
## 10	FALSE						
## 11	FALSE						
## 12	FALSE						
## 13	FALSE						
## 14	FALSE						
## 15	FALSE						
## 16	FALSE						
## 17	FALSE						
## 18	FALSE						
## 19	FALSE						
## 20	FALSE						
## 21	FALSE						
## 22	FALSE						
## 23	FALSE						
## 24	FALSE						
## 25	TRUE						
## 26	TRUE						
## 27	FALSE						
## 28	FALSE						
## 29	FALSE						
## 30	FALSE						
## 31	FALSE						
## 32	FALSE						
## 33	FALSE						
## 34	FALSE						
## 35	FALSE						
## 36	FALSE						
## 37	FALSE						
## 38	FALSE						

## 39	FALSE
## 40	FALSE
## 41	FALSE
## 42	FALSE
## 43	FALSE
## 44	FALSE
## 45	TRUE
## 46	FALSE
## 47	FALSE
## 48	FALSE
## 49	TRUE
## 50	FALSE
## 51	FALSE
## 52	FALSE
## 53	FALSE
## 54	FALSE
## 55	FALSE
## 56	FALSE
## 57	FALSE
## 58	FALSE
## 59	FALSE
## 60	FALSE
## 61	FALSE
## 62	FALSE
## 63	FALSE
## 64	FALSE
## 65	FALSE
## 66	FALSE
## 67	FALSE
## 68	FALSE
## 69	TRUE
## 70	FALSE
## 71	FALSE
## 72	TRUE
## 73	FALSE
## 74	TRUE
## 75	FALSE
## 76	FALSE
## 77	FALSE
## 78	FALSE
## 79	FALSE
## 80	FALSE
## 81	FALSE
## 82	FALSE
## 83	FALSE
## 84	FALSE
## 85	FALSE
## 86	FALSE
## 87	FALSE
## 88	FALSE
## 89	FALSE
## 90	FALSE
## 91	FALSE

## 92	FALSE
## 93	FALSE
## 94	FALSE
## 95	FALSE
## 96	FALSE
## 97	FALSE
## 98	FALSE
## 99	FALSE
## 100	FALSE
## 101	FALSE
## 102	FALSE
## 103	FALSE
## 104	FALSE
## 105	FALSE
## 106	FALSE
## 107	FALSE
## 108	FALSE
## 109	FALSE
## 110	FALSE
## 111	FALSE
## 112	FALSE
## 113	FALSE
## 114	FALSE
## 115	FALSE
## 116	FALSE
## 117	FALSE
## 118	FALSE
## 119	FALSE
## 120	FALSE
## 121	FALSE
## 122	FALSE
## 123	FALSE
## 124	FALSE
## 125	FALSE
## 126	FALSE
## 127	FALSE
## 128	FALSE
## 129	FALSE
## 130	FALSE
## 131	FALSE
## 132	FALSE
## 133	FALSE
## 134	FALSE
## 135	FALSE
## 136	FALSE
## 137	FALSE
## 138	FALSE
## 139	FALSE
## 140	FALSE
## 141	FALSE
## 142	FALSE
## 143	FALSE
## 144	FALSE

## 145	FALSE
## 146	FALSE
## 147	FALSE
## 148	FALSE
## 149	FALSE
## 150	FALSE
## 151	FALSE
## 152	FALSE
## 153	FALSE
## 154	FALSE
## 155	FALSE
## 156	FALSE
## 157	FALSE
## 158	FALSE
## 159	FALSE
## 160	FALSE
## 161	FALSE
## 162	FALSE
## 163	TRUE
## 164	FALSE
## 165	FALSE
## 166	FALSE
## 167	FALSE
## 168	FALSE
## 169	FALSE
## 170	FALSE
## 171	FALSE
## 172	FALSE
## 173	FALSE
## 174	FALSE
## 175	TRUE
## 176	FALSE
## 177	FALSE
## 178	FALSE
## 179	FALSE
## 180	FALSE
## 181	FALSE
## 182	FALSE
## 183	FALSE
## 184	FALSE
## 185	FALSE
## 186	FALSE
## 187	TRUE
## 188	FALSE
## 189	FALSE
## 190	FALSE
## 191	FALSE
## 192	FALSE
## 193	FALSE
## 194	FALSE
## 195	FALSE
## 196	FALSE
## 197	FALSE

```
## 198      FALSE
## 199      FALSE
## 200      FALSE
```

Overall mean

```
dd <- impute(msurf, varname="Income", method="mean")
dd[dd$Income.imp,]
```

```
##      Personid Gender Qualification Age Hours   Income Marital Ethnicity
## 25         25 female         none  45    38 586.8895    <NA>    Maori
## 59         59 female      school  22     6 586.8895   never  European
## 78         78  male      school  44    36 586.8895   never  European
## 90         90 female      school  18     9 586.8895   other  Pacific
## 111        111 female      school  40    21 586.8895 married European
## 113        113  male      school  16    16 586.8895   never  European
## 135        135  male      degree  33    38 586.8895 married European
## 136        136  male      degree  26    38 586.8895   never  European
## 152        152 female      school  18    17 586.8895   never  European
## 176        176 female      school  30    19 586.8895   never  European
```

```
##      Income.imp
## 25          TRUE
## 59          TRUE
## 78          TRUE
## 90          TRUE
## 111         TRUE
## 113         TRUE
## 135         TRUE
## 136         TRUE
## 152         TRUE
## 176         TRUE
```

```
dd <- impute(msurf, varname="Income", method="lm.mean", formula=-1)
dd[dd$Income.imp,]
```

```
##      Personid Gender Qualification Age Hours   Income Marital Ethnicity
## 25         25 female         none  45    38 586.8895    <NA>    Maori
## 59         59 female      school  22     6 586.8895   never  European
## 78         78  male      school  44    36 586.8895   never  European
## 90         90 female      school  18     9 586.8895   other  Pacific
## 111        111 female      school  40    21 586.8895 married European
## 113        113  male      school  16    16 586.8895   never  European
## 135        135  male      degree  33    38 586.8895 married European
## 136        136  male      degree  26    38 586.8895   never  European
## 152        152 female      school  18    17 586.8895   never  European
## 176        176 female      school  30    19 586.8895   never  European
```

```
##      Income.imp
## 25          TRUE
## 59          TRUE
## 78          TRUE
## 90          TRUE
## 111         TRUE
## 113         TRUE
## 135         TRUE
## 136         TRUE
```

```
## 152      TRUE
## 176      TRUE
```

```
mean(msurf$Income, na.rm=TRUE)
```

```
## [1] 586.8895
```

```
mean(dd$Income)
```

```
## [1] 586.8895
```

```
sd(msurf$Income, na.rm=TRUE)
```

```
## [1] 346.505
```

```
sd(dd$Income)
```

```
## [1] 337.6866
```

```
dd <- impute(msurf, varname="Income", method="lm.sim", formula=~1)
dd[dd$Income.imp,]
```

##	Personid	Gender	Qualification	Age	Hours	Income	Marital	Ethnicity
## 25	25	female	none	45	38	707.11435	<NA>	Maori
## 59	59	female	school	22	6	652.63412	never	European
## 78	78	male	school	44	36	531.59532	never	European
## 90	90	female	school	18	9	700.04041	other	Pacific
## 111	111	female	school	40	21	794.18753	married	European
## 113	113	male	school	16	16	-51.21132	never	European
## 135	135	male	degree	33	38	1528.70926	married	European
## 136	136	male	degree	26	38	653.15661	never	European
## 152	152	female	school	18	17	135.98391	never	European
## 176	176	female	school	30	19	-491.85581	never	European

```
##      Income.imp
```

```
## 25      TRUE
```

```
## 59      TRUE
```

```
## 78      TRUE
```

```
## 90      TRUE
```

```
## 111     TRUE
```

```
## 113     TRUE
```

```
## 135     TRUE
```

```
## 136     TRUE
```

```
## 152     TRUE
```

```
## 176     TRUE
```

```
mean(msurf$Income, na.rm=TRUE)
```

```
## [1] 586.8895
```

```
mean(dd$Income)
```

```
## [1] 583.3468
```

```
sd(msurf$Income, na.rm=TRUE)
```

```
## [1] 346.505
```

```
sd(dd$Income)
```

```
## [1] 357.4977
```

```
dd <- impute(msurf, varname="Income", method="lm.mean", formula=~Gender)
dd[dd$Income.imp,]
```

```
##      Personid Gender Qualification Age Hours   Income Marital Ethnicity
## 25          25 female          none  45    38 451.6436    <NA>    Maori
## 59          59 female      school  22     6 451.6436   never   European
## 78          78 male       school  44    36 740.3708   never   European
## 90          90 female      school  18     9 451.6436   other   Pacific
## 111         111 female      school  40    21 451.6436 married European
## 113         113 male       school  16    16 740.3708   never   European
## 135         135 male       degree  33    38 740.3708 married European
## 136         136 male       degree  26    38 740.3708   never   European
## 152         152 female      school  18    17 451.6436   never   European
## 176         176 female      school  30    19 451.6436   never   European
```

```
##      Income.imp
```

```
## 25          TRUE
## 59          TRUE
## 78          TRUE
## 90          TRUE
## 111         TRUE
## 113         TRUE
## 135         TRUE
## 136         TRUE
## 152         TRUE
## 176         TRUE
```

```
tapply(msurf$Income, msurf$Gender, mean, na.rm=TRUE)
```

```
##   female    male
## 451.6436 740.3708
```

```
tapply(dd$Income, dd$Gender, mean)
```

```
##   female    male
## 451.6436 740.3708
```

```
tapply(msurf$Income, msurf$Gender, sd, na.rm=TRUE)
```

```
##   female    male
## 270.5141 360.4080
```

```
tapply(dd$Income, dd$Gender, sd)
```

```
##   female    male
## 262.7465 352.4859
```

```
dd <- impute(msurf, varname="Income", method="lm.sim", formula=~Gender)
dd[dd$Income.imp,]
```

```
##      Personid Gender Qualification Age Hours   Income Marital Ethnicity
## 25          25 female          none  45    38 154.36863    <NA>    Maori
## 59          59 female      school  22     6 893.83677   never   European
## 78          78 male       school  44    36 228.63621   never   European
## 90          90 female      school  18     9 -263.94417   other   Pacific
## 111         111 female      school  40    21 818.90986 married European
## 113         113 male       school  16    16 703.68973   never   European
```

```
## 135      135   male      degree 33    38 845.92680 married European
## 136      136   male      degree 26    38 544.30757   never European
## 152      152 female      school 18    17  38.00251   never European
## 176      176 female      school 30    19  80.35643   never European
```

```
##      Income.imp
## 25          TRUE
## 59          TRUE
## 78          TRUE
## 90          TRUE
## 111         TRUE
## 113         TRUE
## 135         TRUE
## 136         TRUE
## 152         TRUE
## 176         TRUE
```

```
tapply(msurf$Income, msurf$Gender, mean, na.rm=TRUE)
```

```
##   female      male
## 451.6436 740.3708
```

```
tapply(dd$Income, dd$Gender, mean)
```

```
##   female      male
## 442.4068 733.5006
```

```
tapply(msurf$Income, msurf$Gender, sd, na.rm=TRUE)
```

```
##   female      male
## 270.5141 360.4080
```

```
tapply(dd$Income, dd$Gender, sd)
```

```
##   female      male
## 283.9823 357.2095
```

Overall mean imputation

```
mnhanes.mfit1 <- mnhanes |> mice(method="mean", m=1, maxit=1)
```

```
##
## iter imp variable
## 1 1 bmi hyp chl
```

```
## Warning: Number of logged events: 6
```

```
mnhanes.m1 <- complete(mnhanes.mfit1)
md.pattern(mnhanes.m1) # no missing data left
```

```
## /\      /\
## { '---' }
## { 0  0 }
## ==> V <== No need for mice. This data set is completely observed.
## \  \||/ /
##  '-----'
```

	id	age	bmi	hyp	chl	
25						0
	0	0	0	0	0	0

```
##      id age bmi hyp chl
## 25  1   1   1   1   1  0
##      0   0   0   0   0  0
```

Overall mean imputation

```
mnhanes.grandmean <- mnhanes |>
  mice(method="mean", m=1) |>
  complete()
```

```
##
## iter imp variable
## 1  1  bmi  hyp  chl
## 2  1  bmi  hyp  chl
## 3  1  bmi  hyp  chl
## 4  1  bmi  hyp  chl
## 5  1  bmi  hyp  chl
```

```
## Warning: Number of logged events: 18
```

Cell mean imputation

```
mnhanes.cellmean <- do.call(rbind,
  by(mnhanes, mnhanes$age,
    function(dsub) complete(mice(dsub, method="mean", m=1))
```

```
##
## iter imp variable
## 1  1  bmi  chl
## 2  1  bmi  chl
## 3  1  bmi  chl
## 4  1  bmi  chl
```

```
## 5 1 bmi chl
## Warning: Number of logged events: 14
##
## iter imp variable
## 1 1 bmi hyp chl
## 2 1 bmi hyp chl
## 3 1 bmi hyp chl
## 4 1 bmi hyp chl
## 5 1 bmi hyp chl
## Warning: Number of logged events: 19
##
## iter imp variable
## 1 1 bmi hyp
## 2 1 bmi hyp
## 3 1 bmi hyp
## 4 1 bmi hyp
## 5 1 bmi hyp
## Warning: Number of logged events: 14
mnhanes.cellmean <- arrange(mnhanes.cellmean, id)
```

Compare

```
mnhanes.grandmean[is.na(mnhanes$bmi),]
```

```
## id age bmi hyp chl
## 1 1 1 26.5625 1.235294 191.4
## 3 3 1 26.5625 1.000000 187.0
## 4 4 3 26.5625 1.235294 191.4
## 6 6 3 26.5625 1.235294 184.0
## 10 10 2 26.5625 1.235294 191.4
## 11 11 1 26.5625 1.235294 191.4
## 12 12 2 26.5625 1.235294 191.4
## 16 16 1 26.5625 1.235294 191.4
## 21 21 1 26.5625 1.235294 191.4
```

```
mnhanes.cellmean[is.na(mnhanes$bmi),]
```

```
## id age bmi hyp chl
## 1.1 1 1 28.37143 NA 169.0
## 1.3 3 1 28.37143 1.0 187.0
## 3.4 4 3 24.82500 1.5 NA
## 3.6 6 3 24.82500 1.5 184.0
## 2.10 10 2 25.42000 1.4 202.8
## 1.11 11 1 28.37143 NA 169.0
## 2.12 12 2 25.42000 1.4 202.8
## 1.16 16 1 28.37143 NA 169.0
## 1.21 21 1 28.37143 NA 169.0
```

Repeat cell mean imputation using the regression formulation

```
allvarnames <- names(mnhanes)
```

```

blocks <- as.list(allvarnames) # each variable in its own block
names(blocks) <- allvarnames
method <- rep("",length(allvarnames)) # method for each block: start with no imputation
names(method) <- allvarnames
formulas <- as.list(rep("~1", length(allvarnames)))
names(formulas) <- allvarnames
formulas <- lapply(formulas, formula)

method["bmi"] <- "norm.predict" # we're only going to impute bmi, by regression imputation
formulas[["bmi"]] <- formula(~age)

mnhanes.cellmeanregf <- mnhanes |>
  mice(m=1, method=method, blocks=blocks, formulas=formulas) |>
  complete()

```

```

##
## iter imp variable
## 1 1 bmi
## 2 1 bmi
## 3 1 bmi
## 4 1 bmi
## 5 1 bmi

```

```
mnhanes.cellmeanregf
```

```

##      id age      bmi hyp chl
## 1  1  1  28.37143  NA  NA
## 2  2  2  22.70000   1 187
## 3  3  1  28.37143   1 187
## 4  4  3  24.82500  NA  NA
## 5  5  1  20.40000   1 113
## 6  6  3  24.82500  NA 184
## 7  7  1  22.50000   1 118
## 8  8  1  30.10000   1 187
## 9  9  2  22.00000   1 238
## 10 10  2  25.42000  NA  NA
## 11 11  1  28.37143  NA  NA
## 12 12  2  25.42000  NA  NA
## 13 13  3  21.70000   1 206
## 14 14  2  28.70000   2 204
## 15 15  1  29.60000   1  NA
## 16 16  1  28.37143  NA  NA
## 17 17  3  27.20000   2 284
## 18 18  2  26.30000   2 199
## 19 19  1  35.30000   1 218
## 20 20  3  25.50000   2  NA
## 21 21  1  28.37143  NA  NA
## 22 22  1  33.20000   1 229
## 23 23  1  27.50000   1 131
## 24 24  3  24.90000   1  NA
## 25 25  2  27.40000   1 186

```

Compare

```
mnhanes.grandmean[is.na(mnhanes$bmi),]
```

```
##      id age      bmi      hyp    chl
## 1     1  1  26.5625 1.235294 191.4
## 3     3  3  26.5625 1.000000 187.0
## 4     4  4  26.5625 1.235294 191.4
## 6     6  6  26.5625 1.235294 184.0
## 10    10  2  26.5625 1.235294 191.4
## 11    11  1  26.5625 1.235294 191.4
## 12    12  2  26.5625 1.235294 191.4
## 16    16  1  26.5625 1.235294 191.4
## 21    21  1  26.5625 1.235294 191.4
```

```
mnhanes.cellmean[is.na(mnhanes$bmi),]
```

```
##      id age      bmi hyp    chl
## 1.1    1  1  28.37143  NA 169.0
## 1.3    3  1  28.37143  1.0 187.0
## 3.4    4  3  24.82500  1.5   NA
## 3.6    6  3  24.82500  1.5 184.0
## 2.10   10  2  25.42000  1.4 202.8
## 1.11   11  1  28.37143  NA 169.0
## 2.12   12  2  25.42000  1.4 202.8
## 1.16   16  1  28.37143  NA 169.0
## 1.21   21  1  28.37143  NA 169.0
```

```
mnhanes.cellmeanregf[is.na(mnhanes$bmi),]
```

```
##      id age      bmi hyp chl
## 1     1  1  28.37143  NA  NA
## 3     3  1  28.37143   1 187
## 4     4  3  24.82500  NA  NA
## 6     6  3  24.82500  NA 184
## 10    10  2  25.42000  NA  NA
## 11    11  1  28.37143  NA  NA
## 12    12  2  25.42000  NA  NA
## 16    16  1  28.37143  NA  NA
## 21    21  1  28.37143  NA  NA
```

Means of numeric and modes of categorical variables

```
getmeanmode <- function(df) {
  as.data.frame(lapply(df,
    function(x) {
      if(is.factor(x) || is.character(x)) {
        xt <- table(x)
        mval <- names(xt)[which.max(xt)]
      } else if(is.numeric(x)) {
        mval <- mean(x, na.rm=TRUE)
      } else {
        mval <- NA
      }
      return(mval)
    })
}
```

```
getmeanmode(mnhanes)
```

```
##   id age      bmi      hyp   chl
## 1  1    1 26.5625 1.235294 191.4
```

```
sapply(names(mnhanes.m1), function(varname) mnhanes.m1[is.na(mnhanes[,varname]),varname])
```

```
## $id
## factor()
## 25 Levels: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 ... 25
##
## $age
## factor()
## Levels: 1 2 3
##
## $bmi
## [1] 26.5625 26.5625 26.5625 26.5625 26.5625 26.5625 26.5625 26.5625 26.5625
##
## $hyp
## [1] 1.235294 1.235294 1.235294 1.235294 1.235294 1.235294 1.235294 1.235294
##
## $chl
## [1] 191.4 191.4 191.4 191.4 191.4 191.4 191.4 191.4 191.4 191.4
```


Just do the numerical variables (gives a warning because Marital is categorical)

```
msurf.mfit1 <- msurf |> mice(method="mean", m=1, maxit=1)
```

```
##
##   iter imp variable
##    1    1 Income
```

```
## Warning: Number of logged events: 4
```

```
msurf.m1 <- complete(msurf.mfit1)
md.pattern(msurf.m1) # still has missing Marital data
```




```
##      Personid Gender Qualification Age Hours Income Ethnicity Marital
## 190          1      1              1  1    1      1          1      1  0
## 10           1      1              1  1    1      1          1      0  1
##           0      0              0  0    0      0          0     10 10
```

Avoid the error

```
msurf.mfit1 <- msurf[,sapply(msurf,is.numeric)] |> mice(method="mean", m=1, maxit=1)
```

```
##
## iter imp variable
## 1 1 Income
msurf.m1 <- complete(msurf.mfit1)
md.pattern(msurf.m1) # this deletes the Marital column of course
```

```
## /\      /\
## { '---' }
## { 0 0 }
## ==> V <== No need for mice. This data set is completely observed.
## \  \|/  /
## '-----'
```



```
##      Personid Age Hours Income
## 200         1   1     1      1 0
##           0   0     0      0 0
```

Hot deck imputation

```
mnhanes.hdfit1 <- mice(mnhanes, method="sample", m=1, seed=123)
```

```
##
## iter imp variable
## 1 1 bmi hyp chl
## 2 1 bmi hyp chl
## 3 1 bmi hyp chl
## 4 1 bmi hyp chl
## 5 1 bmi hyp chl
```

```
## Warning: Number of logged events: 18
```

```
mnhanes.hd1 <- complete(mnhanes.hdfit1)
mnhanes.hd1[is.na(mnhanes$bmi),]
```

```
##   id age  bmi hyp chl
## 1  1   1 27.5  1 206
## 3  3   1 25.5  1 187
## 4  4   3 33.2  1 131
## 6  6   3 22.5  1 184
## 10 10   2 27.2  2 131
## 11 11   1 28.7  1 184
## 12 12   2 21.7  1 204
## 16 16   1 26.3  1 187
## 21 21   1 27.2  1 184
```

```
table(mnhanes$bmi,exclude=NULL)
```

```
##
## 20.4 21.7 22 22.5 22.7 24.9 25.5 26.3 27.2 27.4 27.5 28.7 29.6 30.1 33.2 35.3
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## <NA>
## 9
```

```
table(mnhanes.hd1$bmi)
```

```
##
## 20.4 21.7 22 22.5 22.7 24.9 25.5 26.3 27.2 27.4 27.5 28.7 29.6 30.1 33.2 35.3
## 1 2 1 2 1 1 2 2 3 1 2 2 1 1 2 1
```

Hot deck within cells [*** not yet done ***]

```
mnhanes.hdfit1 <- mice(mnhanes, method="sample", m=1, blocks=, seed=123)
```

```
##
## iter imp variable
## 1 1 bmi hyp chl
## 2 1 bmi hyp chl
## 3 1 bmi hyp chl
## 4 1 bmi hyp chl
## 5 1 bmi hyp chl
```

Warning: Number of logged events: 18

```
mnhanes.hd1 <- complete(mnhanes.hdfit1)
mnhanes.hd1[is.na(mnhanes$bmi),]
```

```
## id age bmi hyp chl
## 1 1 1 27.5 1 206
## 3 3 1 25.5 1 187
## 4 4 3 33.2 1 131
## 6 6 3 22.5 1 184
## 10 10 2 27.2 2 131
## 11 11 1 28.7 1 184
## 12 12 2 21.7 1 204
## 16 16 1 26.3 1 187
## 21 21 1 27.2 1 184
```

```
table(mnhanes$bmi,exclude=NULL)
```

```
##
## 20.4 21.7 22 22.5 22.7 24.9 25.5 26.3 27.2 27.4 27.5 28.7 29.6 30.1 33.2 35.3
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## <NA>
## 9
```

```
table(mnhanes.hd1$bmi)
```

```
##
## 20.4 21.7 22 22.5 22.7 24.9 25.5 26.3 27.2 27.4 27.5 28.7 29.6 30.1 33.2 35.3
## 1 2 1 2 1 1 2 2 3 1 2 2 1 1 2 1
```

Regression imputation

Use regressions on all available observations?

```
mnhanes.mfit2 <- mice(mnhanes, method="norm.predict", m=1, maxit=1)
```

```
##
## iter imp variable
## 1 1 bmi* hyp chl*

## Warning: Number of logged events: 6
```

```
mnhanes.m2 <- complete(mnhanes.mfit2)
mnhanes.m1[is.na(mnhanes$bmi),]
```

```
## id age bmi hyp chl
## 1 1 1 26.5625 1.235294 191.4
## 3 3 1 26.5625 1.000000 187.0
## 4 4 3 26.5625 1.235294 191.4
## 6 6 3 26.5625 1.235294 184.0
## 10 10 2 26.5625 1.235294 191.4
## 11 11 1 26.5625 1.235294 191.4
## 12 12 2 26.5625 1.235294 191.4
## 16 16 1 26.5625 1.235294 191.4
## 21 21 1 26.5625 1.235294 191.4
```

```
mnhanes.m2[is.na(mnhanes$bmi),]
```

```
## id age bmi hyp chl
## 1 1 1 NA NA NA
## 3 3 1 NA 1 187
## 4 4 3 NA NA NA
## 6 6 3 NA NA 184
## 10 10 2 NA NA NA
## 11 11 1 NA NA NA
## 12 12 2 NA NA NA
## 16 16 1 NA NA NA
## 21 21 1 NA NA NA
```

Multiple imputation

Instead of imputing a **single** value for each missing observation, we impute several possible imputed values - we end up with multiple completed data sets. The variability among estimates from these separate completed data sets gives us additional insight into the effect of our imputations, and the uncertainty we have associated with having missing data.

The confidence we can have regarding the true scale of the uncertainty depends on how strongly we believe in the model we've used to impute.

Hot deck multiple imputation

```
mnhanes.multfit1 <- mice(mnhanes, method="sample", m=5, seed=123)
```

```
##
## iter imp variable
## 1 1 bmi hyp chl
## 1 2 bmi hyp chl
## 1 3 bmi hyp chl
## 1 4 bmi hyp chl
## 1 5 bmi hyp chl
```

```
## 2 1 bmi hyp chl
## 2 2 bmi hyp chl
## 2 3 bmi hyp chl
## 2 4 bmi hyp chl
## 2 5 bmi hyp chl
## 3 1 bmi hyp chl
## 3 2 bmi hyp chl
## 3 3 bmi hyp chl
## 3 4 bmi hyp chl
## 3 5 bmi hyp chl
## 4 1 bmi hyp chl
## 4 2 bmi hyp chl
## 4 3 bmi hyp chl
## 4 4 bmi hyp chl
## 4 5 bmi hyp chl
## 5 1 bmi hyp chl
## 5 2 bmi hyp chl
## 5 3 bmi hyp chl
## 5 4 bmi hyp chl
## 5 5 bmi hyp chl
```

```
## Warning: Number of logged events: 90
```

```
mnhanes.mult1 <- complete(mnhanes.multfit1)
mnhanes.mult1[is.na(mnhanes$bmi),]
```

```
## id age bmi hyp chl
## 1 1 1 35.3 1 118
## 3 3 1 22.0 1 187
## 4 4 3 35.3 1 118
## 6 6 3 24.9 1 184
## 10 10 2 24.9 1 206
## 11 11 1 27.2 1 187
## 12 12 2 29.6 1 184
## 16 16 1 26.3 1 199
## 21 21 1 24.9 2 229
```

Regression mean imputation

```
mice(mnhanes, method="norm.nob", m=1, maxit=1, seed=123)
```

```
##
```

```
## iter imp variable
```

```
## 1 1 bmi* hyp chl*
```

```
## Warning: Number of logged events: 6
```

```
## Class: mids
```

```
## Number of multiple imputations: 1
```

```
## Imputation methods:
```

```
## id age bmi hyp chl
## "" "" "norm.nob" "norm.nob" "norm.nob"
```

```
## PredictorMatrix:
```

```
## id age bmi hyp chl
## id 0 1 1 1 1
## age 1 0 1 1 1
```

```

## bmi 1 1 0 1 1
## hyp 1 1 1 0 1
## chl 1 1 1 1 0
## Number of logged events: 6
## it im dep meth
## 1 1 1 bmi norm.nob
## 2 1 1 bmi norm.nob
## 3 1 1 bmi norm.nob
## 4 1 1 chl norm.nob
## 5 1 1 chl norm.nob
## 6 1 1 chl norm.nob
##
## 1
## 2
## 3 mice detected that your data are (nearly) multi-collinear.\nIt applied a ridge per
## 4
## 5
## 6 mice detected that your data are (nearly) multi-collinear.\nIt applied a ridge per

Stochastic regression imputation: we sample from the fitted distribution
mice(mnhanes, method="norm.nob", m=1, maxit=1, seed=123)

##
## iter imp variable
## 1 1 bmi* hyp chl*

## Warning: Number of logged events: 6

## Class: mids
## Number of multiple imputations: 1
## Imputation methods:
## id age bmi hyp chl
## "" "" "norm.nob" "norm.nob" "norm.nob"
## PredictorMatrix:
## id age bmi hyp chl
## id 0 1 1 1 1
## age 1 0 1 1 1
## bmi 1 1 0 1 1
## hyp 1 1 1 0 1
## chl 1 1 1 1 0
## Number of logged events: 6
## it im dep meth
## 1 1 1 bmi norm.nob
## 2 1 1 bmi norm.nob
## 3 1 1 bmi norm.nob
## 4 1 1 chl norm.nob
## 5 1 1 chl norm.nob
## 6 1 1 chl norm.nob
##
## 1
## 2
## 3 mice detected that your data are (nearly) multi-collinear.\nIt applied a ridge per
## 4
## 5

```

```
## 6 mice detected that your data are (nearly) multi-collinear.\nIt applied a ridge penalty to continu
```


Chapter 10

Package mice

Don't impute a column by setting the method="" for that column

'blocks' are sets of variables, treated similarly

```
mice(msurf, m=5,  
      method=,  
      blocks=)
```

```
##  
## iter imp variable  
## 1 1 Income  
## 1 2 Income  
## 1 3 Income  
## 1 4 Income  
## 1 5 Income  
## 2 1 Income  
## 2 2 Income  
## 2 3 Income  
## 2 4 Income  
## 2 5 Income  
## 3 1 Income  
## 3 2 Income  
## 3 3 Income  
## 3 4 Income  
## 3 5 Income  
## 4 1 Income  
## 4 2 Income  
## 4 3 Income  
## 4 4 Income  
## 4 5 Income  
## 5 1 Income  
## 5 2 Income  
## 5 3 Income  
## 5 4 Income  
## 5 5 Income
```

```
## Warning: Number of logged events: 4
```

```
## Class: mids
```

```

## Number of multiple imputations: 5
## Imputation methods:
##      Personid      Gender Qualification      Age      Hours
##      " "          " "          " "          " "          " "
##      Income      Marital      Ethnicity
##      "pmm"       " "          " "
## PredictorMatrix:
##      Personid Gender Qualification Age Hours Income Marital Ethnicity
## Personid      0      0              0  1      1      1      0      0
## Gender        1      0              0  1      1      1      0      0
## Qualification  1      0              0  1      1      1      0      0
## Age           1      0              0  0      1      1      0      0
## Hours         1      0              0  1      0      1      0      0
## Income        1      0              0  1      1      0      0      0
## Number of logged events: 4
##   it im dep      meth      out
## 1  0  0      constant      Gender
## 2  0  0      constant Qualification
## 3  0  0      constant      Marital
## 4  0  0      constant      Ethnicity

```

Mean imputation/cell mean imputation: method = 'mean' Median imputation/cell median imputation: Hot deck cell imputation: method = 'sample' within cells Regression imputation

Chapter 11

Designing Surveys

The design of surveys is a massive topic, and we can't cover much here. The key principle is that everything that is done in the design of a survey should be done to meet the Objectives (the research goals) of the survey. Nothing else matters.

For our purposes calculation of the sample size required when designing a survey is the most significant single decision to make since it determines the standard errors for the survey estimates. In real life the decision about sample size is often severely constrained by available resources.

11.1 Another interpretation of the Deff

Assume that for sample size n the variance of an estimator in a specific design has the approximate formula

$$\mathbf{Var}[\widehat{T}_{\text{design}}] = \frac{c_{\text{design}}}{n}$$

then using the definition of the Deff, the ratio of the variance with respect to SRSWOR, then

$$\mathbf{Deff}[\widehat{T}_{\text{design}}] = \frac{\mathbf{Var}[\widehat{T}_{\text{design}}]}{\mathbf{Var}[\widehat{T}_{\text{SRSWOR}}]} \quad (11.1)$$

$$= \frac{c_{\text{design}}/n}{c_{\text{SRSWOR}}/n} \quad (11.2)$$

$$= \frac{c_{\text{design}}}{c_{\text{SRSWOR}}} \quad (11.3)$$

So if we want to find the sample sizes at which the variances in the design of interest and SRSWOR are equal, then we set

$$V_{\text{design}} = V_{\text{SRSWOR}} \quad (11.4)$$

$$\frac{c_{\text{design}}}{n_{\text{design}}} = \frac{c_{\text{SRSWOR}}}{n_{\text{SRSWOR}}} \quad (11.5)$$

$$n_{\text{design}} = \frac{c_{\text{design}}}{c_{\text{SRSWOR}}} n_{\text{SRSWOR}} \quad (11.6)$$

$$= \mathbf{Deff}[\widehat{T}_{\text{design}}] n_{\text{SRSWOR}} \quad (11.7)$$

This means that a rough way to compute the sample size required in a complex design is to first calculate the sample size we'd need in SRSWOR, and then scale it by the Deff for the complex design.

If the Deff is 2 (as it may be in a cluster design for example), then we need twice the sample size an SRSWOR would need in order to achieve the same standard error.

11.2 Computing the Required Sample Size

1. Identify your **design estimate** – this is the estimate that will underlie a headline you could imagine being written about your survey.
2. Identify the desired MOE, m , for this design estimate, noting whether you want this MOE to apply to the overall estimate, or for that estimate in each of, say, H subgroups/strata.
3. Compute the sample size required, n_1 , assuming that you're going to select a Simple Random Sample and ignoring the fpc.

For estimation of a mean $\bar{Y}$ of some characteristic y :

$$n_1 = \left(\frac{Z}{MOE} \right)^2 s^2$$

where Z is appropriate for the desired level of confidence and s is the estimated population standard deviation of the characteristic y .

For estimation of a proportion P :

$$n_1 = \left(\frac{Z}{MOE} \right)^2 p(1-p)$$

where p is the estimated proportion, or 0.5 if really unknown.

4. If the sample size is a substantial proportion of the population or subgroup size N , then apply the fpc adjustment:

$$n_2 = \frac{n_1}{1 + \frac{n_1}{N}}$$

5. If you have H subgroups and are using equal allocation then multiply by H

$$n_3 = Hn_2$$

or alternatively sum over the n_2 values in each of the strata you've defined.

6. Apply the design effect (Deff) appropriate for the actual design you'll be using.

$$n_4 = \text{Deff} \times n_3$$

7. Adjust for anticipated non-response: if response rate ϕ is expected then

$$n_5 = \frac{n_4}{\phi}$$

(If non-response is likely to vary between strata, then apply this correction to n_2 , within each stratum.)

References

Appendix A

Placeholder Appendix

Nothing to see here.

Bibliography

- Cochran, W. G. (1977). *Sampling Techniques*. Wiley, New York, third edition.
- Kish, L. (1965). *Survey Sampling*. Wiley, New York.
- Little, R. J. A. and Rubin, D. B. (2002). *Statistical Analysis with Missing Data*. Wiley, Hoboken, second edition.
- Lohr, S. L. (1999). *Sampling: Design and Analysis*. Duxbury Press, Pacific Grove, CA, USA.
- Lumley, T. (2004). Analysis of complex survey samples. *Journal of Statistical Software*, 9(1):1–19. R package version 2.2.
- Lumley, T. S. (2010). *Complex Surveys: A Guide to Analysis Using R (Wiley Series in Survey Methodology)*. Wiley.
- Särndal, C.-E., Swensson, B., and Wretmann, J. (1992). *Model Assisted Survey Sampling*. Springer-Verlag, New York.
- Scheaffer, R. L., III, W. M., and Ott, R. L. (2006). *Elementary Survey Sampling*. Duxbury, Belmont, sixth edition.
- Zealand, S. N. (1995). *A Guide to Good Survey Design*. Statistics New Zealand, Wellington, second edition.